

#### **Practical Concurrent and Parallel Programming VI**

### **Performance and Scalability**

Raúl Pardo

## Assignment 1 grading is completed



- Make sure you have a grade on assignment 1
  - If not, contact us!
- At the end of the semester, we will simply check that you got a 100 in at least 5 assignments

#### Oral feedback



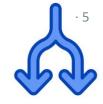
- Remember that during oral feedback sessions you will need to explain your solutions
  - So please revisit your solutions before the oral feedback session so that you remember what you did

#### Exam answers



- For the exam, scarce answers may be considered insufficient
- Answers should be self-contained
  - Make sure that you include the context and all your assumptions are clearly stated

## Agenda



- Performance versus scalability
- Scalability, speed-up and loss (of scalablity) classification
   Example: QuickSort
- Executors and Future
   Example: Dot Product
- Lock striping
  - A case study with Hash maps

## Previously on PCPP...



#### Week 5

Speedup for quicksort: 3.6 using 8 threads 2.9 using 4 threads

Speedup for counting primes: 3.9 using 8 threads 2.3 using 4 threads



# Performance versus scalability

## Performance and scalability

# \*\*

#### Performance (of software)

- Latency: time till first result (response time)
- Throughput: results per second

#### Scalability (one way to improve performance)

Improve throughput/latency by adding more resources

One may sacrifice performance for scalability Maybe OK to be slower on 1 core, if faster on 2 or 4 or ...

Goetz chapter 11

## What limits performance?



Suggestions?

## What limits performance?

# •10

#### **CPU-bound**

- Eg. counting prime numbers
- To speed up, add more CPUs (cores) (exploitation)

#### Input/output-bound

- Eg. reading from network
- To speed up, use more tasks (inherent)

#### **Synchronization-bound**

- Eg. Algorithm using shared data structure
- To speed up, improve shared data structure (Much of this lecture)



# Scalability, speed-up and loss classification

- Example: QuickSort

1 2 43 78 19 54 33 21 64 52 17 53



.13

- 1 2 43 78 19 54 33 21 64 52 17 53
- 1 2 43 78 19 54 33 21 64 52 17 53

•14

- 1 2 43 78 19 54 33 21 64 52 17 53
- 1 2 43 78 19 54 33 21 64 52 17 53
- 1 2 43 78 19 54 33 21 64 52 17 53



 1
 2
 43
 78
 19
 54
 33
 21
 64
 52
 17
 53

 1
 2
 43
 78
 19
 54
 33
 21
 64
 52
 17
 53

 1
 2
 43
 78
 19
 54
 33
 21
 64
 52
 17
 53

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022



1	2	43	78	19	54	33	21	64	52	17	53
1	2	43	78	19	54	33	21	64	52	17	53
1	2	<b>43</b>	78	19	54	33	21	64	52	17 <b>†</b>	53



1	2	43	78	19	54	33	21	64	52	17	53
1	2	43	78	19	54	33	21	64	52	17	53
1	2	43 <b>†</b>	78	19	54	33	21	64	52	17 <b>†</b>	53
1	2	17 <b>↑</b>	78	19	54	33	21	64	52	<b>43</b>	53



1	2	43	78	19	54	33	21	64	52	17	53
1	2	43	78	19	54	33	21	64	52	17	53
1	2	<b>43</b>	78	19	54	33	21	64	52	17 <b>†</b>	53
1	2	17 <b>†</b>	78	19	54	33	21	64	52	<b>43</b>	53
1	2	17	78 <b>†</b>	19	54	33	21	64	52	43	53



1	2	43	78	19	54	33	21	64	52	17	53
1	2	43	78	19	54	33	21	64	52	17	53
1	2	43 <b>†</b>	78	19	54	33	21	64	52	17 <b>†</b>	53
1	2	17 <b>†</b>	78	19	54	33	21	64	52	43 †	53
1	2	17	78 <b>†</b>	19	54	33	21 1	64	52	43	53

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022



1	2	43	78	19	<b>54</b>	33	21	64	52	17	53
1	2	43	78	19	54	33	21	64	52	17	53
1	2	43 †	78	19	54	33	21	64	52	<b>17</b>	53
1	2	17 <b>†</b>	78	19	54	33	21	64	52	<b>43</b>	53
1	2	17	78 <b>†</b>	19	54	33	21 †	64	52	43	53
1	2	17	21	19	33	54	78	64	52	43	53

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022



1	2	43	78	19	54	33	21	64	52	17	53
1	2	43	78	19	54	33	21	64	52	17	53
1	2	43 †	78	19	54	33	21	64	52	17 <b>†</b>	53
1	2	17 <b>†</b>	78	19	54	33	21	64	52	<b>43</b>	53
1	2	17	78 <b>†</b>	19	54	33	21 <b>†</b>	64	52	43	53
1	2	17	21	19	33	54	78	64	52	43	53

1 2 17 21 19 33 54 78 64 52 43 53

<sup>©</sup> Raúl Pardo Jimenez and Jørgen Staunstrup – F2022



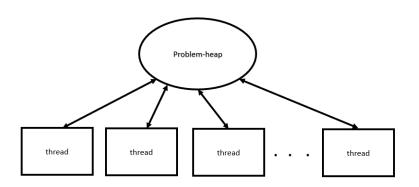
1	2	43	78	19	54	33	21	64	52	17	53	
1	2	43	78	19	54	33	21	64	52	17	53	
1	2	<b>43</b>	78	19	54	33	21	64	52	<b>17</b>	53	
1	2	17 <b>†</b>	78	19	54	33	21	64	52	<b>43</b>	53	
1	2	17	78 <b>†</b>	19	54	33	21 †	64	52	43	53	
1	2	17	21	19	33	54	78	64	52	43	53	

17 21 19 33 54 78 64 52 43 53

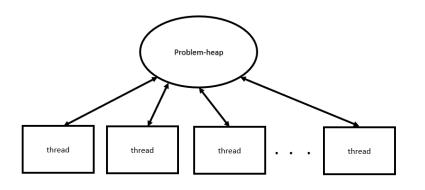
Two parts can be

sorted independently



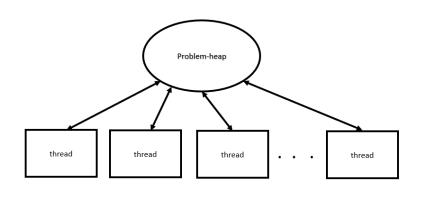






```
class Problem {
    public int[] arr;
    public int low, high;
    ...
    }
}
class ProblemHeap {
    list<Problem> heap= new List<Problem>;
    ...
}
```

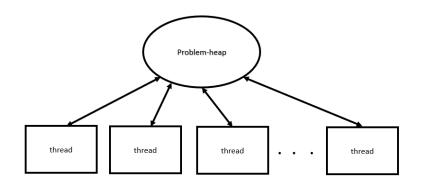




```
class Problem {
    public int[] arr;
    public int low, high;
    ...
    }
}
class ProblemHeap {
    list<Problem> heap= new List<Problem>;
    ...
}
```

```
private static void qsort(Problem problem, ProblemHeap heap) {
  int[] arr= problem.arr;
  int a= problem.low;
  int b= problem.high;
```

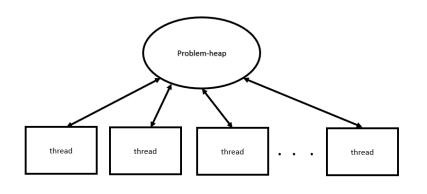




```
class Problem {
    public int[] arr;
    public int low, high;
    ...
    }
}
class ProblemHeap {
    list<Problem> heap= new List<Problem>;
    ...
}
```

```
private static void qsort(Problem problem, ProblemHeap heap) {
  int[] arr= problem.arr;
  int a= problem.low;
  int b= problem.high;
  ... // solve a subtask of quicksort
```

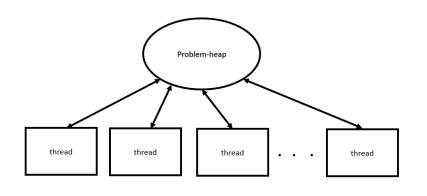




```
class Problem {
    public int[] arr;
    public int low, high;
    ...
    }
}
class ProblemHeap {
    list<Problem> heap= new List<Problem>;
    ...
}
```

```
private static void qsort(Problem problem, ProblemHeap heap) {
  int[] arr= problem.arr;
  int a= problem.low;
  int b= problem.high;
   ... // solve a subtask of quicksort
  heap.add(new Problem(arr, a, j); //qsort(arr, a, j);
```





```
class Problem {
    public int[] arr;
    public int low, high;
    ...
    }
}
class ProblemHeap {
    list<Problem> heap= new List<Problem>;
    ...
}
```

```
private static void qsort(Problem problem, ProblemHeap heap) {
  int[] arr= problem.arr;
  int a= problem.low;
  int b= problem.high;
    ... // solve a subtask of quicksort
  heap.add(new Problem(arr, a, j); //qsort(arr, a, j);
  heap.add(new Problem(arr, i, b));//qsort(arr, i, b);
}
```



```
public static void problemHeapStart(int threadCount, int pSize, int[] intArray) {
 ProblemHeap heap= new ProblemHeap(threadCount);
 heap.add(new Problem(intArray, 0, pSize-1));
 for (int t=0; t<threadCount; t++) {</pre>
```



```
public static void problemHeapStart(int threadCount, int pSize, int[] intArray) {
 ProblemHeap heap= new ProblemHeap(threadCount);
 heap.add(new Problem(intArray, 0, pSize-1));
 for (int t=0; t<threadCount; t++) {</pre>
       while (newProblem != null) { // when newProblem == null alg stops
         qsort(newProblem, heap);
         newProblem= heap.getProblem();
```



```
public static void problemHeapStart(int threadCount, int pSize, int[] intArray) {
 ProblemHeap heap= new ProblemHeap(threadCount);
 heap.add(new Problem(intArray, 0, pSize-1));
 for (int t=0; t<threadCount; t++) {</pre>
   threads[t] = new Thread( () -> { try {
       Problem newProblem= heap.getProblem();
       while (newProblem != null) { // when newProblem == null alg stops
         qsort(newProblem, heap);
         newProblem= heap.getProblem();
   });
```



```
public static void problemHeapStart(int threadCount, int pSize, int[] intArray) {
 ProblemHeap heap= new ProblemHeap(threadCount);
 heap.add(new Problem(intArray, 0, pSize-1));
 for (int t=0; t<threadCount; t++) {</pre>
   threads[t] = new Thread( () -> { try {
       Problem newProblem= heap.getProblem();
       while (newProblem != null) { // when newProblem == null alg stops
         gsort(newProblem, heap);
         newProblem= heap.getProblem();
       }catch (InterruptedException exn) { } //needed because getProblem may wait
   });
```



We use Mark8Setup to measure runtime

```
Benchmark.Mark8Setup("Problem heap quicksort",
         String.format("%2d", threadCount),
         new Benchmarkable() {
```

Goetz p. 224



We use Mark8Setup to measure runtime

```
Benchmark.Mark8Setup("Problem heap quicksort",
         String.format("%2d", threadCount),
         new Benchmarkable() {
           public void setup() {
              shuffle(intArray);
              problemHeapStart(threadCount, pSize, intArray);
```

Goetz p. 224



We use Mark8Setup to measure runtime

```
Benchmark.Mark8Setup("Problem heap quicksort",
         String.format("%2d", threadCount),
         new Benchmarkable() {
           public void setup() {
              shuffle(intArray);
              problemHeapStart(threadCount, pSize, intArray);
           public double applyAsDouble(int i) {
              problemHeapFinish(threadCount, intArray); return 0.0;
```

Goetz p. 224

• 36

We use Mark8Setup to measure runtime

```
Benchmark.Mark8Setup("Problem heap quicksort",
         String.format("%2d", threadCount),
         new Benchmarkable() {
           public void setup() {
              shuffle(intArray);
              problemHeapStart(threadCount, pSize, intArray);
           public double applyAsDouble(int i) {
              problemHeapFinish(threadCount, intArray); return 0.0;
```

Code in ProblemHeapSortingBenchmarkable.java

# -37

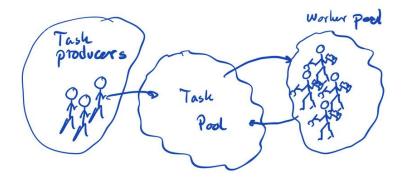
#### **Motivation**

- Threads are expensive to start executors reuse threads
- Problem heap breaking a problem down to smaller problems (tasks)

# -38

#### **Motivation**

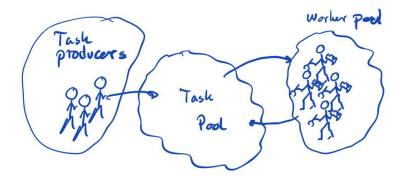
- Threads are expensive to start executors reuse threads
- Problem heap breaking a problem down to smaller problems (tasks)



# •39

#### **Motivation**

- Threads are expensive to start executors reuse threads
- Problem heap breaking a problem down to smaller problems (tasks)

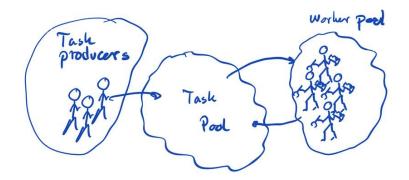


Task producers, and Workers are threads

# •40

#### **Motivation**

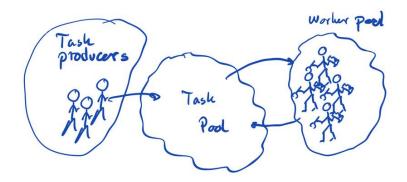
- Threads are expensive to start executors reuse threads
- Problem heap breaking a problem down to smaller problems (tasks)



Task producers, and Workers are threads
Workers may themselves produce new tasks

#### **Motivation**

- Threads are expensive to start executors reuse threads
- Problem heap breaking a problem down to smaller problems (tasks)



Task producers, and Workers are threads
Workers may themselves produce new tasks

The Task pool and Worker pool together is called an *Executor service* 

Reuse of threads

Thread(runnable1).start();

Thread(runnable2).start(); new Thread(runnable3).start();

ExecutorService pool;

Threads are expensive!

pool.execute(runnable1); pool.execute(runnable2);



# pool.execute(runnable2); IT UNIVERSITY OF COPENHAGEN





- Tasks are a central concept for executors
- When designing a program using executors, first think about the tasks to be executed
  - Like for threads, tasks can be conveniently defined in their own class
- Ideally, tasks should be independent



- Tasks are a central concept for executors
- When designing a program using executors, first think about the tasks to be executed
  - Like for threads, tasks can be conveniently defined in their own class
- Ideally, tasks should be independent





IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022



```
class QuicksortTask implements Runnable {
```



```
class QuicksortTask implements Runnable {
  private Problem p;
```



```
class QuicksortTask implements Runnable {
  private Problem p;
  private ExecutorService pool;
```



```
class QuicksortTask implements Runnable {
  private Problem p;
  private ExecutorService pool;
  private static int threshold;
```



```
class QuicksortTask implements Runnable {
  private Problem p;
  private ExecutorService pool;
  private static int threshold;
  ...
```



```
class QuicksortTask implements Runnable {
  private Problem p;
  private ExecutorService pool;
  private static int threshold;
  ...

@Override
  public void run() { qsort(p, pool, ...); }
```



```
class QuicksortTask implements Runnable {
  private Problem p;
  private ExecutorService pool;
  private static int threshold;
  ...

@Override
  public void run() { qsort(p, pool, ...); }

public static void qsort(ExecutorService pool, ...) {
```



```
class QuicksortTask implements Runnable {
  private Problem p;
  private ExecutorService pool;
  private static int threshold;
  ...

@Override
  public void run() { qsort(p, pool, ...); }

public static void qsort(ExecutorService pool, ...) {
  if (a < b) {</pre>
```



```
class QuicksortTask implements Runnable {
private Problem p;
private ExecutorService pool;
private static int threshold;
@Override
public void run() { qsort(p, pool, ...); }
public static void qsort(ExecutorService pool, ...) {
 if (a < b) {
   ... // quicksort on Problem p
  ... // subarray already ordered
```



```
class QuicksortTask implements Runnable {
private Problem p;
private ExecutorService pool;
private static int threshold;
 @Override
public void run() { qsort(p, pool, ...); }
public static void qsort(ExecutorService pool, ...) {
 if (a < b) {
   ... // quicksort on Problem p
   if ((j-a)>=threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
     // subarray already ordered
```



```
class QuicksortTask implements Runnable {
private Problem p;
private ExecutorService pool;
private static int threshold;
 @Override
public void run() { qsort(p, pool, ...); }
public static void qsort(ExecutorService pool, ...) {
 if (a < b) {
   ... // quicksort on Problem p
  if ((j-a)>=threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else qsort( new Problem(...);
  ... // subarray already ordered
```



```
class QuicksortTask implements Runnable {
private Problem p;
private ExecutorService pool;
private static int threshold;
 @Override
public void run() { qsort(p, pool, ...); }
public static void gsort(ExecutorService pool, ...) {
 if (a < b) {
   ... // quicksort on Problem p
  if ((j-a)>=threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else gsort( new Problem(...);
  if ((b-i)>= threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else qsort(new Problem(...);
  ... // subarray already ordered
```



```
class QuicksortTask implements Runnable {
private Problem p;
                                                         Are Quicksort tasks
private ExecutorService pool;
private static int threshold;
                                                        independent of each
                                                                 other?
@Override
public void run() { qsort(p, pool, ...); }
public static void qsort(ExecutorService pool, ...) {
 if (a < b) {
   ... // quicksort on Problem p
  if ((j-a)>=threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else gsort( new Problem(...);
  if ((b-i)>= threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else qsort(new Problem(...);
 ... // subarray already ordered
```



```
class QuicksortTask implements Runnable {
private Problem p;
                                                        Are Quicksort tasks
private ExecutorService pool;
private static int threshold;
                                                       independent of each
                                                                other?
@Override
public void run() { qsort(p, pool, ...); }
public static void qsort(ExecutorService pool, ...) {
 if (a < b) {
   ... // quicksort on Problem p
  if ((j-a)>=threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else gsort( new Problem(...);
  if ((b-i)>= threshold) pool.execute(new QuicksortTask(new Problem(...), pool, ...));
  else qsort(new Problem(...);
 ... // subarray already ordered
                                                        Code in QuicksortTask.java
```

#### Quicksort Executor class

- •60
- Kick-off class for the program
- It initializes the Executor service

- 61
- Kick-off class for the program
- It initializes the Executor service

- There are several type of thread pools:
  - newFixedThreadPool
  - newCachedThreadPool (useful for testing)
  - newSingleThreadExecutor
  - ..



- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention



- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention

Thread 1



- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention

Thread 1

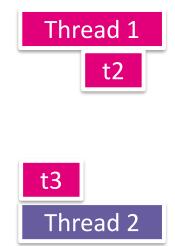


- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention





- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention





- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention

Thread 1



- Instead of having a single task pool, each worker has its own task pool
  - This tasks pools are implemented as queues
- Workers add tasks to their own queues
  - Minimizes thread contention in a single thread resource
- If worker can work but its queue is empty, it steals a task from another worker task pool
  - From the end of the queue
  - To further minimize contention

Thread 1

Thread 2

In what type of algorithms may work stealing queues be useful?

#### Performance of Executor Quicksort



```
Executor quicksort 1 98003405.0 ns
Executor quicksort 2 53568593.9 ns
Executor quicksort 4 36397241.3 ns
Executor quicksort 8 21714103.7 ns
Executor quicksort 16 22237307.4 ns
Executor quicksort 32 22510681.9 ns
```

Speedup = 4.5

A bit better speed-up than using native Threads (slide 3)

#### Quicksort



1	2	43	78	19	54	33	21	64	52	17	53	
1	2	43	78	19	54	33	21	64	52	17	53	
1	2	43 †	78	19	54	33	21	64	52	<b>17</b>	53	
1	2	17 <b>†</b>	78	19	54	33	21	64	52	<b>43</b>	53	
1	2	17	78 <b>†</b>	19	54	33	21 †	64	52	43	53	
1	2	17	21	19	54	33	78	64	52	43	53	

1 2 17 21 19 33 54 78 64 52 43 53

Two parts can be

sorted independently

# What limits scalability?





#### Example: growing a crop

- 4 months growth + 1 month harvest if done by 1 person
- Growth (sequential) cannot be speeded up
- Using 30 people to harvest, takes 1/30 month = 1 day
- Speed-up using many harvesters: 5/(4+1/30) = 1.24 times faster

# What limits scalability?



#### Example: growing a crop

- 4 months growth + 1 month harvest if done by 1 person
- Growth (sequential) cannot be speeded up
- Using 30 people to harvest, takes 1/30 month = 1 day
- Speed-up using many harvesters: 5/(4+1/30) = 1.24 times faster

#### Amdahl's law (Goetz 11.2)

$$F$$
 = sequential fraction of problem =  $4/5 = 0.8$ 

$$N = \text{number of threads (people)} = 30$$

speed 
$$up \le \frac{1}{F + \frac{(1-F)}{N}} = \frac{1}{0.8 + 0.2/30} = 1.24$$



- Starvation loss
  - Minimize the time that the problem heap is empty

# Other types of "loss" of scalability



- Starvation loss
  - Minimize the time that the problem heap is empty
- Separation loss (best threshold)
  - Find a good threshold to distribute workload evenly

# Other types of "loss" of scalability



- Starvation loss
  - Minimize the time that the problem heap is empty
- Separation loss (best threshold)
  - Find a good threshold to distribute workload evenly
- Saturation loss (locking common data structure)
  - Minimize high thread contention in the problem

# Other types of "loss" of scalability



- Starvation loss
  - Minimize the time that the problem heap is empty
- Separation loss (best threshold)
  - Find a good threshold to distribute workload evenly
- Saturation loss (locking common data structure)
  - Minimize high thread contention in the problem
- Braking loss
  - Stop all threads as soon as the problem is solved



- Starvation loss
  - Minimize the time that the problem heap is empty
- Separation loss (best threshold)
  - Find a good threshold to distribute workload evenly
- Saturation loss (locking common data structure)
  - Minimize high thread contention in the problem
- Braking loss
  - Stop all threads as soon as the problem is solved

Møller-Nielsen, P and Staunstrup, J, Problem-heap. A paradigm for multiprocessor algorithms. *Parallel Computing*, 4:63-74, 1987



# **Executors and Future**

- Example: Dot Product

```
-80
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
```

```
-81
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
class QuicksortTask implements Runnable {
    public void run() { qsort(p, ...); }
}
```

```
- 82
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
class QuicksortTask implements Runnable {
    public void run() { qsort(p, ...); }
}
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
```



```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
class QuicksortTask implements Runnable {
    public void run() { qsort(p, ...); }
}
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
```

```
public class QuicksortExecutor {
   //main (thread)
   setUpQS( ... ) // Initializes pool, executes initial task
   finishQS( ... ) //wait for all tasks to be completed
}
```



```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
class QuicksortTask implements Runnable {
    public void run() { qsort(p, ...); }
}
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
```

```
public class QuicksortExecutor {
   //main (thread)
   setUpQS( ... ) // Initializes pool, executes initial task
   finishQS( ... ) //wait for all tasks to be completed
}
```

How do we known when all tasks are finished?



```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
class QuicksortTask implements Runnable {
    public void run() { qsort(p, ...); }
}
```

```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
```

```
public class QuicksortExecutor {
   //main (thread)
   setUpQS( ... ) // Initializes pool, executes initial task
   finishQS( ... ) //wait for all tasks to be completed
}
```

How do we known when all tasks are finished?

Need for a signal from the workers to the main thread!

# Shut down



The ExecutorService has methods to shut down the pool.

```
// Executor body
...
pool.shutdown();
```

# Shut down



The ExecutorService has methods to shut down the pool.

```
// Executor body
...
pool.shutdown();
```

The challenge is when to shut down

# Shut down



The ExecutorService has methods to shut down the pool.

```
// Executor body
...
pool.shutdown();
```

The challenge is when to shut down

In the Quicksort example, a counter is used to track the number of pending tasks. When there are no pending tasks we can shut down.

## Solution 1: shared counter + barrier



# Using the counter + barrier



```
public static void qsort(.., CyclicBarrier done, AtomicInteger count) {
if (a < b) {
  if ((j-a)>= threshold) count.incrementAndGet();
  if ((b-i)>= threshold) count.incrementAndGet();
  if ((j-a) >= threshold) {
    pool.execute(new QuicksortTask(new Problem(arr, a, j), pool, c) );
  } else qsort(...) // sequentially
  if ((b-i)>= threshold) {
    pool.execute(new solveProblem(new Problem(arr, i, b), pool, c) );
  } else qsort(...) // sequentially
  if (count.decrementAndGet() == 0) { done.await(); pool.shutdown(); }
```



```
public static void qsort(.., CyclicBarrier done, AtomicInteger count) {
if (a < b) {
  if ((j-a)>= threshold) count.incrementAndGet();
  if ((b-i)>= threshold) count.incrementAndGet();
                                                       Why not incrementing in the
                                                        branch creating the task?
  if ((j-a) \ge threshold) {
    pool.execute(new QuicksortTask(new Problem(arr, a, j), pool, c) );
  } else qsort(...) // sequentially
  if ((b-i)>= threshold) {
    pool.execute(new solveProblem(new Problem(arr, i, b), pool, c) );
  } else qsort(...) // sequentially
  if (count.decrementAndGet() == 0) { done.await(); pool.shutdown(); }
```



```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
class QuicksortTask implements Runnable {
    public void run() { qsort(p, ...); }
}
```

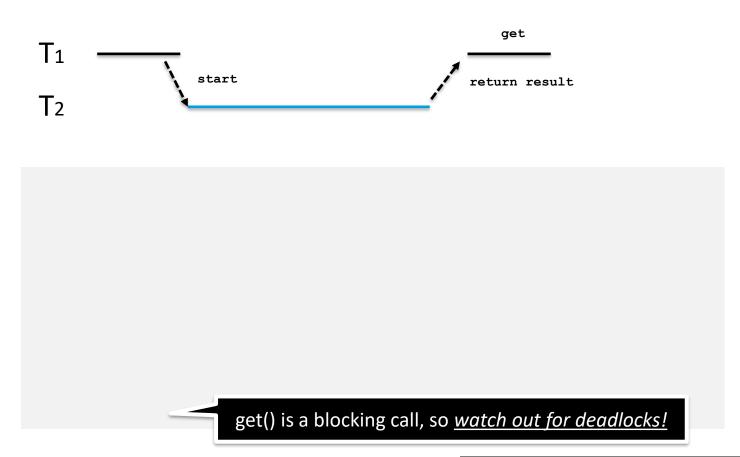
```
class QuicksortTask implements Runnable {
  public void run() { qsort(p, ...); }
}
```

```
public class QuicksortExecutor {
   //main (thread)
   setUpQS( ... ) // Initializes pool, executes initial task
   finishQS( ... ) //wait for all tasks to be completed
}
```

```
private static void finishQS(CyclicBarrier done) {
  try { done.await(); }
  catch (InterruptedException | BrokenBarrierException e) { ... ); }
}
```

# Solution 2: Future





## Solution 2: Future





## Solution 2: Future



```
T1 start return result
```



Code in **futureExample.java** 



Both are used to specify the code of a thread.

\_\_\_\_



Both are used to specify the code of a thread.

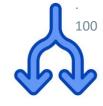
- Runnable cannot return a result
  - Overrides run()
- <u>Callable</u> returns a result (via a Future)
  - Overrides call()



Both are used to specify the code of a thread.

- Runnable cannot return a result
  - Overrides run()
- <u>Callable</u> returns a result (via a Future)
  - Overrides call()

As illustrated by the Quicksort example, Runnables may use shared data (e.g., to deliver a result)



Both are used to specify the code of a thread.

- Runnable cannot return a result
  - Overrides run()
- <u>Callable</u> returns a result (via a Future)
  - Overrides call()

As illustrated by the Quicksort example, Runnables may use shared data (e.g., to deliver a result)

Futures are an example of message passing (coming weeks)



Both are used to specify the code of a thread.

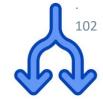
- Runnable cannot return a result
  - Overrides run()
- <u>Callable</u> returns a result (via a Future)
  - Overrides call()

Could Callables use shared data as well?

As illustrated by the Quicksort example, Runnables may use shared data (e.g., to deliver a result)

Futures are an example of message passing (coming weeks)

# dotProduct – Task class

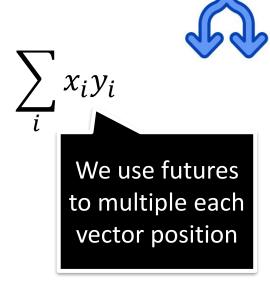


Given two vectors x, y of equal size, their dot product equals  $\sum x_i y_i$ 

$$\sum_{i} x_i y_i$$

# dotProduct – Task class

Given two vectors x, y of equal size, their dot product equals



## dotProduct – Task class

Given two vectors x, y of equal size, their dot product equals

```
public class DotProductTask implements Callable<Integer> {
    final int pos;
    final int[] x, y;
    public DotProductTask(int[] x, int[] y, int pos) {
        this.x
                 = x;
        this.y = y;
        this.pos = pos;
    @Override
    public Integer call() {
        return x[pos] * y[pos];
```

```
\sum_{i} x_{i}y_{i}
We use futures to multiple each vector position
```



Given two vectors x, y of equal size, their dot product equals  $\sum x_i y_i$ 

```
...
List<DotProductTask> tasks = new ArrayList<DotProductTask>();

// Randomly initialize arrays ...
```



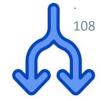
Given two vectors x, y of equal size, their dot product equals  $\sum_{i=1}^{n} x_i y_i$ 

```
List<DotProductTask> tasks = new ArrayList<DotProductTask>();
// Randomly initialize arrays ...
// Create the list of tasks (Futures) to execute
for (int i = 0; i < N; i++)
    tasks.add(new DotProductTask(x,y,i));
```



Given two vectors x, y of equal size, their dot product equals  $\sum_{i=1}^{n} x_i y_i$ 

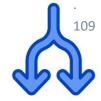
```
List<DotProductTask> tasks = new ArrayList<DotProductTask>();
// Randomly initialize arrays ...
// Create the list of tasks (Futures) to execute
for (int i = 0; i < N; i++)
    tasks.add(new DotProductTask(x,y,i));
// Add all futures to the execution pool at once
List<Future<Integer>> futures = pool.invokeAll(tasks);
```



Given two vectors x, y of equal size, their dot product equals  $\sum_{i=1}^{n} x_i y_i$ 

```
List<DotProductTask> tasks = new ArrayList<DotProductTask>();
// Randomly initialize arrays ...
// Create the list of tasks (Futures) to execute
for (int i = 0; i < N; i++)
    tasks.add(new DotProductTask(x,y,i));
. . .
// Add all futures to the execution pool at once
List<Future<Integer>> futures = pool.invokeAll(tasks);
for (Future<Integer> f : futures) {
     result += f.get(); // Wait for each future to be executed
                        // and add partial result
```

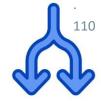
#### dotProduct – Future execution



Given two vectors x, y of equal size, their dot product equals  $\sum_{i=1}^{n} x_i y_i$ 

```
List<DotProductTask> tasks = new ArrayList<DotProductTask>();
// Randomly initialize arrays ...
// Create the list of tasks (Futures) to execute
for (int i = 0; i < N; i++)
    tasks.add(new DotProductTask(x,y,i));
. . .
// Add all futures to the execution pool at once
List<Future<Integer>> futures = pool.invokeAll(tasks);
for (Future<Integer> f : futures) {
     result += f.get(); // Wait for each future to be executed
                        // and add partial result
pool.shutdown(); // We are sure to be done, so we shut down the pool
```

#### dotProduct – Future execution



Given two vectors x, y of equal size, their dot product equals  $\sum_{i=1}^{n} x_i y_i$ 

```
List<DotProductTask> tasks = new ArrayList<DotProductTask>();
// Randomly initialize arrays ...
// Create the list of tasks (Futures) to execute
for (int i = 0; i < N; i++)
    tasks.add(new DotProductTask(x,y,i));
. . .
// Add all futures to the execution pool at once
List<Future<Integer>> futures = pool.invokeAll(tasks);
for (Future<Integer> f : futures) {
     result += f.get(); // Wait for each future to be executed
                        // and add partial result
                                                                 Code in FuturesDotProduct.java
pool.shutdown(); // We are sure to be done, so we shut down the pool
```



# Lock striping A case study with Hash maps



A *collection* is simply an object that groups multiple elements into a single unit

Package: java.util



A collection is simply an object that groups multiple elements into a single unit

Package: java.util

Examples: ArrayList, HashMap, TreeSet, ...

https://docs.oracle.com/javase/tutorial/collections/intro/index.html



A *collection* is simply an object that groups multiple elements into a single unit

Package: java.util

Examples: ArrayList, HashMap, TreeSet, ...

https://docs.oracle.com/javase/tutorial/collections/intro/index.html

Methods: add, remove, size, contains, ...



A *collection* is simply an object that groups multiple elements into a single unit

Package: java.util

Examples: ArrayList, HashMap, TreeSet, ...

https://docs.oracle.com/javase/tutorial/collections/intro/index.html

Methods: add, remove, size, contains, ...

Many of the classes have synchronized/concurrent implementations

https://www.baeldung.com/java-synchronized-collections



```
import java.util.*;
public class syncCollectionExample {
 public static void main(String[] args) {
                                            new syncCollectionExample(); }
 public syncCollectionExample() {
```



```
import java.util.*;
public class syncCollectionExample {
  public static void main(String[] args) {    new syncCollectionExample(); }
  public syncCollectionExample() {
    ArrayList<String> a= new ArrayList<String>();
    a.add("A"); ...
```



```
import java.util.*;
public class syncCollectionExample {
  public static void main(String[] args) {    new syncCollectionExample(); }
  public syncCollectionExample() {
    ArrayList<String> a= new ArrayList<String>();
    a.add("A"); ...
    Collection<String> synColl = Collections.synchronizedCollection(a);
```

```
119
```

```
import java.util.*;
public class syncCollectionExample {
  public static void main(String[] args) {    new syncCollectionExample(); }
  public String getLast(ArrayList<String> 1) {
    int last= 1.size()-1;
    return l.get(last);
  public syncCollectionExample() {
    ArrayList<String> a= new ArrayList<String>();
    a.add("A"); ...
    Collection<String> synColl = Collections.synchronizedCollection(a);
```



```
import java.util.*;
public class syncCollectionExample {
  public static void main(String[] args) {    new syncCollectionExample(); }
  public String getLast(ArrayList<String> 1) {
    int last= 1.size()-1;
    return l.get(last);
  public static void delete(ArrayList<String> 1) {
    int last= 1.size()-1;
    1.remove(last);
  public syncCollectionExample() {
    ArrayList<String> a= new ArrayList<String>();
    a.add("A"); ...
    Collection<String> synColl = Collections.synchronizedCollection(a);
```

# Thread-safety (from week 3)



It is very important to note that for a program p:

p only accesses thread-safe classes

⇒

p is a thread-safe program

#### Making the synchronized ArrayList thread safe



```
import java.util.*;
public class syncCollectionExample {
  public static void main(String[] args) {
                                            new syncCollectionExample(); }
  public String getLast(ArrayList<String> 1) {
    synchronized(1) {
      int last= 1.size()-1;
      return l.get(last);
  public static void delete(ArrayList<String> 1) {
    synchronized(1) {
      int last= 1.size()-1;
      1.remove(last);
  public syncCollectionExample() {
```





and used by many threads?



and used by many threads?

for example:

a bank

Facebook updates

• • •



and used by many threads?

for example:

a bank

Facebook updates

• • •

Would not work if everything I "synchronized"



and used by many threads?

for example:

a bank

Facebook updates

• • •

Would not work if everything I "synchronized"

What can we do?



and used by many threads?

for example:

a bank

Facebook updates

• • •

Would not work if everything I "synchronized"

What can we do?

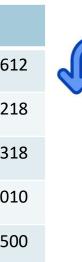
Reduce locking !!



Key value pairs: <k1, v1>, <k2, v2>, ...

Key value pairs: <k1, v1>, <k2, v2>, ...

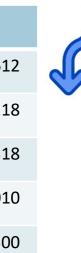
Key	Value	
Peter	20487612	(
Anna	51251218	
_ena	34458318	
Holger	89545010	
_isa	94959500	



Key value pairs: <k1, v1>, <k2, v2>, ...

```
class HashMap<K,V> {
  ... // data structure
 public V get(K k) { ... }
 public V put(K k, V v) { ... }
 public boolean containsKey(K k) { ... }
 public int size() { return cachedSize; }
 public V remove(K k) { ... }
```

Key	Value	
Peter	20487612	
Anna	51251218	
-ena	34458318	
Holger	89545010	
_isa	94959500	



Key value pairs: <k1, v1>, <k2, v2>, ...

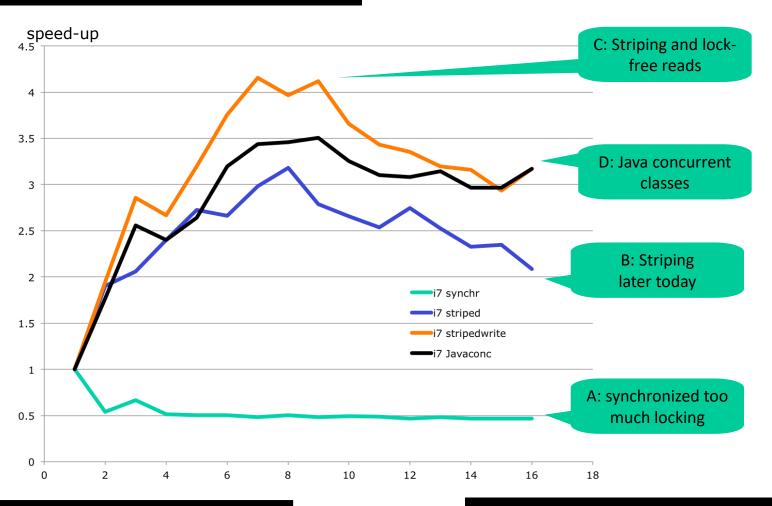
```
class HashMap<K,V> {
    ... // data structure
    public V get(K k) { ... }
    public V put(K k, V v) { ... }
    public boolean containsKey(K k) { ... }
    public int size() { return cachedSize; }
    public V remove(K k) { ... }
    ...
}
```

How to make it thread-safe?

еу	Value	
eter	20487612	
nna	51251218	
ena	34458318	
olger	89545010	
isa	94959500	

# Scaling a HashMap

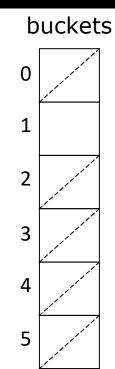




IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022

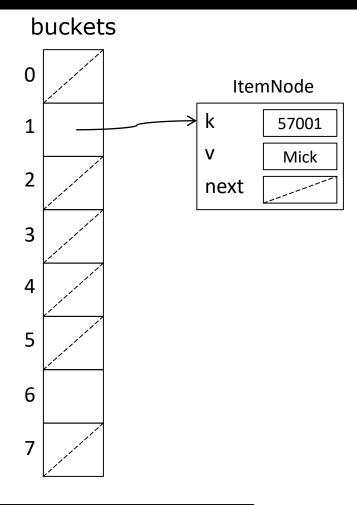




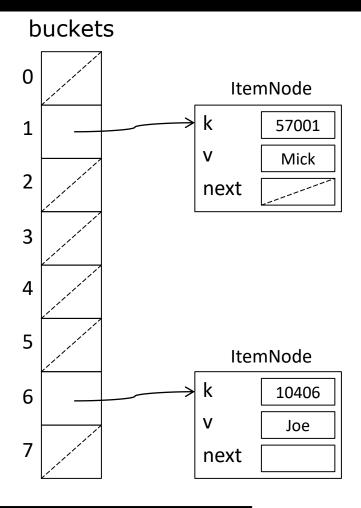
6

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022

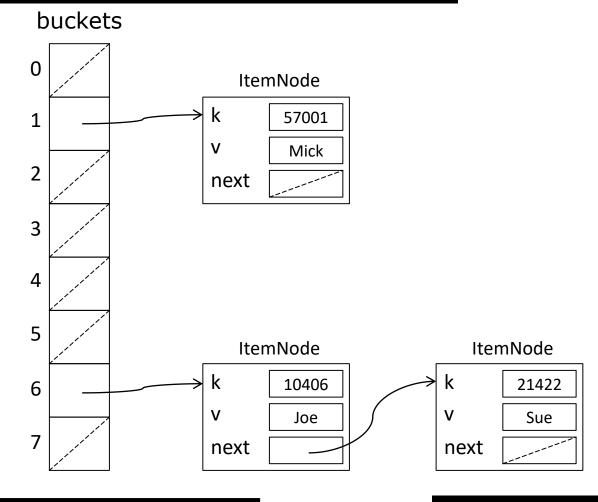




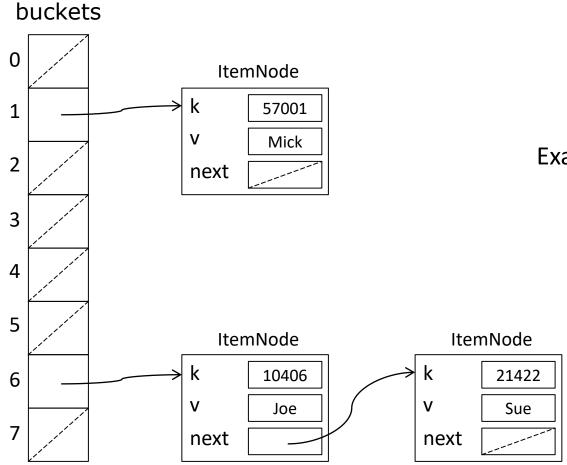






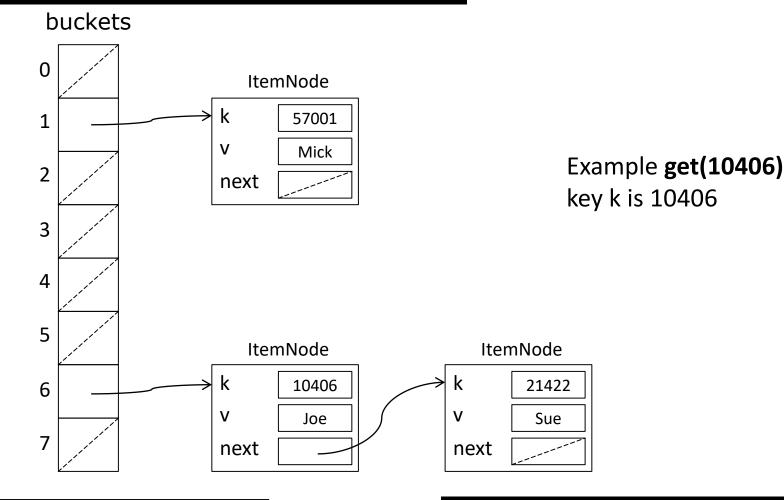




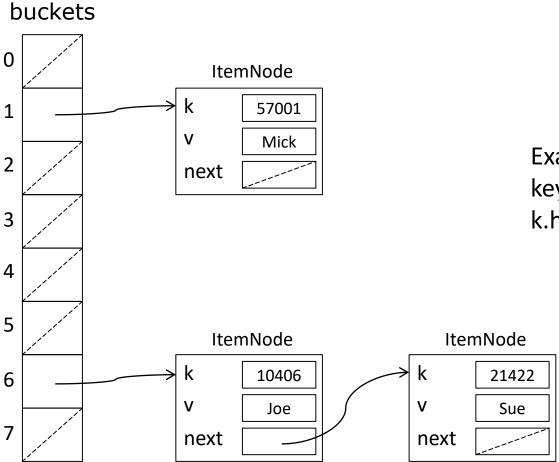


Example get(10406)





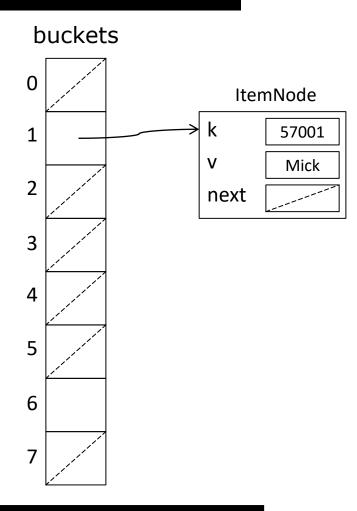




Example get(10406) key k is 10406 k.hashCode() is 6

# HaspMap put

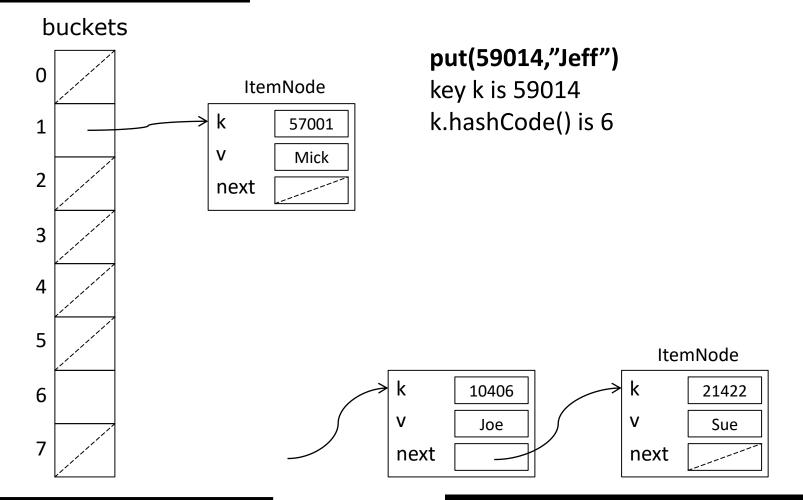




© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022

## HaspMap put



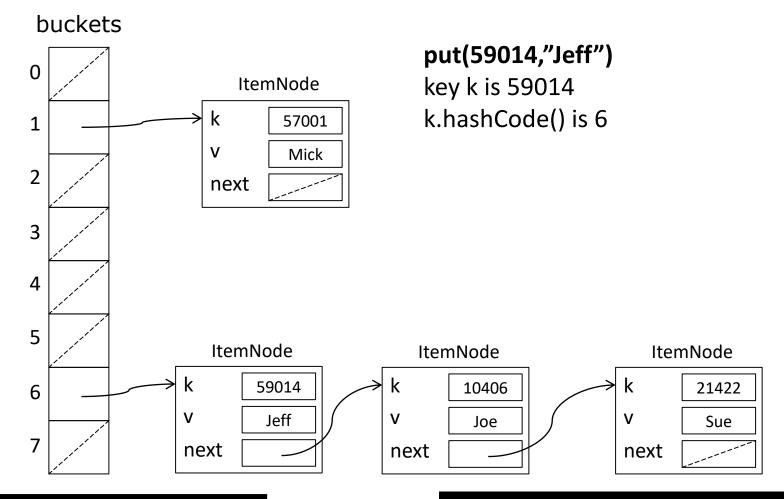


IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022

# HaspMap put





### Synchronized implementation



```
static class ItemNode<K,V> {
   private final K k;
   private V v;
   private ItemNode<K,V> next;
   public ItemNode(K k, V v, ItemNode<K,V> next) { ... }
}
```

```
class SynchronizedMap<K,V> {
   public synchronized V get(K k) { ... }
   public synchronized boolean containsKey(K k) { ... }
   public synchronized int size() { return cachedSize; }
   public synchronized V put(K k, V v) { ... }
   public synchronized V remove(K k) { ... }
}
```

### Synchronized implementation



```
static class ItemNode<K,V> {
   private final K k;
   private V v;
   private ItemNode<K,V> next;
   public ItemNode(K k, V v, ItemNode<K,V> next) { ... }
}
```

```
class SynchronizedMap<K,V> {
  private ItemNode<K,V>[] buckets; // guarded by this
  private int cachedSize; // guarded by this
  public synchronized V get(K k) { ... }
  public synchronized boolean containsKey(K k) { ... }
  public synchronized int size() { return cachedSize; }
  public synchronized V put(K k, V v) { ... }
  public synchronized V remove(K k) { ... }
}
```

- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)



- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock

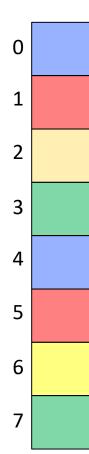


- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks



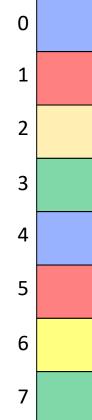
- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks
- -guard every 4<sup>th</sup> bucket with the same lock

- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks
- -guard every 4<sup>th</sup> bucket with the same lock

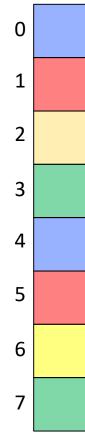




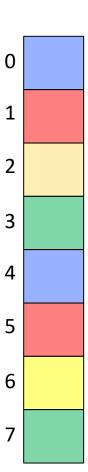
- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks
- -guard every 4<sup>th</sup> bucket with the same lock
- -locks[0] guards bucket 0, 4, 8, ... (stripe 0)



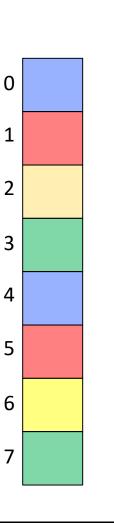
- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks
- -guard every 4<sup>th</sup> bucket with the same lock
- -locks[0] guards bucket 0, 4, 8, ... (stripe 0)
- -locks[1] guards bucket 1, 5, 9, ... (stripe 1) et



- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks
- -guard every 4<sup>th</sup> bucket with the same lock
- -locks[0] guards bucket 0, 4, 8, ... (stripe 0)
- -locks[1] guards bucket 1, 5, 9, ... (stripe 1) et
- -With high probability
- -two operations will work on different stripes

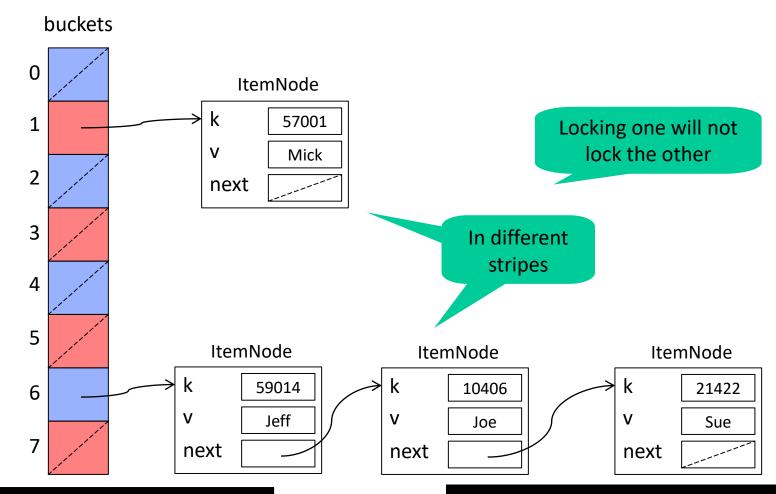


- Guarding the table with a single lock works
- -... but does not scale well (actually **very** badly)
- Idea: Each bucket could have its own lock
- In practice
- -use fewer, to illustrate we use 4, locks
- -guard every 4<sup>th</sup> bucket with the same lock
- -locks[0] guards bucket 0, 4, 8, ... (stripe 0)
- -locks[1] guards bucket 1, 5, 9, ... (stripe 1) et
- -With high probability
- -two operations will work on different stripes
- -hence will take different locks
- Less lock contention, better scalability



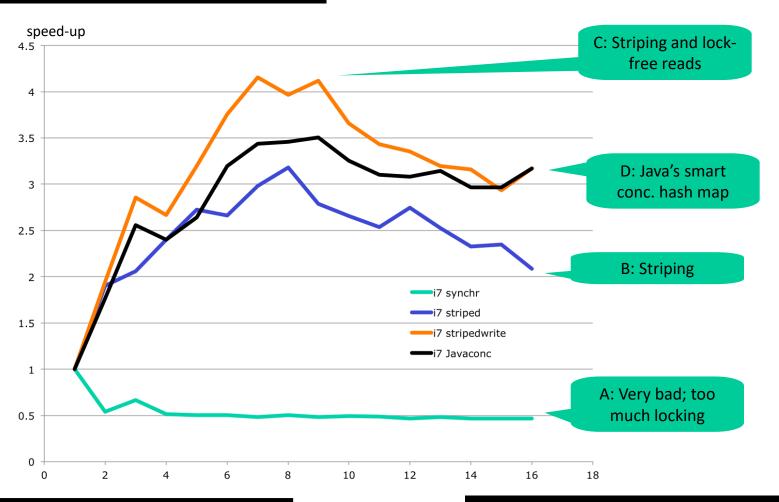
#### Bucket idea





# Reducing locking





IT UNIVERSITY OF COPENHAGEN © Raúl Pardo Jimenez and Jørgen Staunstrup – F2022

#### Ultimate scalability



A web-shop, Facebook, ...

We must give up thread safety,

but still maintain some sort of consistency

#### Ultimate scalability



A web-shop, Facebook, ...

We must give up thread safety,

but still maintain some sort of consistency

Week 13