



# Practical Concurrent and Parallel Programming X

## Lock-Free Data Structures

Raúl Pardo

Based on slides of PCPP 2019

# Troubles running Gradle?




<https://www.menti.com/alkyczp33jsu>




# Previously on PCPP...

## Reducing locking



IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2022





- Compare-And-Swap (CAS)
  - Lock-free Counter
  - Atomic libraries
  - CAS based lock implementation
- Lock-free stack
- Lock-free queue
- ABA problem

# HW support for atomic compound operations



- Early processors had atomic compound operations to implement mutexes
  - *test-and-set*
  - Not suitable for implementing advanced lock-free data structures
- Modern processors provide special instructions to for managing concurrent access to shared variables
  - *store-conditional*, **compare-and-swap (CAS)**

See Herlihy Sections 5.1 – 5.8  
(outside the scope of the course)



# Compare-And-Swap (CAS)



- **`v.compareAndSwap(a, b)`**
  - Compares the value of `v` and `a`, and, if they are equal `v` is set to `b`, otherwise it does nothing. In either case, it returns the current value in `v`, i.e., the value when CAS was executed
  - This instruction is executed atomically by the hardware

# CAS pseudo-code

.7

```
class MyAtomicInteger {
    private int value;    // Visibility ensured by locking

    public synchronized int compareAndSwap(int oldValue, int newValue) {
        int valueInRegister = this.value;
        if (this.value == oldValue)
            this.value = newValue;
        return valueInRegister;
    }

    public synchronized boolean compareAndSet(int oldValue, int newValue) {
        return oldValue == this.compareAndSwap(oldValue, newValue);
    }

    public synchronized int get() { return this.value; }
    ...
}
```

Illustrative implementation of CAS

Common alternative.  
Also abbreviated as CAS  
(optional exercise: implement  
natively compareAndSet)

- Only to *illustrate* CAS semantics
  - Internally, locks are implemented using CAS
  - Not the other way around

# AtomicInteger operations via CAS



- Standard AtomicInteger operations can now be implemented using the CAS operations *without blocking*
- This is an example of *optimistic* concurrency (or non-blocking computation)
  - In a nutshell, trying several times until the operation succeeds

```
class MyAtomicInteger {  
...  
    public int addAndGet(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return newValue;  
    }  
  
    public int getAndAdd(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return oldValue;  
    }  
...  
}
```



# AtomicInteger operations via CAS



- Standard AtomicInteger operations can now be implemented using the CAS operations *without blocking*
- This is an example of *optimistic* concurrency (or non-blocking computation)
  - In a nutshell, trying several times until the operation succeeds

```
class MyAtomicInteger {  
    ...  
    public int addAndGet(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return newValue;  
    }  
  
    public int getAndAdd(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return oldValue;  
    }  
    ...  
}
```

Is this the same  
as busy-wait?



- *Busy-wait* is an *alternative to blocking* a thread to wait until some condition holds or to enter the critical section
- The main difference with `lock()` or `await()` is that the thread does not transition to the “blocked” state
- Generally, busy-wait is a bad idea,
  - Threads may consume computing resources to check a condition that has not been updated
  - In this course, we will never ask you to use busy-wait
  - Exercise solutions using busy-wait will be rejected
- Very rarely busy-wait may be preferred over blocking the thread
  - When the thread waits for a very short time it might be more efficient to use busy-wait
  - However, as we have discussed, reasoning about how it takes for a thread to do anything is pointless in concurrency

```
...
// state variables
int i = 0;
Lock l = new ReentrantLock();
...

// method example
public void method(...) {
    l.lock()
    try{
        // busy-wait
        while(i>0) {
            // do nothing
        }
    }
    catch (InterruptedException e) {...}
    finally {l.unlock();}
}
```

# Progress conditions for non-blocking DS

· 10




There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
  - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)

# Progress conditions for non-blocking DS

· 10



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
  - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)
2. Lock-free: A method of an object implementation is lock-free if executing the method guarantees that some method call (including concurrent) finishes in a finite number of steps
  - Somebody's operations are guaranteed to complete in a bounded number of my steps
  - Most non-blocking data structures are lock-free, e.g., Treap's stack (see next slides)



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
  - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)
2. Lock-free: A method of an object implementation is lock-free if executing the method guarantees that some method call (including concurrent) finishes in a finite number of steps
  - Somebody's operations are guaranteed to complete in a bounded number of my steps
  - Most non-blocking data structures are lock-free, e.g., Treap's stack (see next slides)
3. Obstruction-free: A method of an object implementation is obstruction-free if, from any point after which it executes in isolation, it finishes in a finite number of steps;
  - My operations are guaranteed to complete in a bounded number of steps (if I get to execute them)



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
  - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)
2. Lock-free: A method of an object implementation is lock-free if executing the method guarantees that some method call (including concurrent) finishes in a finite number of steps
  - Somebody's operations are guaranteed to complete in a bounded number of my steps
  - Most non-blocking data structures are lock-free, e.g., Treap's stack (see next slides)
3. Obstruction-free: A method of an object implementation is obstruction-free if, from any point after which it executes in isolation, it finishes in a finite number of steps;
  - My operations are guaranteed to complete in a bounded number of steps (if I get to execute them)

$$wait_{free} \Rightarrow lock_{free} \Rightarrow obstruction_{free}$$

# Why non-blocking $\neq$ busy-wait?



- Threads cannot wait forever because other thread finished incorrectly
  - Even in obstruction-free, completion must be guaranteed when the thread runs in isolation
- All progress definitions for non-blocking programs are part of an algorithm composed by multiple threads
  - Threads collaborate with each other



- Standard AtomicInteger operations can now be implemented using the CAS

What type of non-blocking progress applies to this AtomicInteger?

- In a nutshell, trying several times until the operation succeeds

```
class MyAtomicInteger {
...
    public int addAndGet(int delta) {
        int oldValue, newValue;
        do {
            oldValue = get();
            newValue = oldValue + delta;
        } while (!compareAndSet(oldValue, newValue));
        return newValue;
    }

    public int getAndAdd(int delta) {
        int oldValue, newValue;
        do {
            oldValue = get();
            newValue = oldValue + delta;
        } while (!compareAndSet(oldValue, newValue));
        return oldValue;
    }
...
}
```





- In Java, **AtomicXX** CAS operations have the same memory semantics of volatile read/write
  - <https://github.com/AdoptOpenJDK/openjdk-jdk16u/blob/8ddb1d8453017eccd153b7bc4bc6e23de3ba4959/src/java.base/share/classes/jdk/internal/misc/Unsafe.java#L1434>

# Implementing Locks using CAS



## 1. Simple TryLock

- Non-blocking tryLock, the lock may be acquired only once
- Regular unlock

## 2. Reentrant TryLock

- Non-blocking tryLock, the lock may be acquired repeatedly by the same thread
- Regular (reentrant) unlock

## 3. Simple Lock

- Blocking lock, the lock may be acquired only once
- Regular unlock

## 4. Reentrant Lock

- Blocking lock, the lock may be acquired repeatedly by the same thread
- Regular (reentrant) unlock

See `TestCasLocks.java`

# 1. SimpleTryLock, non-blocking



```
class SimpleTryLock {  
  
    // Refers to holding thread, null iff unheld  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
  
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        return holder.compareAndSet(null, current);  
    }  
  
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (!holder.compareAndSet(current, null))  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

If the lock is free (holder == null), takes it and return true. Otherwise, holder is unmodified and returns false.

Sets holder to null. If CAS returns false throws an exception indicating that this thread is not holding the lock.

# 1. SimpleTryLock, non-blocking




```
class SimpleTryLock {  
  
    // Refers to holding thread, null iff unheld  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
  
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        return holder.compareAndSet(null, current);  
    }  
  
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (!holder.compareAndSet(current, null))  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

If the lock is free (holder == null), takes it and return true. Otherwise, holder is unmodified and returns false.

How is CAS used here to determine that the lock is not held by the thread executing unlock?

Sets holder to null. If CAS returns false throws an exception indicating that this thread is not holding the lock.

## 2. ReentrantTryLock, non-blocking



```
class ReentrantTryLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    private volatile int holdCount = 0;
```


```
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount++;  
            return true;  
        } else if (holder.compareAndSet(null, current)) {  
            holdCount = 1;  
            return true;  
        }  
        return false;  
    }  
}
```

If the calling thread already holds the lock, we increase the counter

If not, we try to acquire the lock and set the count to 1

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount--;  
            if (holdCount == 0)  
                holder.compareAndSet(current, null);  
            return;  
        } else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

## 2. ReentrantTryLock, non-blocking



```
class ReentrantTryLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    private volatile int holdCount = 0;

    public boolean tryLock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current) {
            holdCount++;
            return true;
        } else if (holder.compareAndSet(null, current)) {
            holdCount = 1;
            return true;
        }
        return false;
    }


    public void unlock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current) {
            holdCount--;
            if (holdCount == 0)
                holder.compareAndSet(current, null);
            return;
        } else
            throw new RuntimeException("Not lock holder");
    }
}
```

Do we have a possible  
data race here?

If the calling thread already holds  
the lock, we increase the counter

If not, we try to acquire the lock  
and set the count to 1

## 2. ReentrantTryLock, non-blocking



```
class ReentrantTryLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    private volatile int holdCount = 0;
```

```
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount++;  
            return true;  
        } else if (holder.compareAndSet(null, current)) {  
            holdCount = 1;  
            return true;  
        }  
        return false;  
    }  
}
```

If the calling thread already holds the lock, we increase the counter

If not, we try to acquire the lock and set the count to 1

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount--;  
            if (holdCount == 0)  
                holder.compareAndSet(current, null);  
            return;  
        } else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

If the calling thread holds the hold we decrement the counter

If the counter is 0 then we release the lock

If the calling thread does not hold the lock we throw an exception

# 3. SimpleLock, blocking




```
class SimpleLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    // The FIFO queue of threads waiting for this lock
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();

    public void lock() {
        final Thread current = Thread.currentThread();
        waiters.add(current);
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {
            LockSupport.park(this);
        }
        waiters.remove();
    }

    public void unlock() {
        final Thread current = Thread.currentThread();
        if (holder.compareAndSet(current, null))
            LockSupport.unpark(waiters.peek());
        else
            throw new RuntimeException("Not lock holder");
    }
}
```



# 3. SimpleLock, blocking



```
class SimpleLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    // The FIFO queue of threads waiting for this lock  
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();
```

```
    public void lock() {  
        final Thread current = Thread.currentThread();  
        waiters.add(current);  
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {  
            LockSupport.park(this);  
        }  
        waiters.remove();  
    }
```

First we add the thread to a waiting queue


We check whether the current thread is at the head of the queue and try to take the lock

Finally, we remove the thread from the waiting list

If not successful, we put the thread to wait (see Javadoc LockSupport)

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.compareAndSet(current, null))  
            LockSupport.unpark(waiters.peek());  
        else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

# 3. SimpleLock, blocking



```
class SimpleLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    // The FIFO queue of threads waiting for this lock  
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();
```

```
    public void lock() {  
        final Thread current = Thread.currentThread();  
        waiters.add(current);  
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {  
            LockSupport.park(this);  
        }  
        waiters.remove();  
    }
```

First we add the thread to a waiting queue

We check whether the current thread is at the head of the queue and try to take the lock

Finally, we remove the thread from the waiting list

If not successful, we put the thread to wait (see Javadoc LockSupport)

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.compareAndSet(current, null))  
            LockSupport.unpark(waiters.peek());  
        else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

We simply try to release the lock

If successful, we wake up one of the waiting threads, if any

As before, we throw an exception if the calling thread is not the thread holding the lock

# 3. SimpleLock, blocking (interruptions)



```
class SimpleLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    // The FIFO queue of threads waiting for this lock
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();

    public void lock() {
        final Thread current = Thread.currentThread();
        boolean wasInterrupted = false;
        waiters.add(current);
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {
            LockSupport.park(this);
            if (Thread.interrupted())
                wasInterrupted = true;
        }
        waiters.remove();
        if (wasInterrupted)
            current.interrupt();
    }
    ...
}
```

We use a flag to record whether the thread has been interrupted during execution

Immediately After taking the lock we check if the thread was interrupted

If the thread was interrupted we call interrupt to propagate the interruption.

## 4. SimpleReentrantLock, blocking

· 19



```
class MyReentrantLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();
    private volatile int holdCount = 0;


    public void lock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current)
            holdCount++;
        else {
            boolean wasInterrupted = false;
            waiters.add(current);
            while (waiters.peek() != current || !holder.compareAndSet(null, current)) {
                LockSupport.park(this);
                if (Thread.interrupted())
                    wasInterrupted = true;
            }
            holdCount = 1;
            waiters.remove();
            if (wasInterrupted)
                current.interrupt();
        }
    }
    ...
}
```

Here we simply combine what we have seen in threads 2 and 3

- Pros
  - A CAS operation is faster than acquiring a lock
  - An unsuccessful CAS operation does not cause thread de-scheduling (blocking)
- Cons
  - CAS operations result in high memory overhead

Will CAS based implementations always scale well?  
(that is, better than lock based)

# Scalability of PRNGs (unrealistic contention)



See `TestPseudoRandom.java`

Throughput of  
computing on  
ThreadLocal objects

Throughput of  
computing on  
shared objects


# Scalability of PRNGs (unrealistic contention)



· 21

Locality has larger impact on scalability than “lock-freedom”

See `TestPseudoRandom.java`



Throughput of computing on ThreadLocal objects

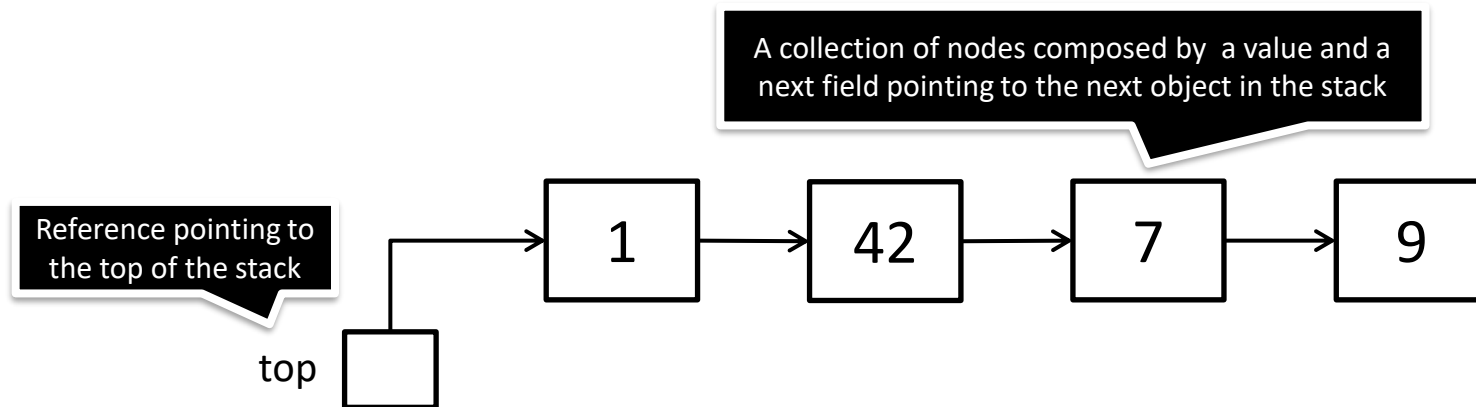
Throughput of computing on shared objects

# Lock-free data structures



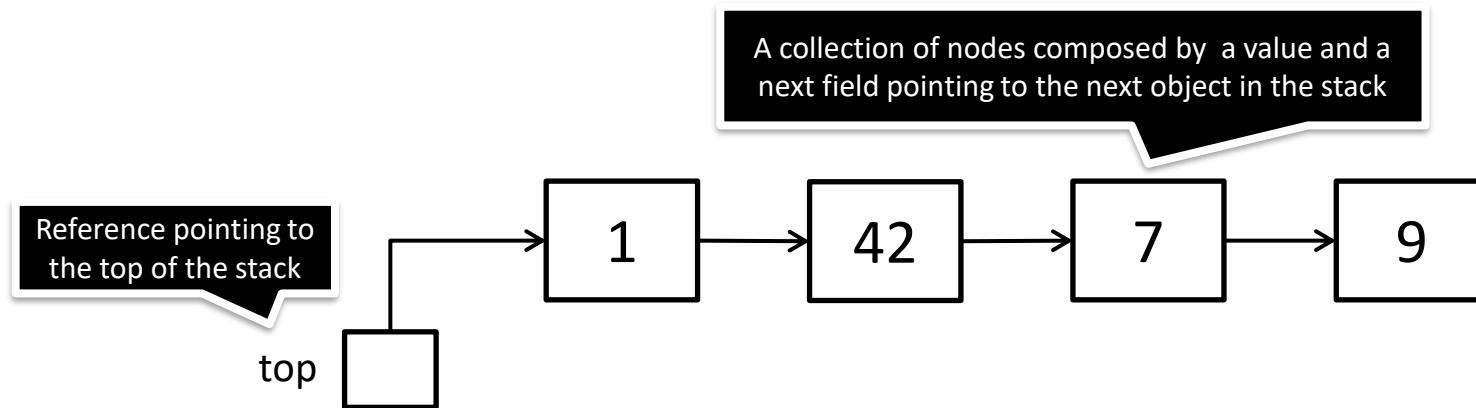


- A stack is a data structure following a LIFO (*last-in-first-out*) policy
  - `push()` – adds an element to the top of the stack
  - `pop()` – removes the top of the stack
- It is typically implemented as a linked list





- A stack is a data structure following a LIFO (*last-in-first-out*) policy
  - push() – adds an element to the top of the stack
  - pop() – removes the top of the stack
- It is typically implemented as a linked list



# Lock-free Stack


Introduced by  
Treiber in 1986

· 24



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```

See `TestLockFreeStack.java`




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

newHead




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```






# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts




# Lock-free Heap | push()

· 25



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts



# Lock-free Heap | pop()

· 26




```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead,newHead));  
        return oldHead.value;  
    }  
}
```

newHead



oldHead




# Lock-free Heap | pop()

· 26



```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```




# Lock-free Heap | pop()

· 26



```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead,newHead));  
        return oldHead.value;  
    }  
}
```






# Lock-free Heap | pop()

· 26




```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```



# Lock-free Heap | pop()




```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```



Would there be any problem if another thread is executing pop() concurrently?

- A queue is a data structure following a FIFO (*first-in-first-out*) policy
  - enqueue() – adds an element to the tail of the queue
  - dequeue() – removes an element from the head of the queue
- It is typically implemented as a linked list



# Lock-free queue

- Michael-Scott lock free queue, introduced in 1996 (see optional readings)
- Implemented in ConcurrentLinkedQueue in java.concurrent.\* by Doug Lea et. al. (see [here](#))
  - The version on the right is not the JDK implementation


```
class MSQueue<T> implements UnboundedQueue<T> {
    private final AtomicReference<Node<T>> head, tail;

    public MSQueue() {
        Node<T> dummy = new Node<T>(null, null);
        head = new AtomicReference<Node<T>>(dummy);
        tail = new AtomicReference<Node<T>>(dummy);
    }

    public void enqueue(T item) {
        Node<T> node = new Node<T>(item, null);
        while (true) {
            Node<T> last = tail.get(), next = last.next.get();
            if (last == tail.get()) {
                if (next == null) {
                    // In quiescent state, try inserting new node
                    if (last.next.compareAndSet(next, node)) {
                        // Insertion succeeded, try advancing tail
                        tail.compareAndSet(last, node);
                        return;
                    }
                } else {
                    // Queue in intermediate state, advance tail
                    tail.compareAndSet(last, next);
                }
            }
        }
    }

    public T dequeue() {
        while (true) {
            Node<T> first = head.get(), last = tail.get(), next = first.next.get();
            if (first == head.get()) {
                if (first == last) {
                    if (next == null)
                        return null;
                    else
                        tail.compareAndSet(last, next);
                } else {
                    T result = next.item;
                    if (head.compareAndSet(first, next))
                        return result;
                }
            }
        }
    }
}
```


See `TestMSQueue.java`



# Lock-free queue | initialization




```
class MSQueue<T> implements UnboundedQueue<T> {  
    private final AtomicReference<Node<T>> head, tail;  
  
    public MSQueue() {  
        Node<T> dummy = new Node<T>(null, null);  
        head = new AtomicReference<Node<T>>(dummy);  
        tail = new AtomicReference<Node<T>>(dummy);  
    }  
    ...  
}
```



# Lock-free queue | enqueue



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
...  
}
```




# Lock-free queue | enqueue

· 30



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
...  
}
```




# Lock-free queue | enqueue

· 30



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```






# Lock-free queue | enqueue

· 30



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```




# Lock-free queue | enqueue

· 30



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```

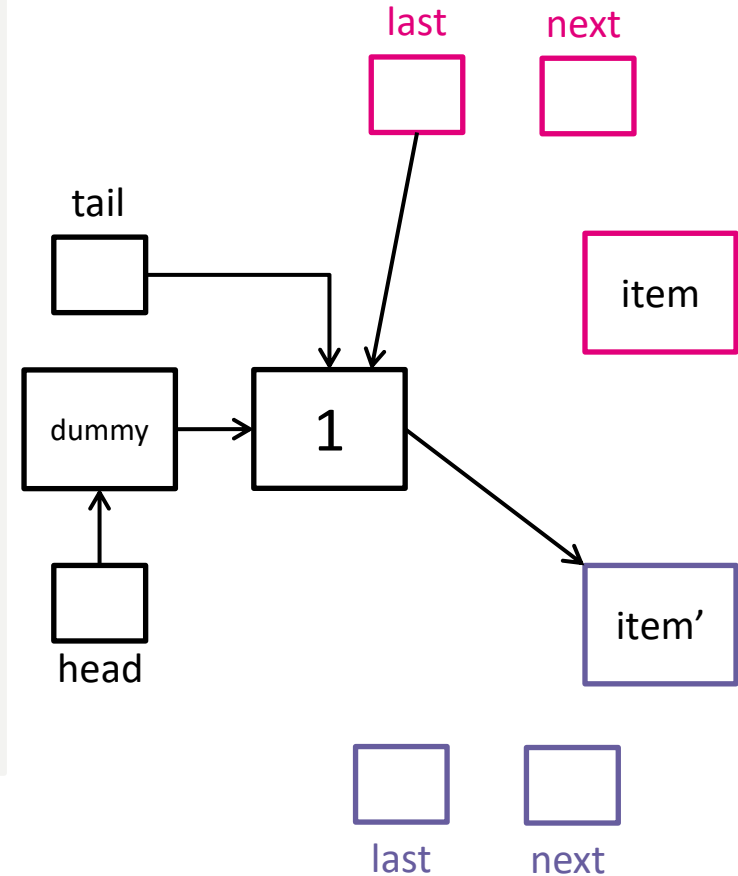


# Lock-free queue | enqueue

· 31



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
...  
}
```



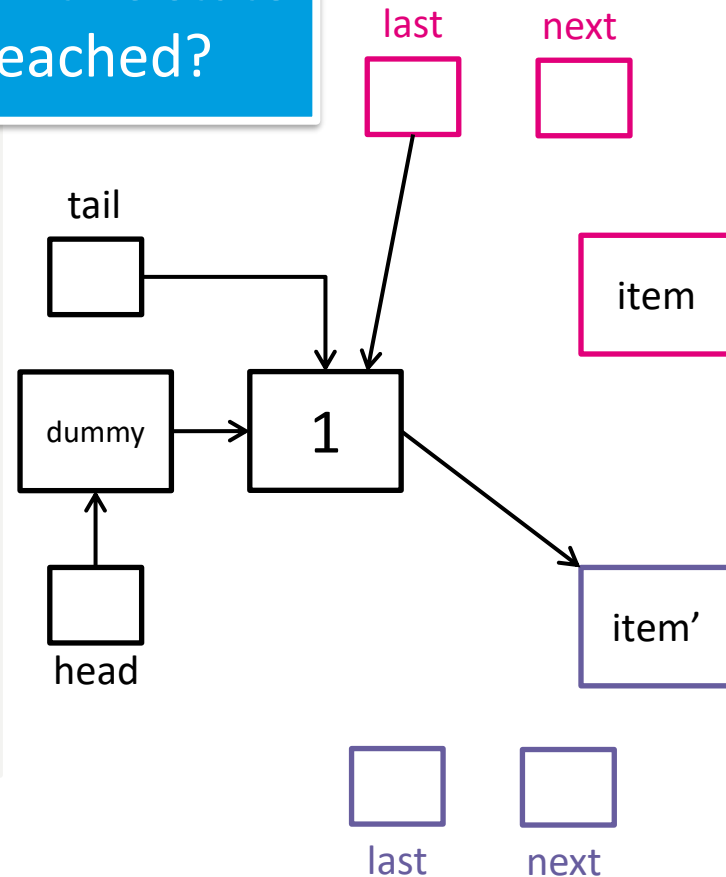
# Lock-free queue | enqueue

· 31



How can this state be reached?

```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
    ...  
}
```




# Lock-free queue | enqueue

· 31



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
public void enqueue(T item) {  
    Node<T> node = new Node<T>(item, null);  
    while (true) {  
        Node<T> last = tail.get();  
        Node<T> next = last.next.get();  
        if (last == tail.get()) {  
            if (next == null) {  
                // In quiescent state, try inserting new node  
                if (last.next.compareAndSet(next, node)) {  
                    // Insertion succeeded, try advancing tail  
                    tail.compareAndSet(last, node);  
                    return;  
                }  
            } else {  
                // Queue in intermediate state, advance tail  
                tail.compareAndSet(last, next);  
            }  
        }  
    }  
}  
...  
}
```

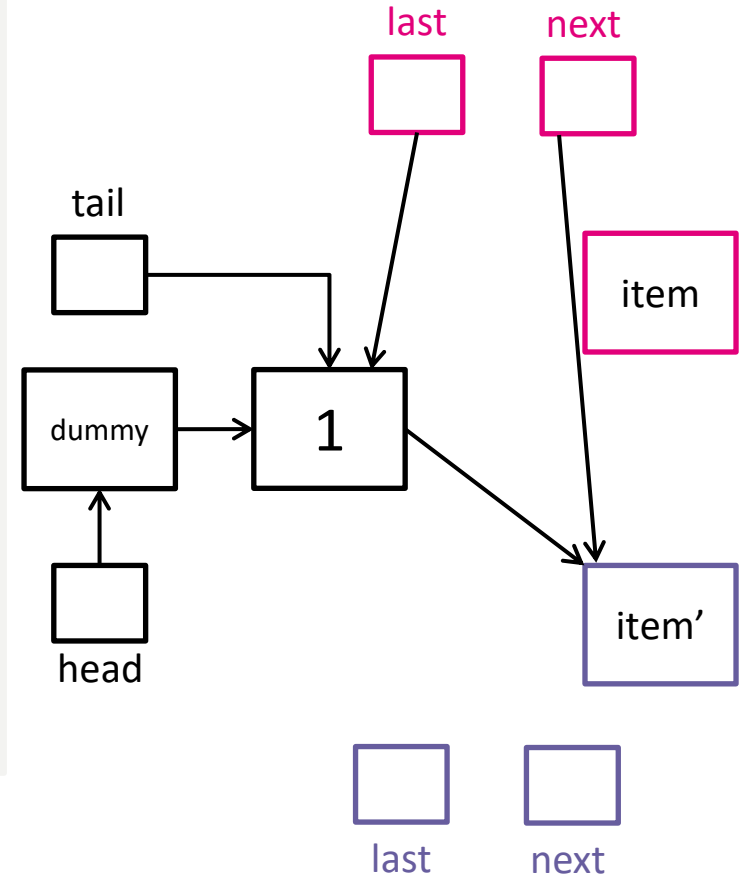


# Lock-free queue | enqueue

· 31



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```

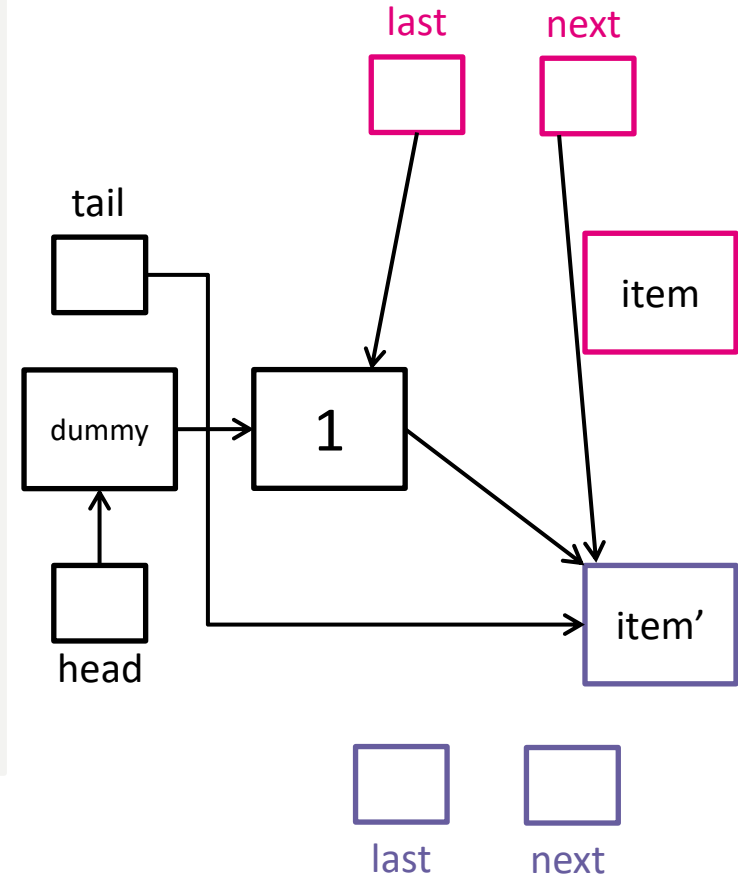


# Lock-free queue | enqueue

· 31



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```




# Lock-free queue | enqueue

· 31



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```





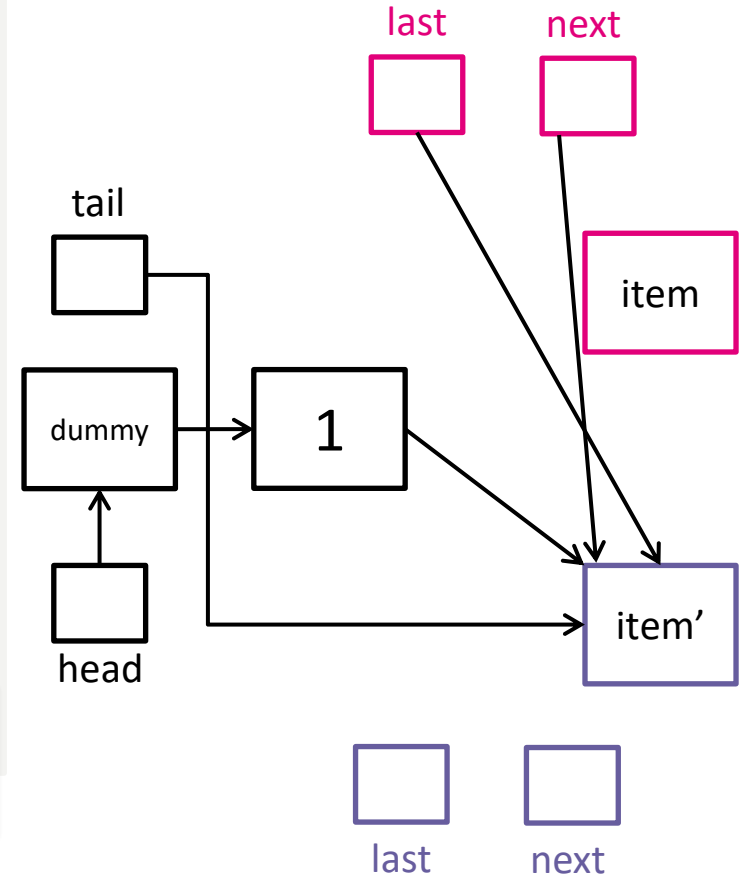
# Lock-free queue | enqueue

· 31



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
public void enqueue(T item) {  
    Node<T> node = new Node<T>(item, null);  
    while (true) {  
        Node<T> last = tail.get();  
        Node<T> next = last.next.get();  
        if (last == tail.get()) {  
            if (next == null) {  
                // In quiescent state, try inserting new node  
                if (last.next.compareAndSet(next, node)) {  
                    // Insertion succeeded, try advancing tail  
                    tail.compareAndSet(last, node);  
                    return;  
                }  
            } else {  
                // Queue in intermediate state, advance tail  
                tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```

In case another thread is enqueueing, and didn't update the tail, the current thread helps by advancing the tail




# Lock-free queue | enqueue

· 32



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```

If the tail has been changed, then the thread restarts right away




# Lock-free queue | dequeue

· 33



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```




# Lock-free queue | dequeue

· 33



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```

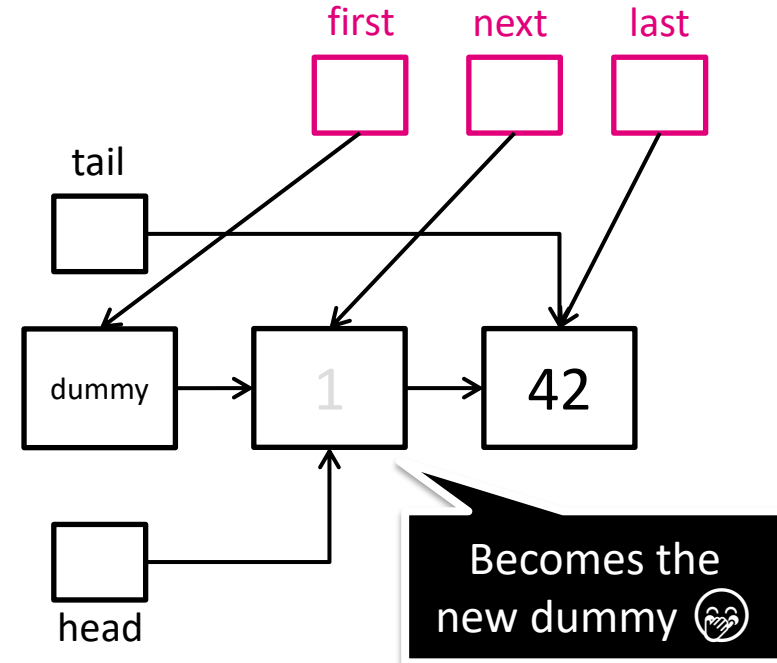


# Lock-free queue | dequeue

· 33



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```




# Lock-free queue | dequeue

· 34



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```




# Lock-free queue | dequeue

· 34



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```




# Lock-free queue | dequeue

· 34



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```

If the next field of the head is not null (because another thread push in the meantime), then the calling thread helps advancing the tail and tries to pop again.






# Lock-free queue | dequeue

· 35



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```




# Lock-free queue | dequeue

· 35



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```

If the next field of the head is pointing to null, then we return null. Why is this correct?




# Lock-free queue | dequeue

· 36



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```


How can a state where  $\text{first} \neq \text{head.get()}$  be reached?



# ABA Problem




1. Thread 1 starts popping A



# ABA Problem




1. Thread 1 starts popping A
2. Before thread 1 finishes, thread 2 pops A and B




# ABA Problem



1. Thread 1 starts popping A
2. Before thread 1 finishes, thread 2 pops A and B
3. Thread 2 pushes A back *// recovered from memory (same as thread 1 was popping)*



1. Thread 1 starts popping A
2. Before thread 1 finishes, thread 2 pops A and B
3. Thread 2 pushes A back *// recovered from memory (same as thread 1 was popping)*
4. Thread 1 finishes popping A
  1. Incorrectly, as it thinks that A is the same as before thread 2 operations





- It is a memory allocation issue which affects mainly languages without garbage collection (e.g., C and C++)
- Languages such as Java do not have this problem because garbage collection ensures that newly created objects are fresh
  - Step 3 in the previous slides would have created a new A object



- Compare-And-Swap (CAS)
  - Lock-free Counter
  - Atomic libraries
  - CAS based lock implementation
- Lock-free stack
- Lock-free queue
- ABA problem

# Scalability in Mastodon servers



· 43

