



# Practical Concurrent and Parallel Programming XII

## Message Passing II

Raúl Pardo

- Actors model (revisited)
  - Bounded Buffer
  - Primer
- Dynamic topology
- Fault-tolerance
  - Supervision
- Adaptive load balancing
  - Scatter-Gather
- Changing behaviour

# What is an Actor? (Bird's eye, revisited)




- An actor can be seen as a sequential unit of computation
  - Although, formally, the model allows for parallelism within the actor, one can safely assume that there are not concurrency issues within the actor.
  - You can think of an actor as a thread
- Actors can send messages to other actors



# Actor – Specification (revisited)



- An actor is an abstraction of a thread (intuitively)
- An actors can only execute any of these 4 actions
  1. Receive messages from other actors
  2. Send asynchronous messages to other actors
  3. Create new actors
  4. Change its behaviour (local state and/or message handlers)
- Actors do not share memory
  - They only have access to:
    - Their *local state* (local memory)
    - Their *mailbox* (multiset of fixed size with received messages)
    - By default, the mailbox is of unbounded size



# Producer-consumer problem | Intuition



- Perhaps more intuitive example

Consumers

Producers



Shared data structure of fixed size

# Producer-consumer problem | Intuition



Consumers

Producers



Shared data structure of fixed size

# Bounded Buffer with Actors



## Producer -> Buffer

- Put(ActorRef<Producer> ref, int elem)

## Consumer -> Buffer

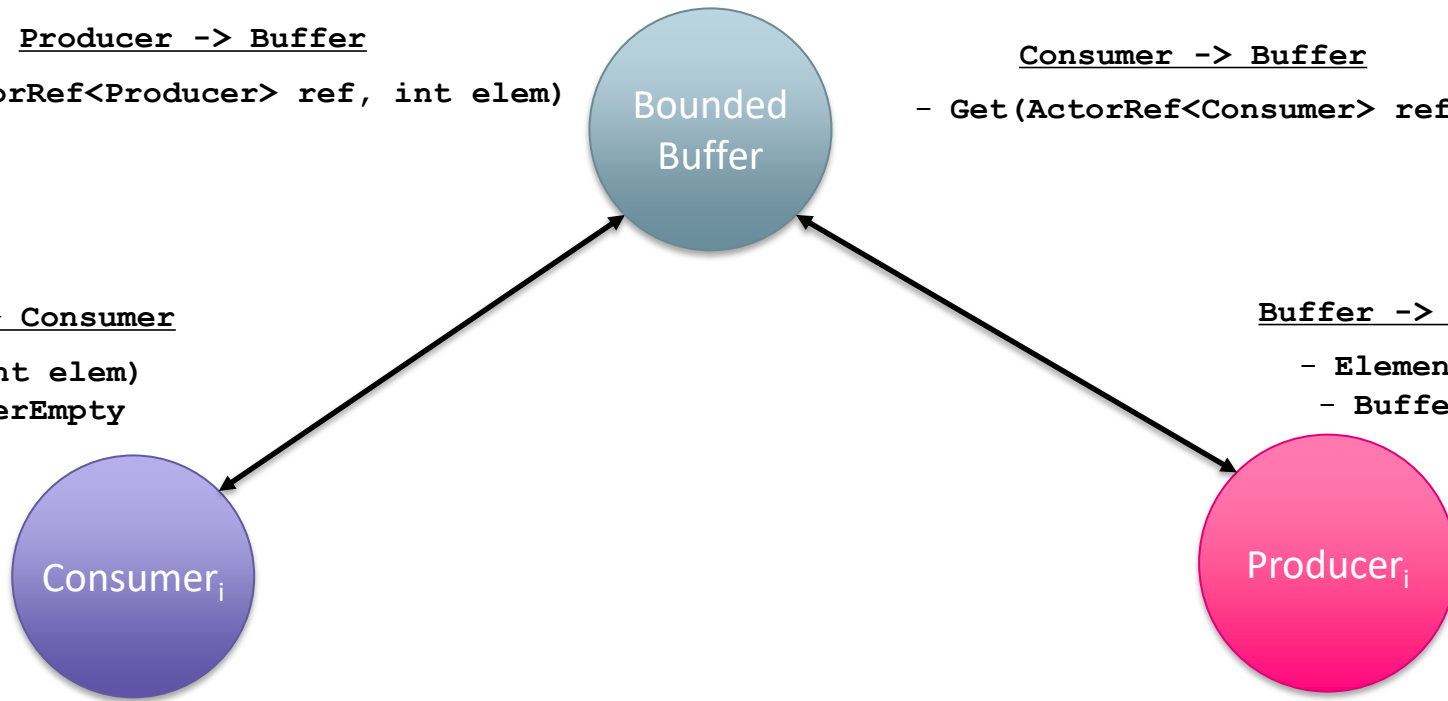
- Get(ActorRef<Consumer> ref)

## Buffer -> Consumer

- Get(int elem)  
- BufferEmpty

## Buffer -> Producer

- ElementAdded  
- BufferFull



# Bounded Buffer with Actors



## Producer -> Buffer

- `Put(ActorRef<Producer> ref, int elem)`

## Consumer -> Buffer

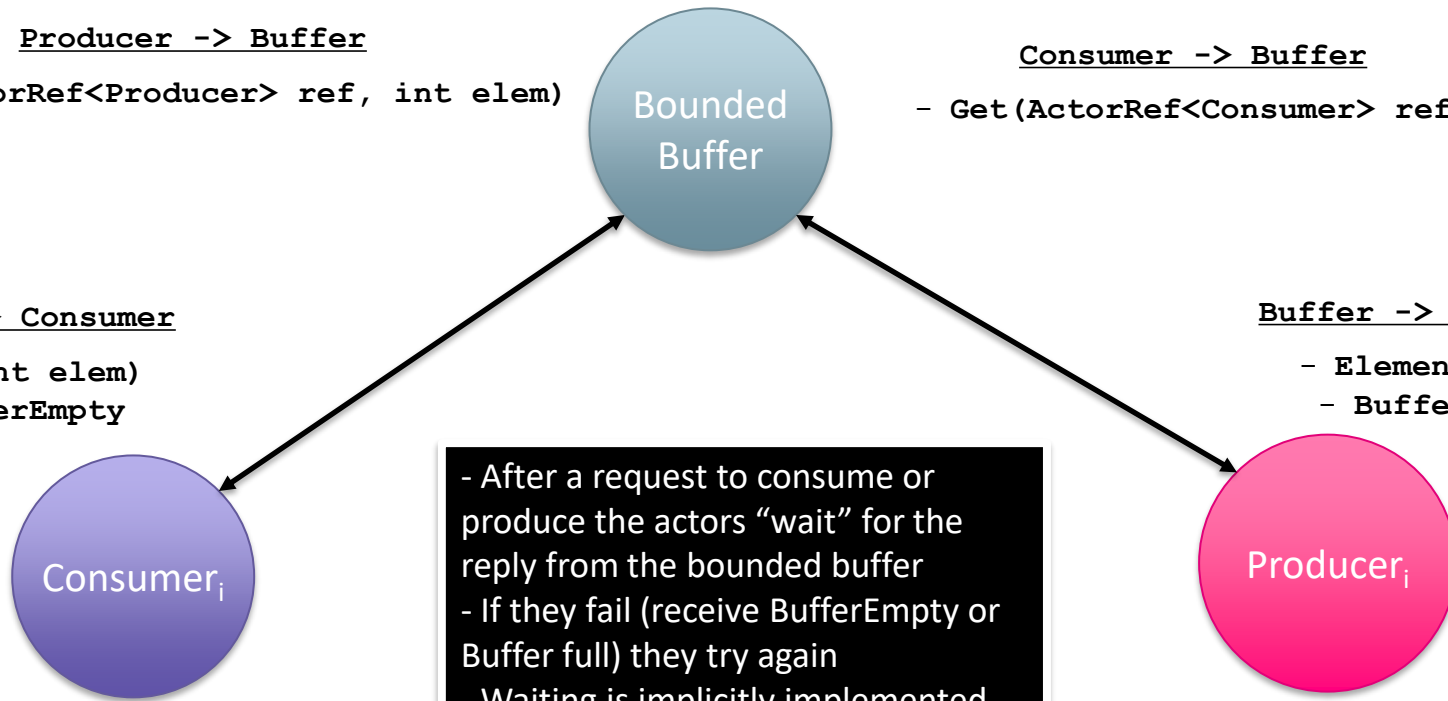
- `Get(ActorRef<Consumer> ref)`

## Buffer -> Consumer

- `Get(int elem)`
- `BufferEmpty`

## Buffer -> Producer

- `ElementAdded`
- `BufferFull`



- After a request to consume or produce the actors “wait” for the reply from the bounded buffer  
- If they fail (receive `BufferEmpty` or `Buffer full`) they try again  
- Waiting is implicitly implemented as producers and consumers wait for the answer of the bounded buffer, and the buffer replies with error only if the actors cannot make progress



# Bounded Buffer (size 1) – execution example



· 8

Producer<sub>1</sub>

Producer<sub>2</sub>

Buffer (1)

Consumer<sub>1</sub>

Consumer<sub>2</sub>

# Bounded Buffer (size 1) – execution example



.8

Producer<sub>1</sub>

Producer<sub>2</sub>

Buffer (1)

Consumer<sub>1</sub>

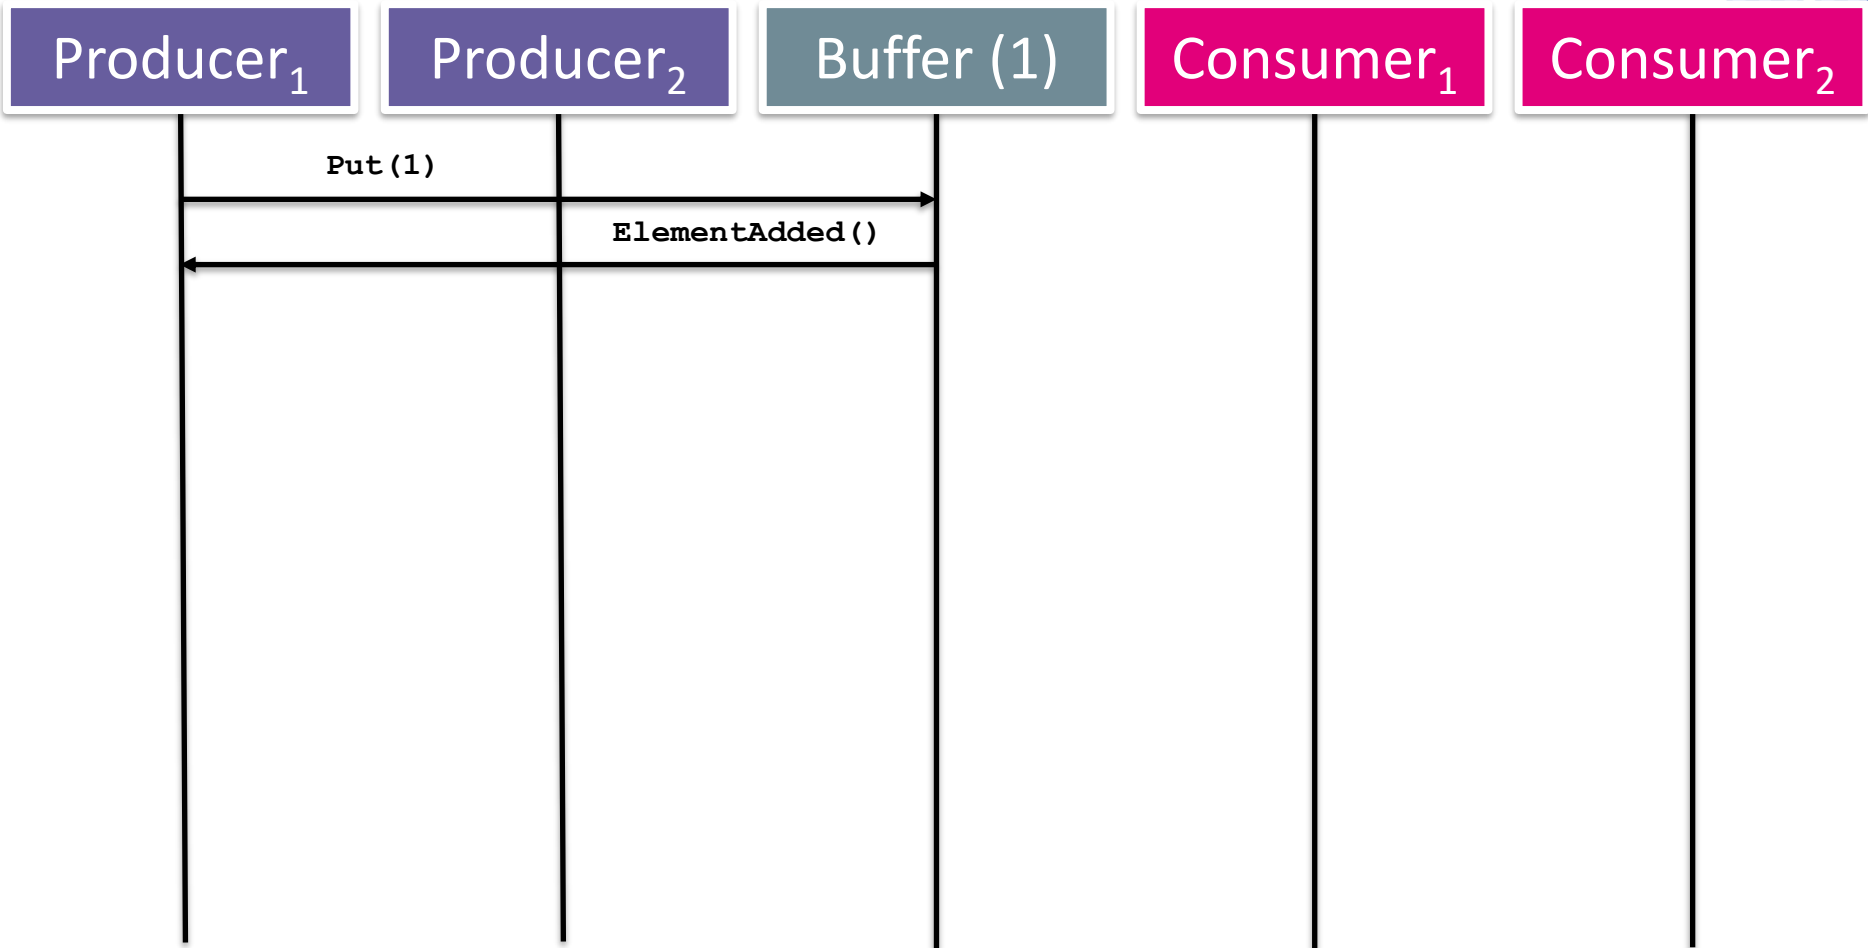
Consumer<sub>2</sub>

Put (1)

# Bounded Buffer (size 1) – execution example



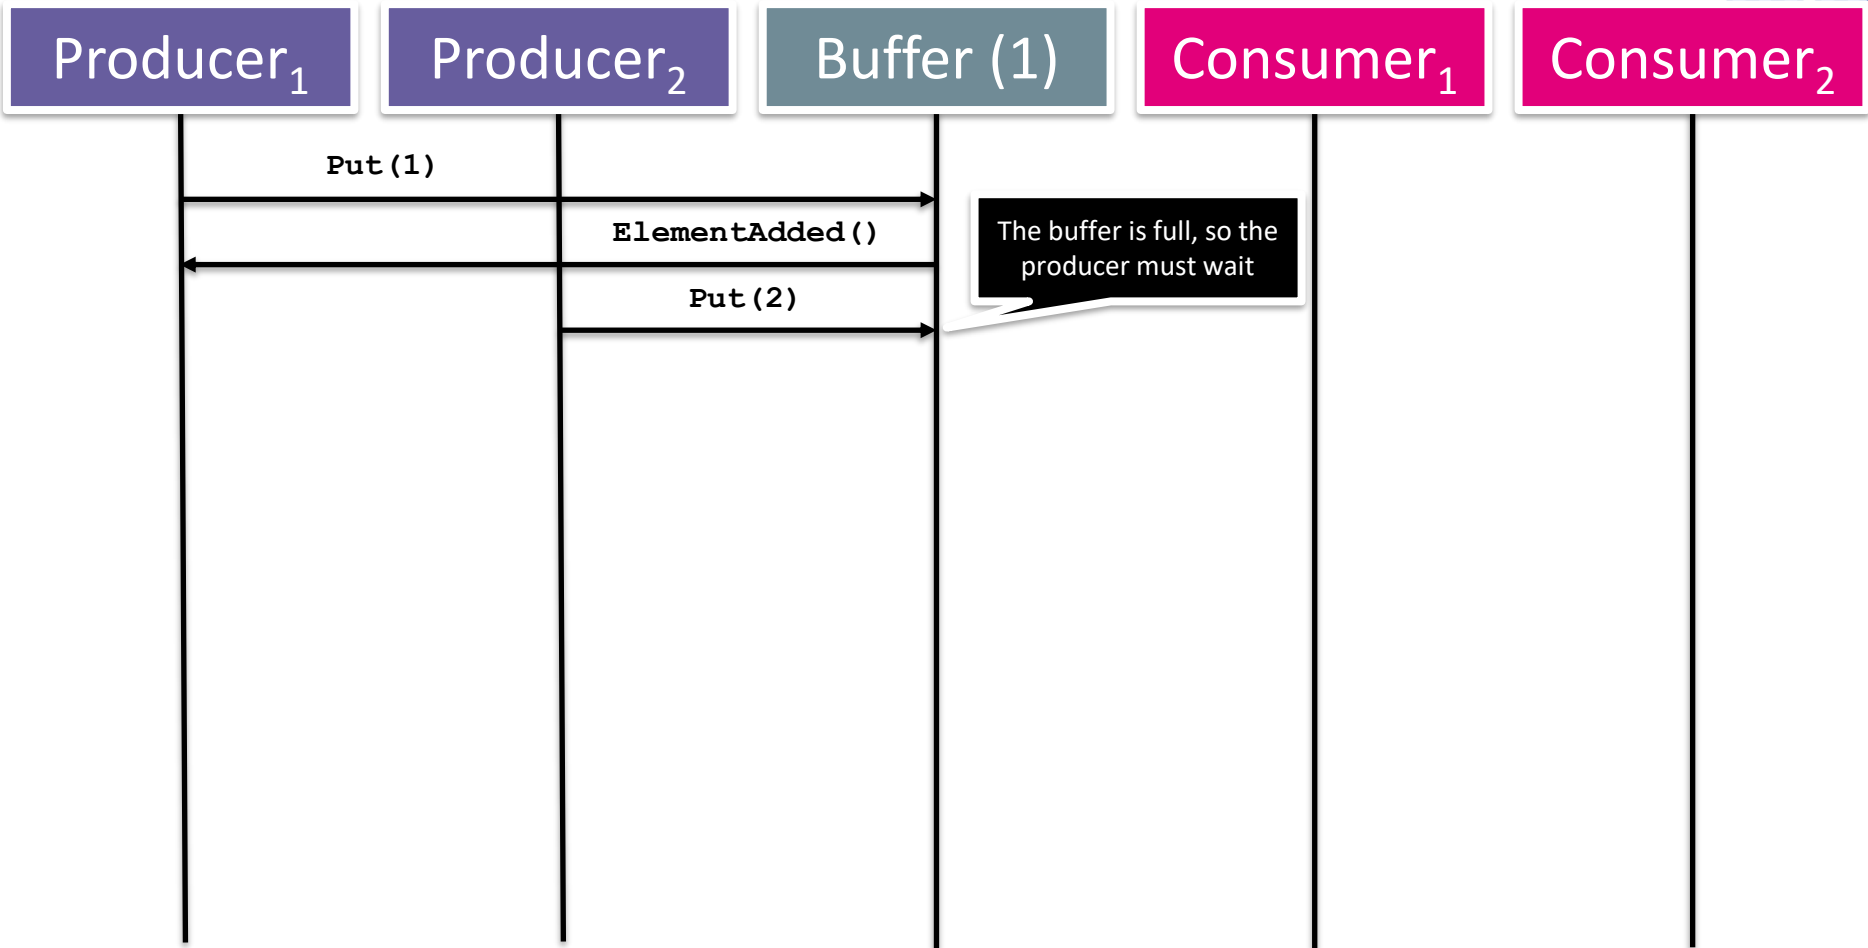
· 8



# Bounded Buffer (size 1) – execution example



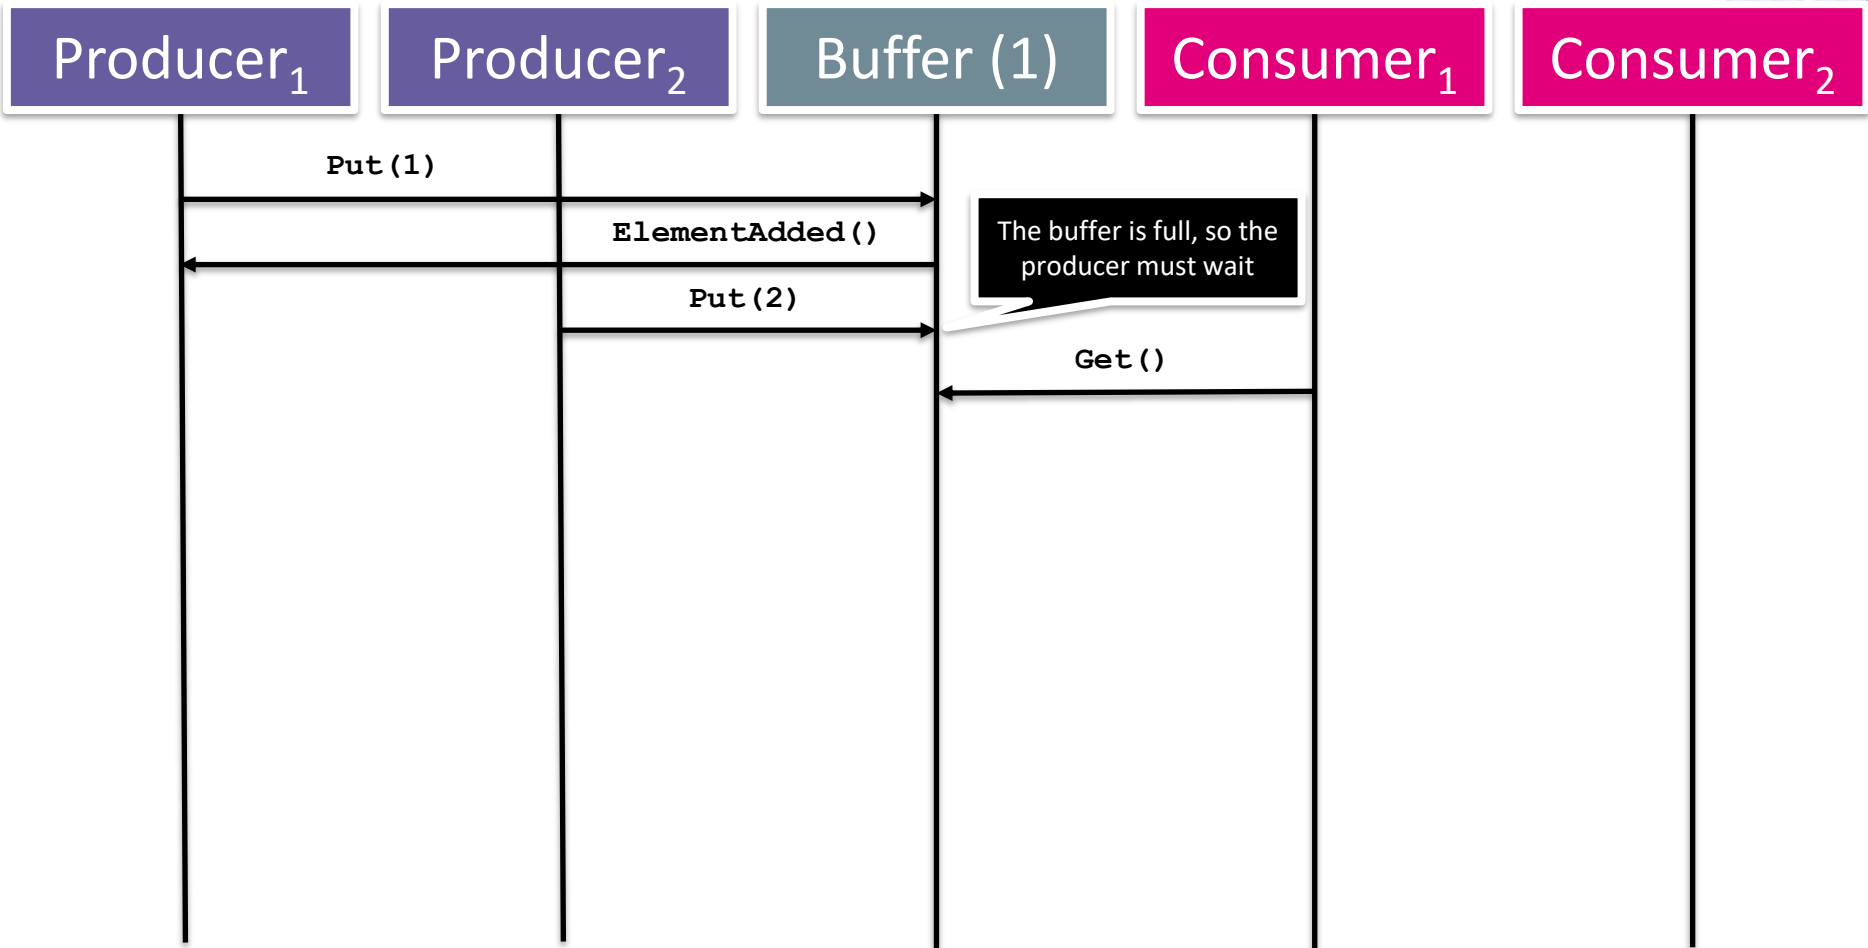
.8



# Bounded Buffer (size 1) – execution example



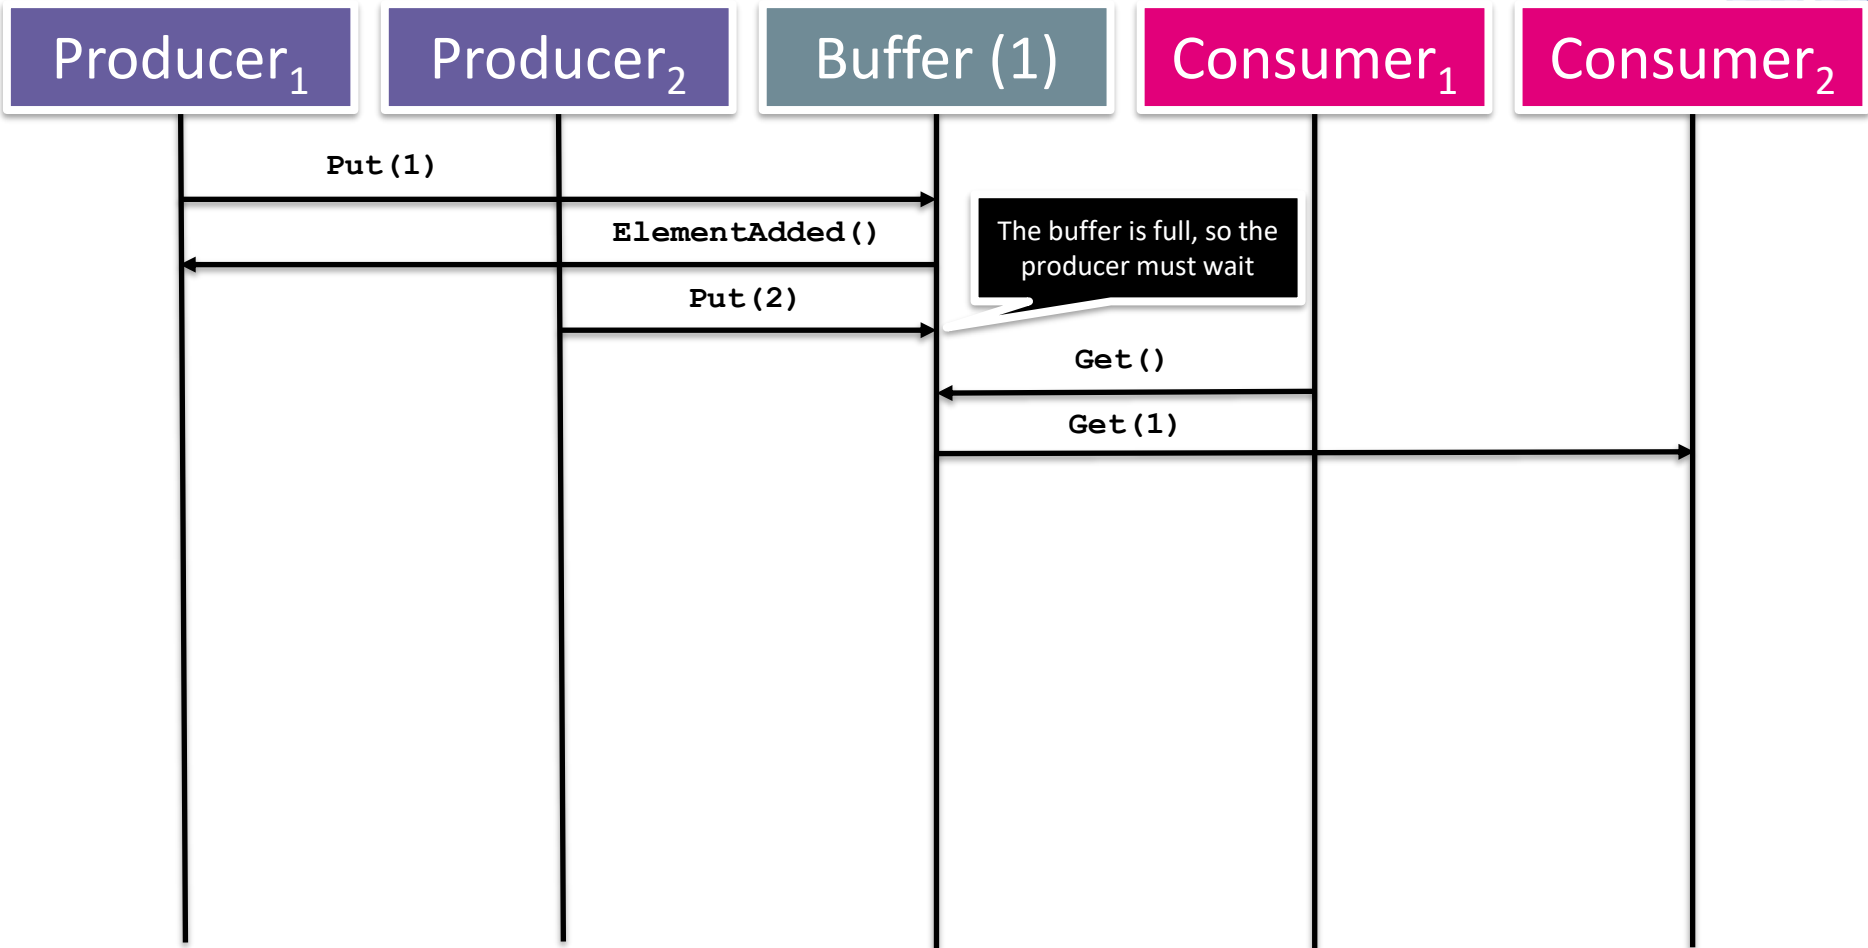
.8



# Bounded Buffer (size 1) – execution example



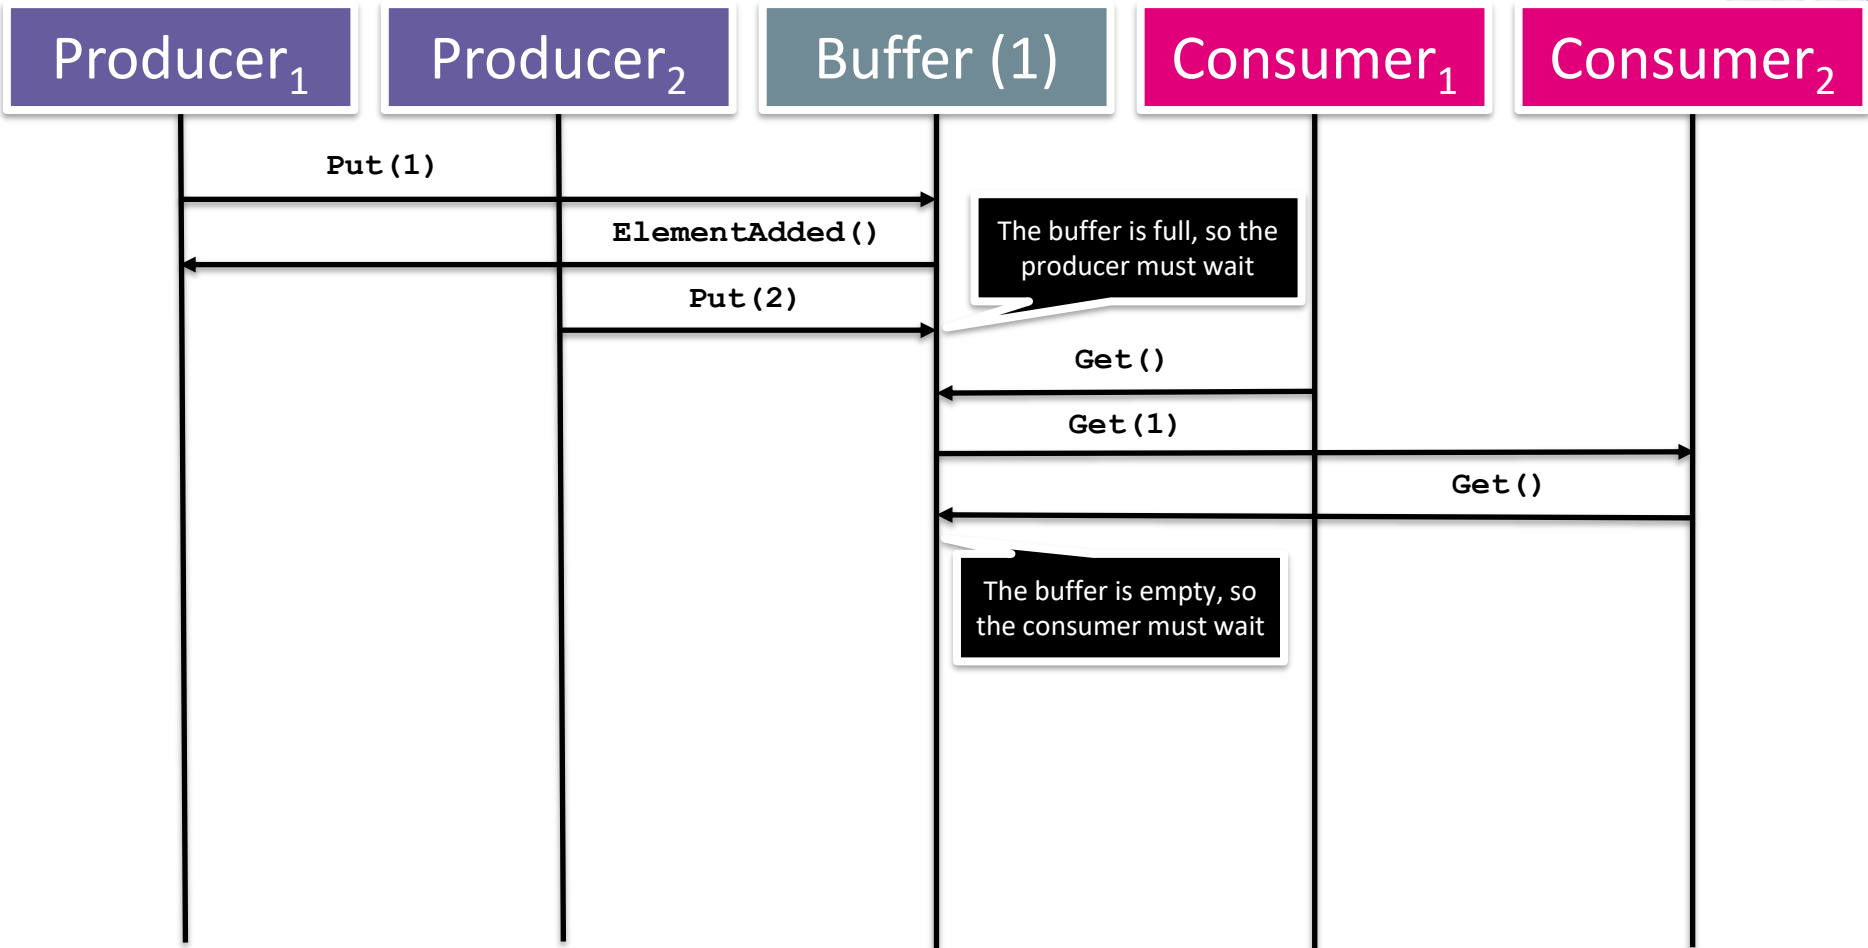
8



# Bounded Buffer (size 1) – execution example



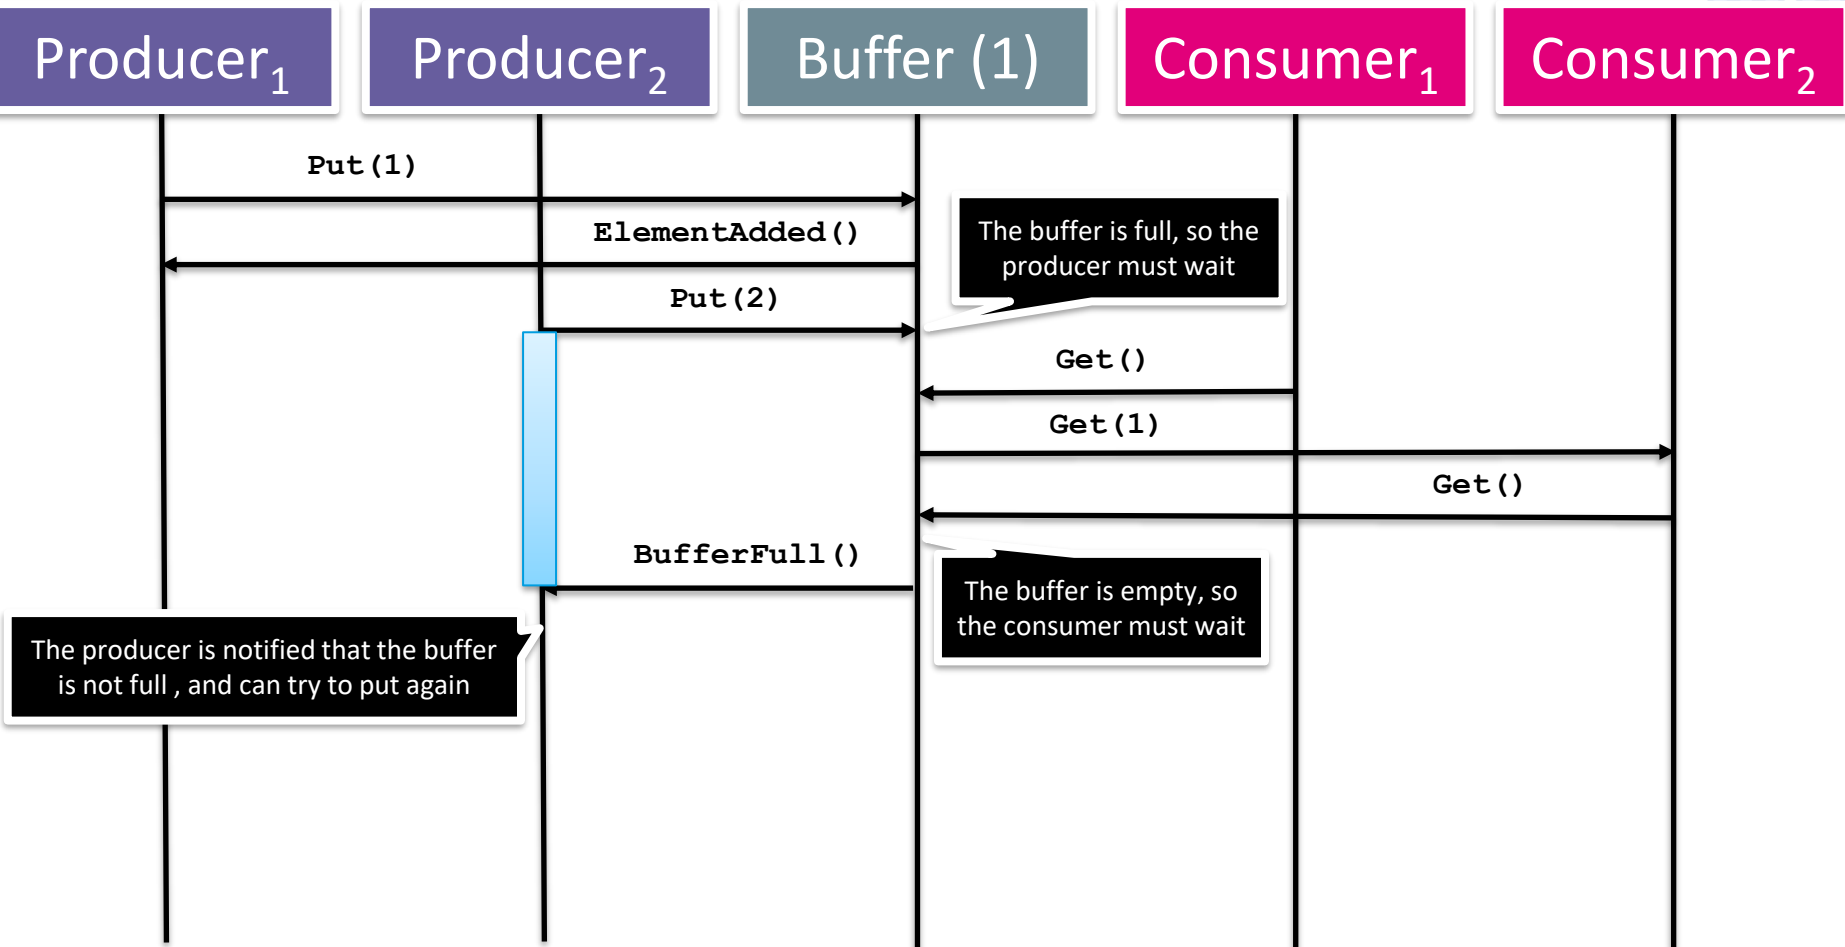
.8



# Bounded Buffer (size 1) – execution example



8

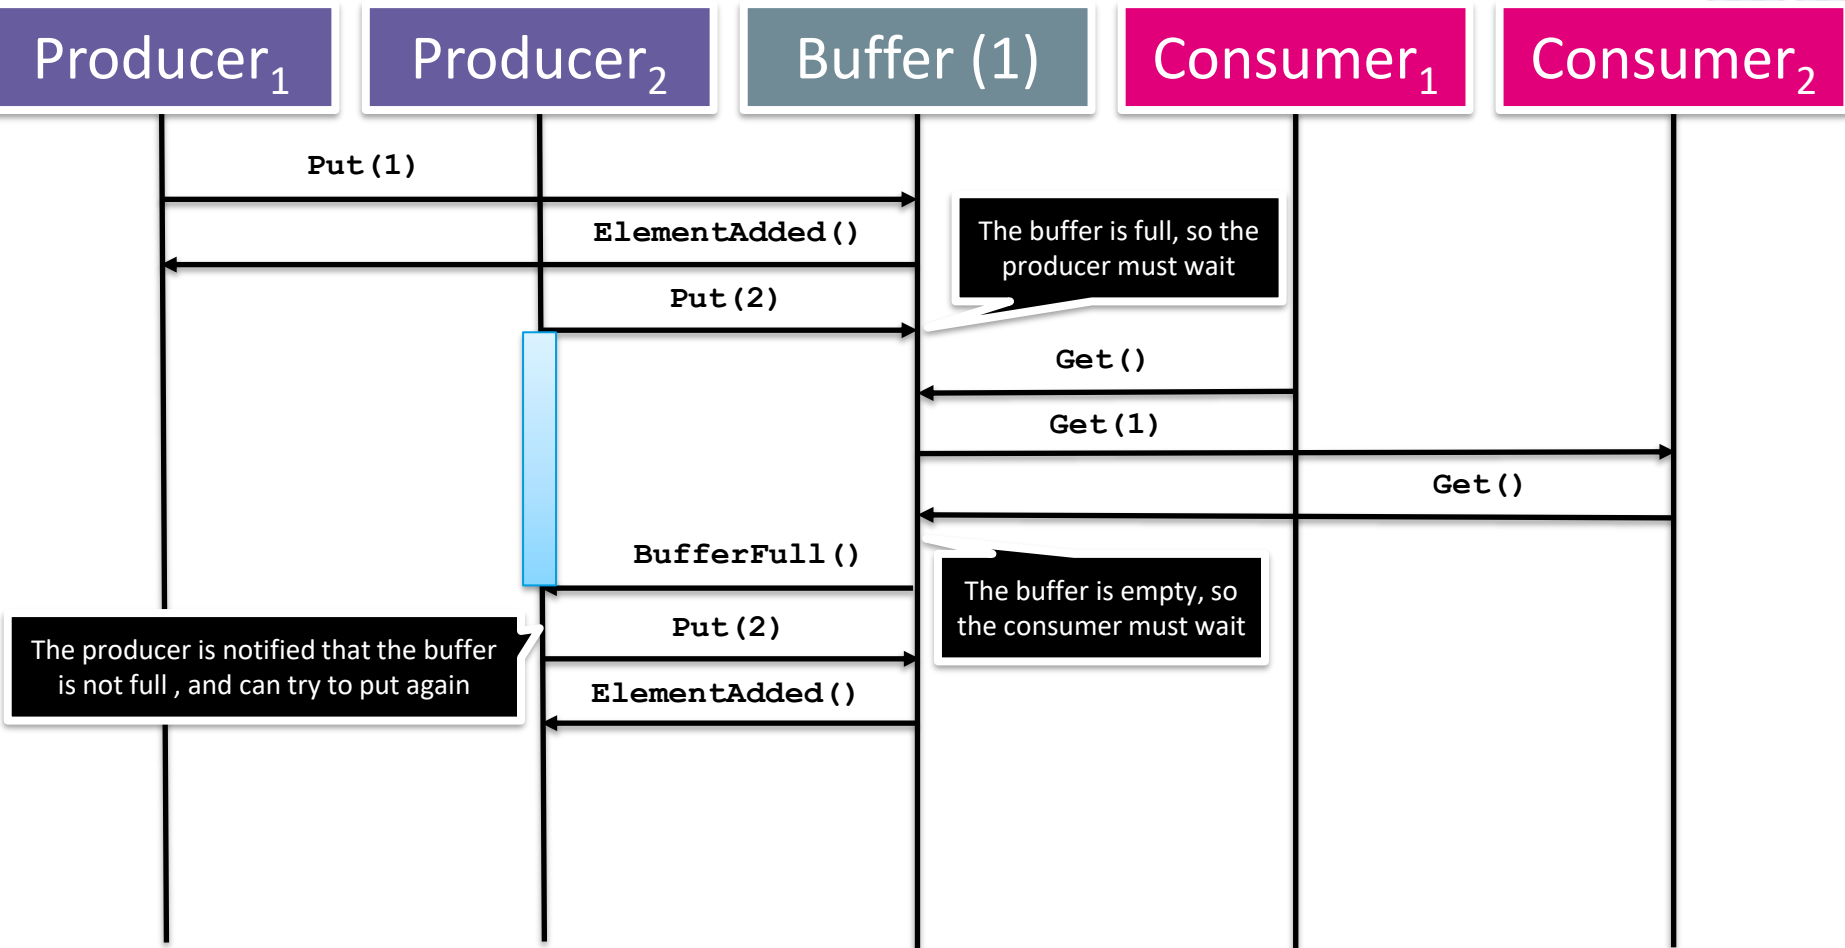




# Bounded Buffer (size 1) – execution example



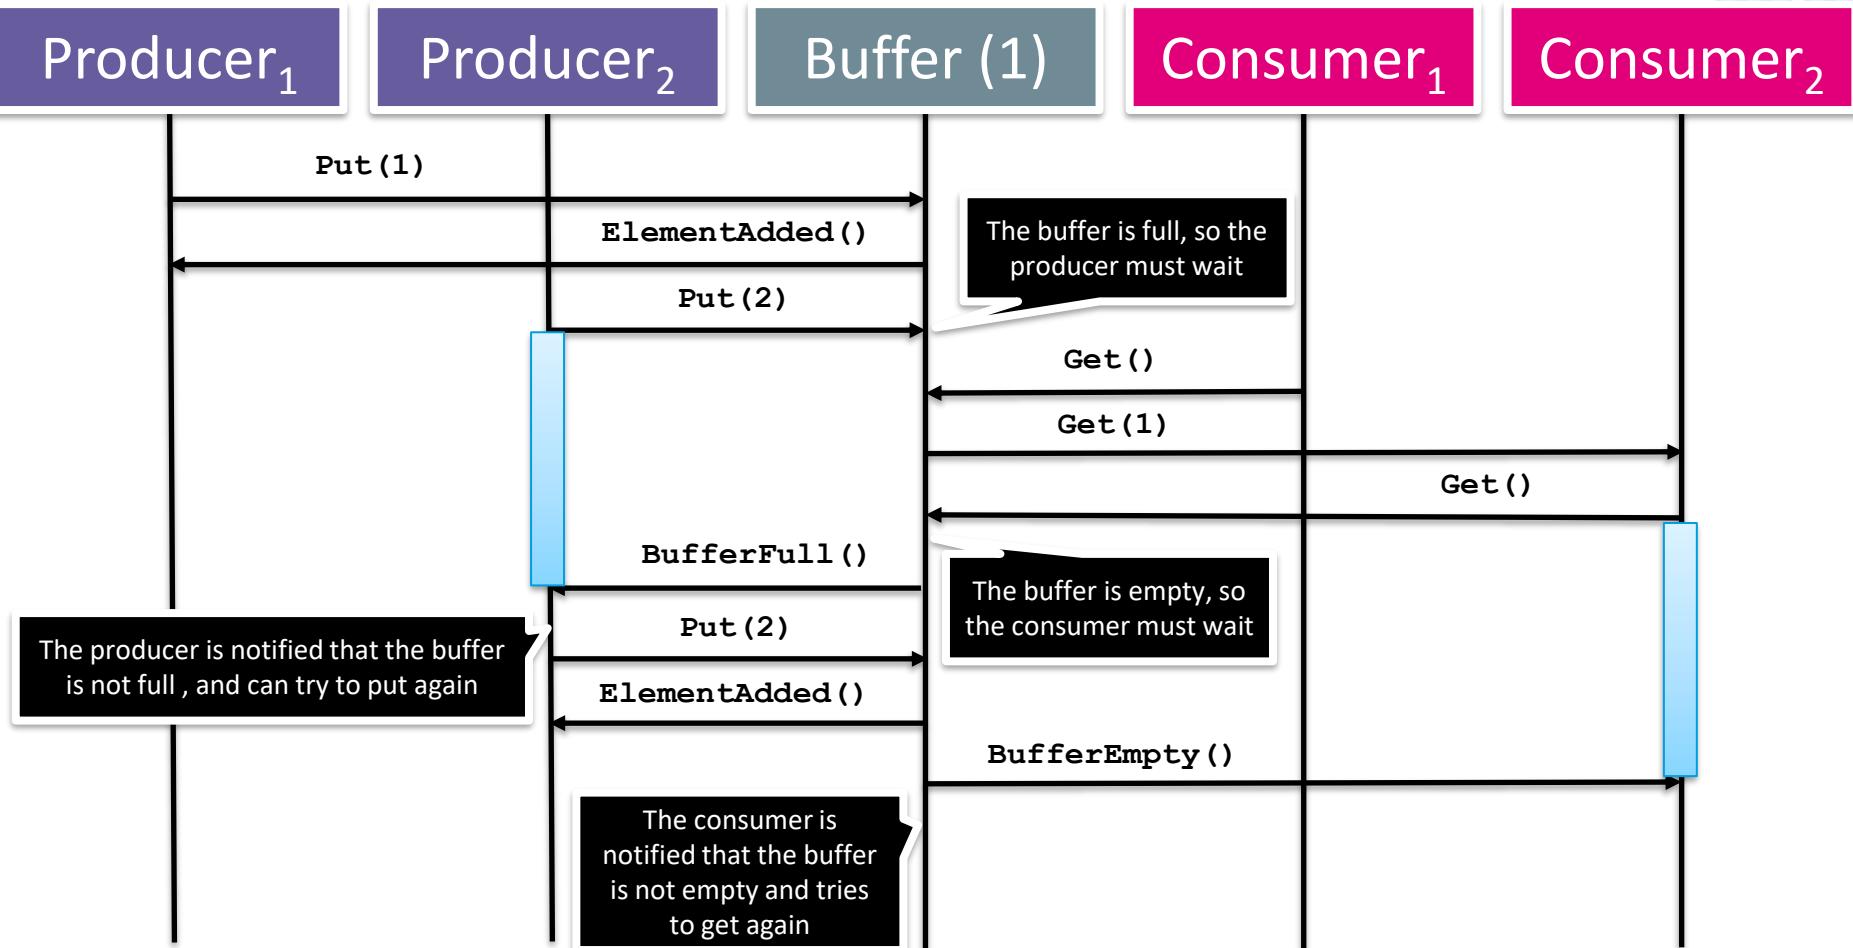
· 8



# Bounded Buffer (size 1) – execution example



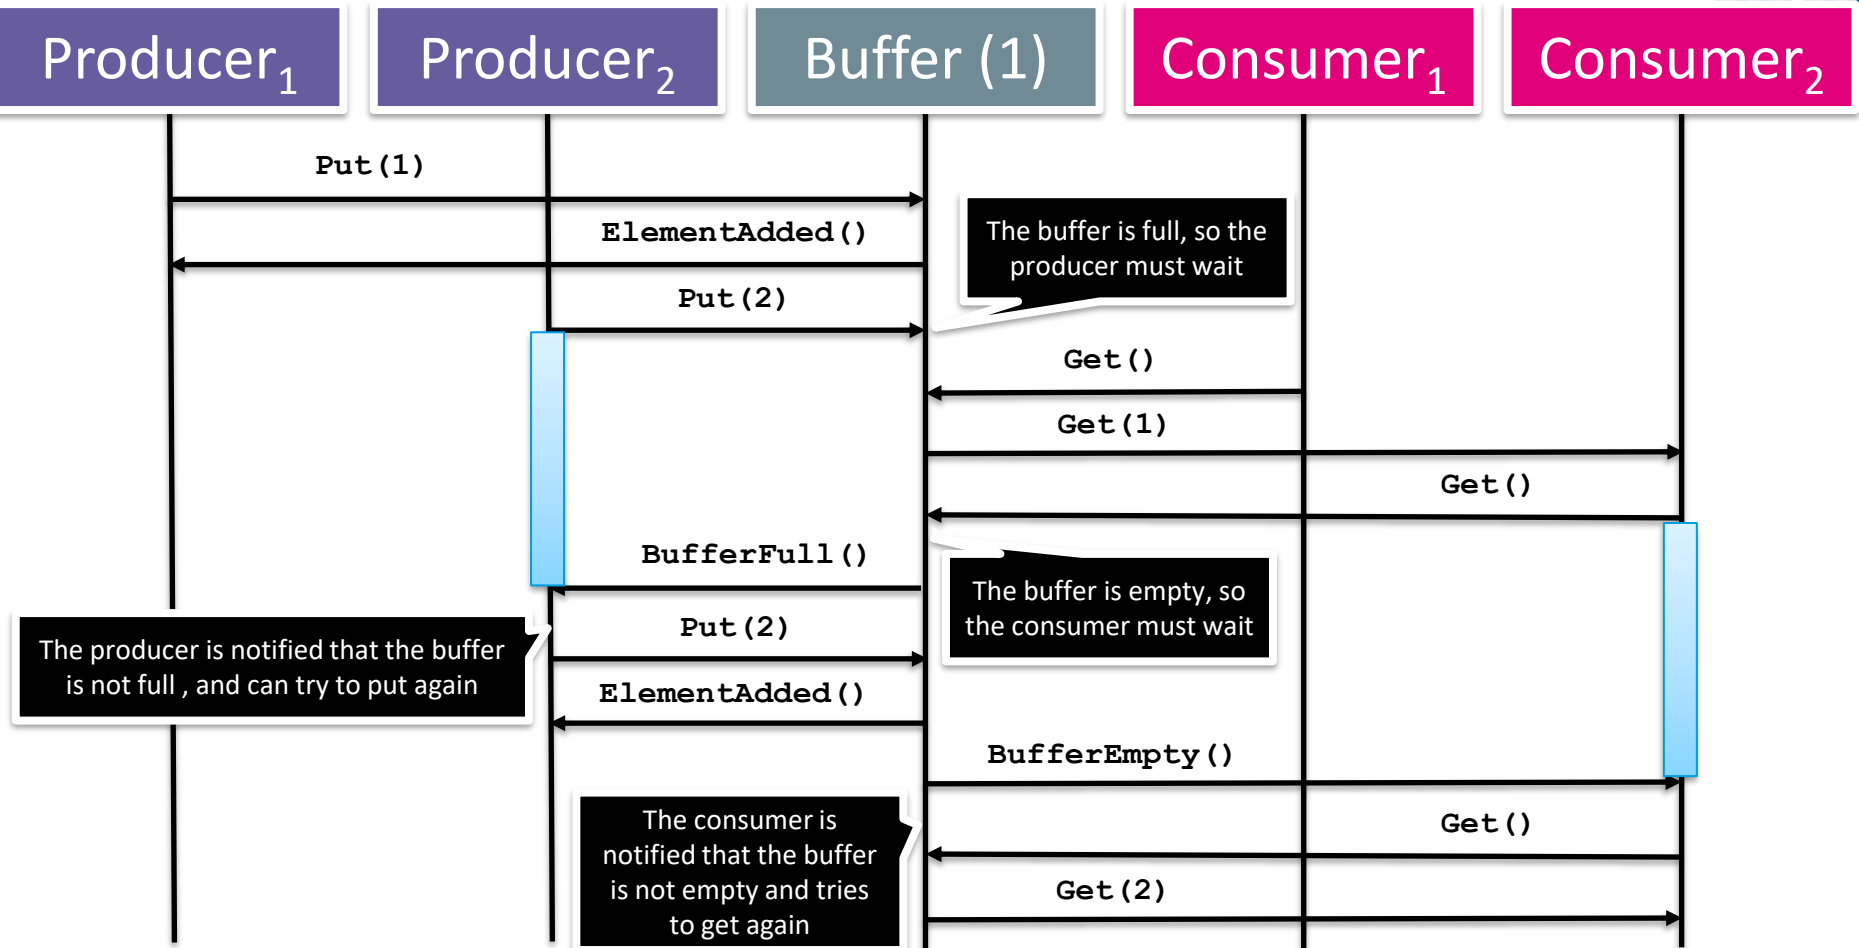
8



# Bounded Buffer (size 1) – execution example

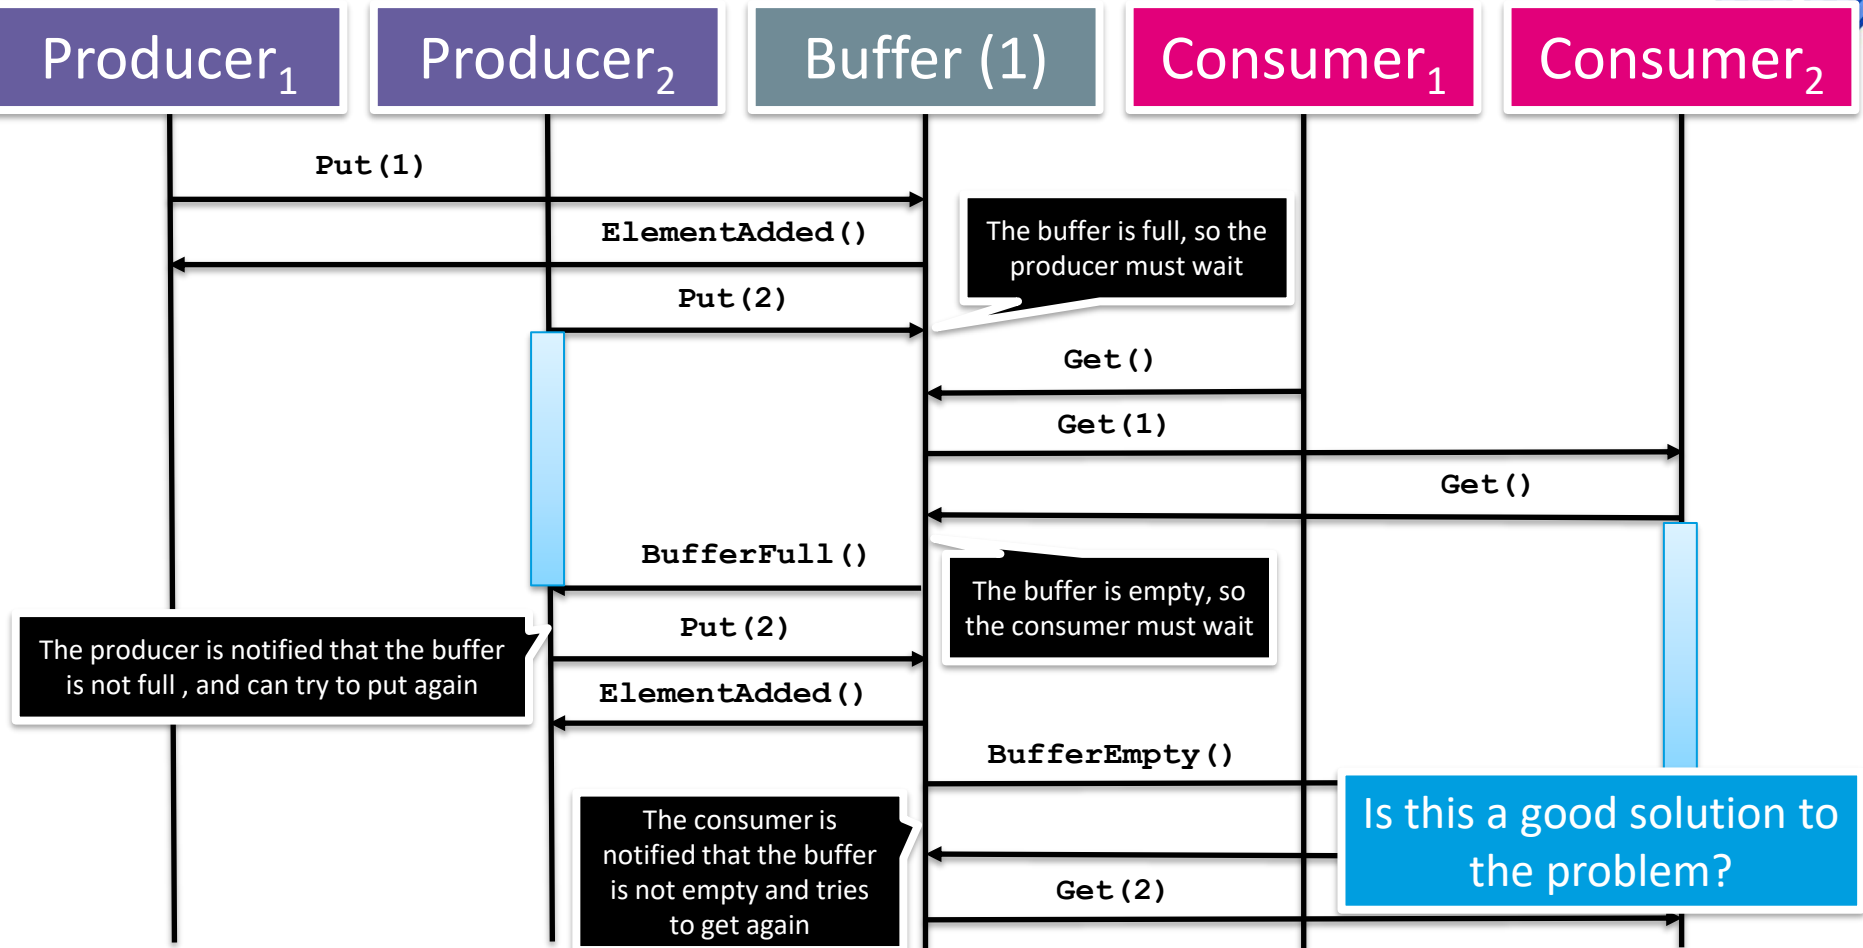


· 8



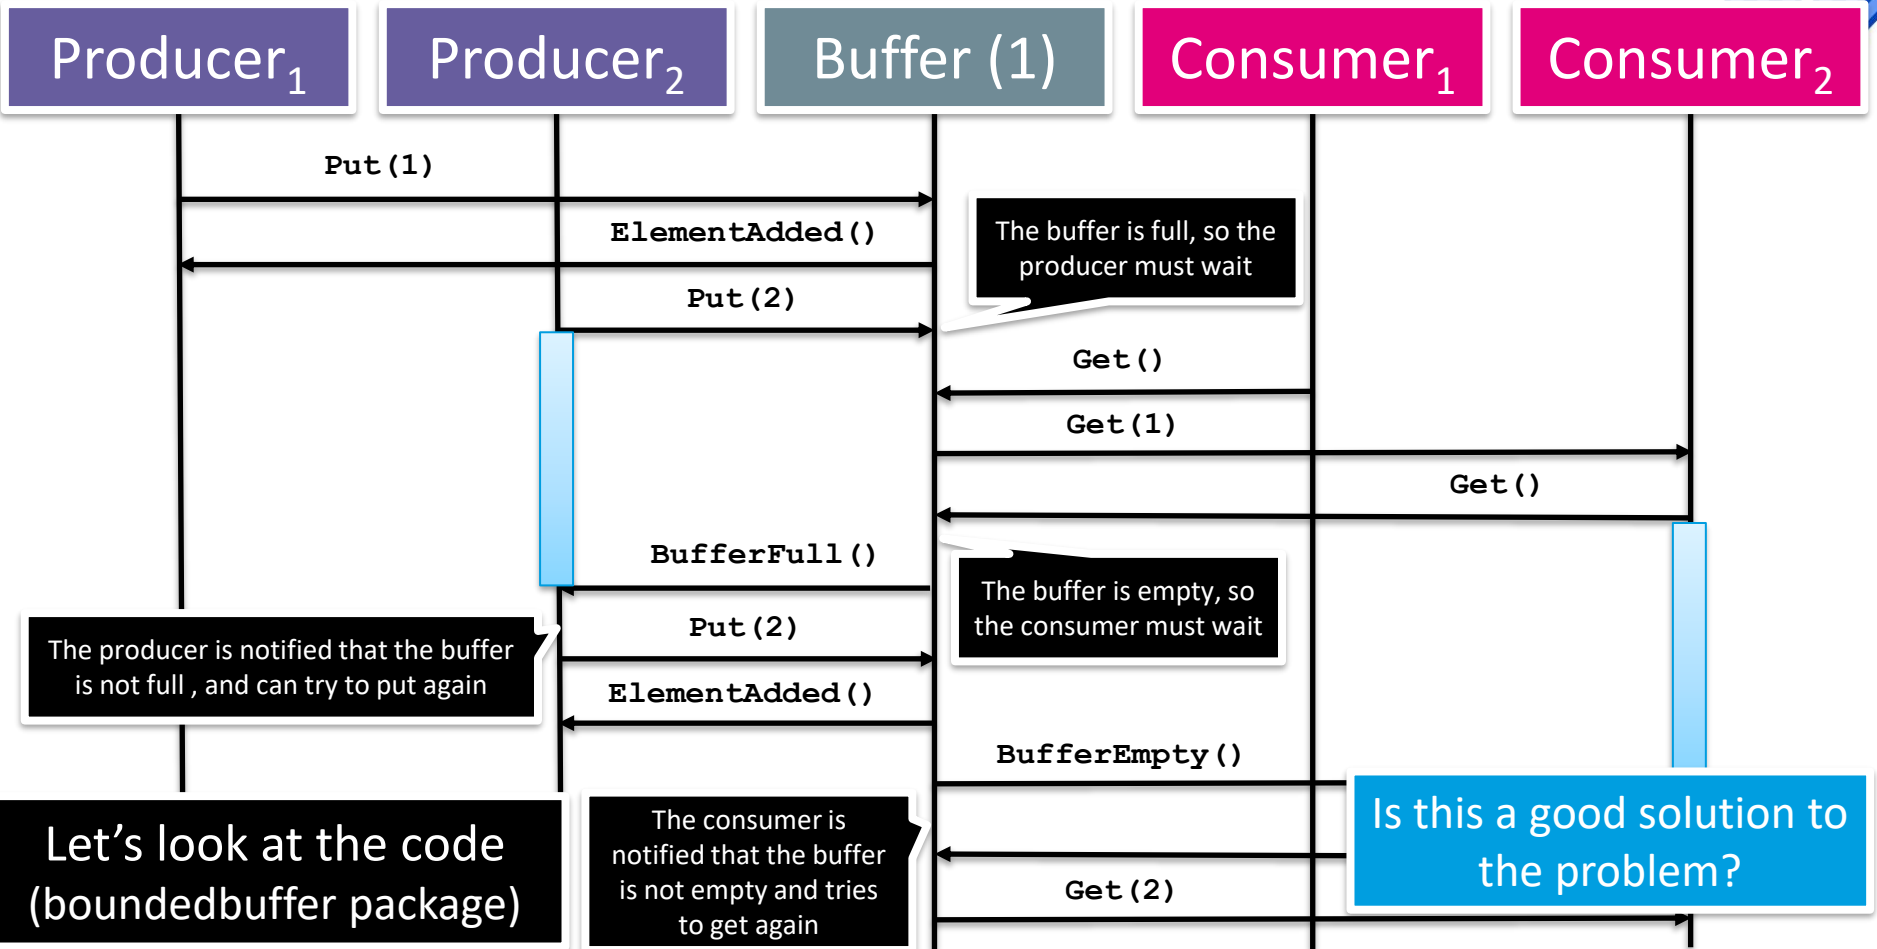
# Bounded Buffer (size 1) – execution example

.8



# Bounded Buffer (size 1) – execution example

.8




# Bounded Buffer (interesting) executions



Assuming FIFO mailboxes  
(Akka's default)

- Consider this execution
  1. Producer1 sends `put(1)` to the buffer
  2. Consumer1 sends `get()` to the buffer
  3. ...

Is it guaranteed that  
Consumer 1 will get the  
produced element?



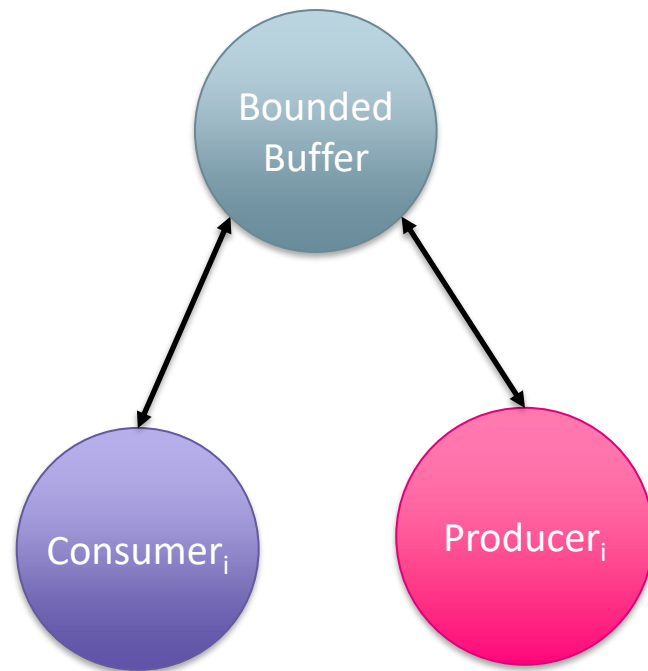
# Bounded Buffer (interesting) executions



Assuming FIFO mailboxes  
(Akka's default)

- Consider this execution
  1. Producer1 sends `put(1)` to the buffer
  2. Producer2 sends `put(2)` to the buffer
  3. Consumer1 sends `get()` to the buffer
  4. ...

Is it guaranteed that Consumer1  
will get either 1 or 2?



# Bounded Buffer (interesting) executions

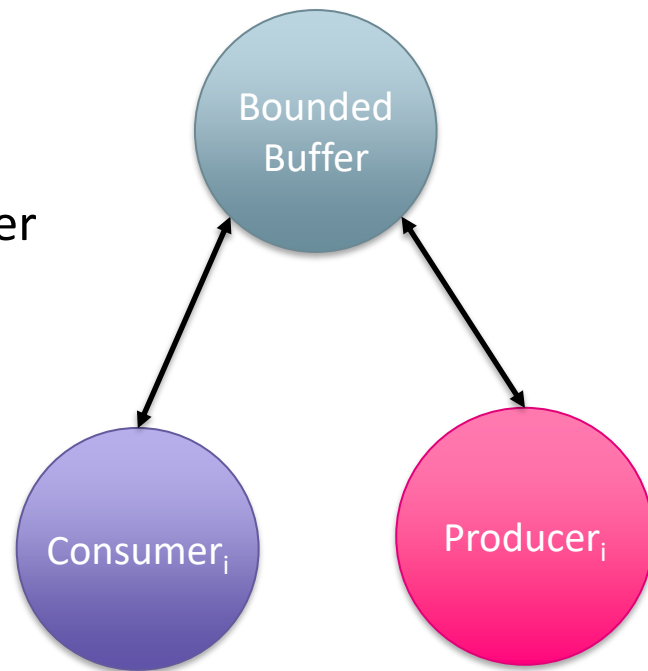


Assuming FIFO mailboxes  
(Akka's default)

- Consider this execution

1. Producer1 sends put(1) to the buffer
2. Producer2 sends put(2) to the buffer
3. Producer1 receives ElementAdded() from the buffer
4. Consumer1 sends get() to the buffer
5. ...

Is it guaranteed that  
Consumer 1 will get the  
produced element?






# Happened-before in distributed systems

· 12

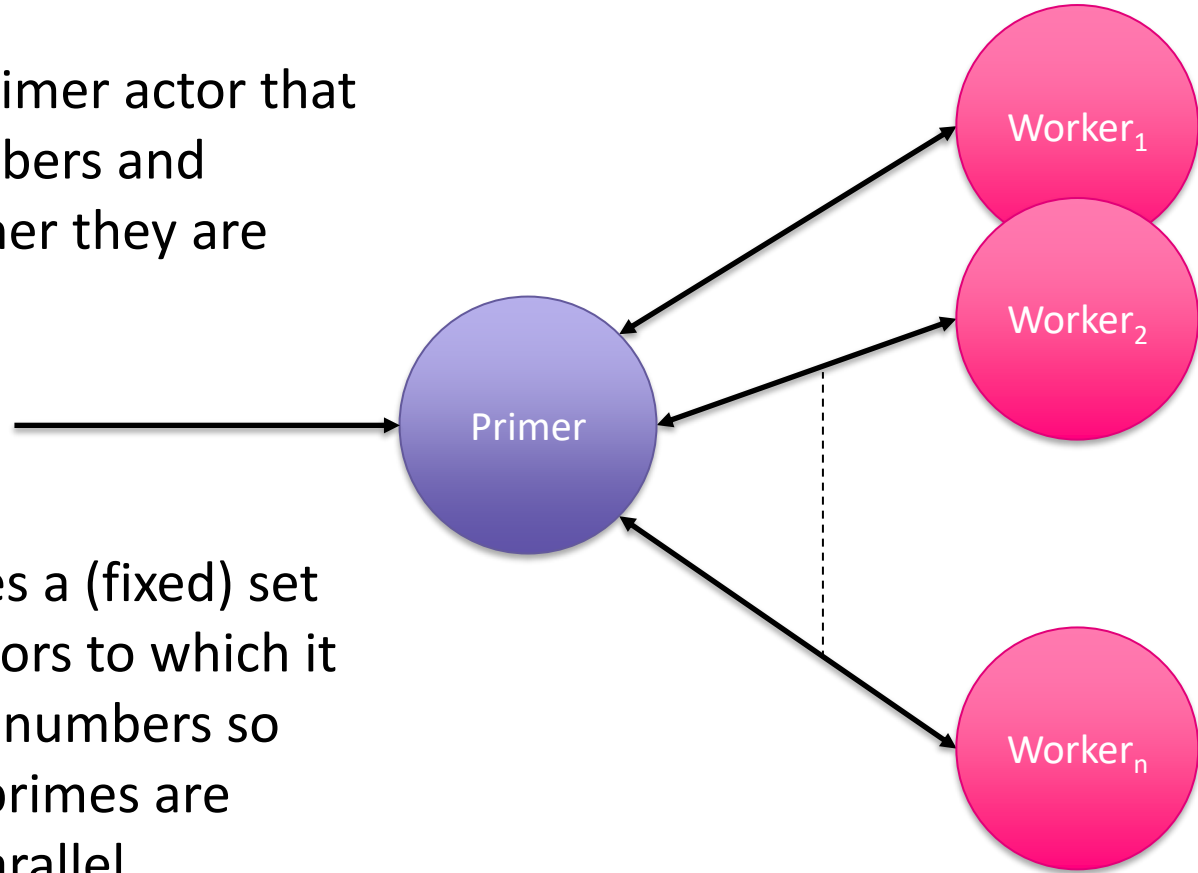


- Note that in the previous questions the behaviour of the systems depends on the reception of messages
- Thus, the happened-before relation defined by Lamport is useful in reasoning about actor systems
  - An action  $a$  happens-before an action  $b$  if they belong to the same actor and  $a$  was executed before  $b$
  - A  $\text{send}(m)$  action happens-before its corresponding  $\text{receive}(m)$
- Note the similarity with the happens-before relation of the Java memory model
  - We reason about message exchange instead of locking (inherent coordination problems remain, i.e., “semantic” deadlock & starvation)
  - Visibility issues disappear as actors only access local memory




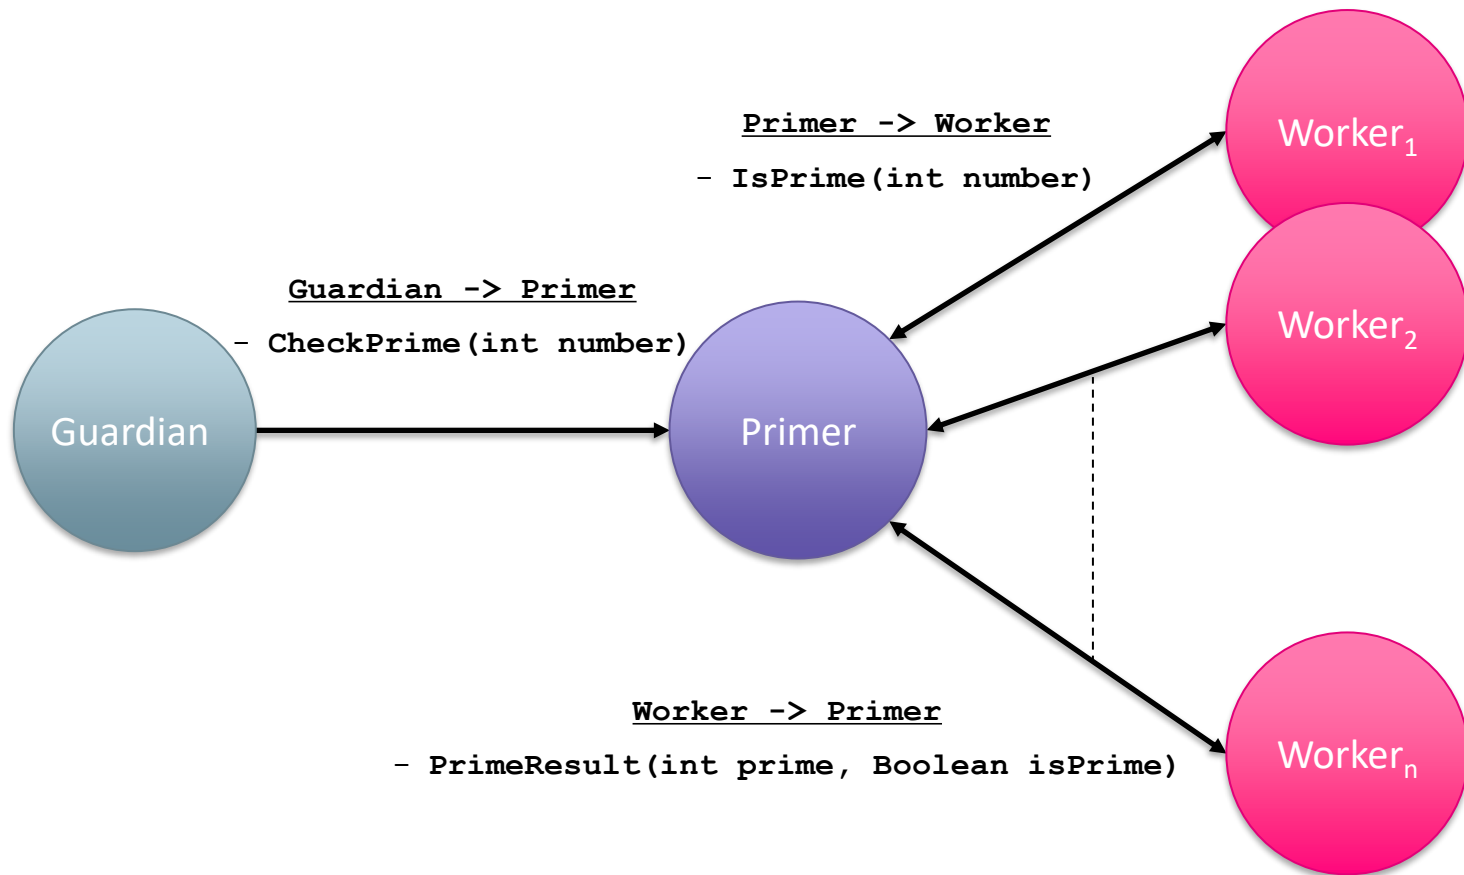


- Consider a Primer actor that receives numbers and checks whether they are prime



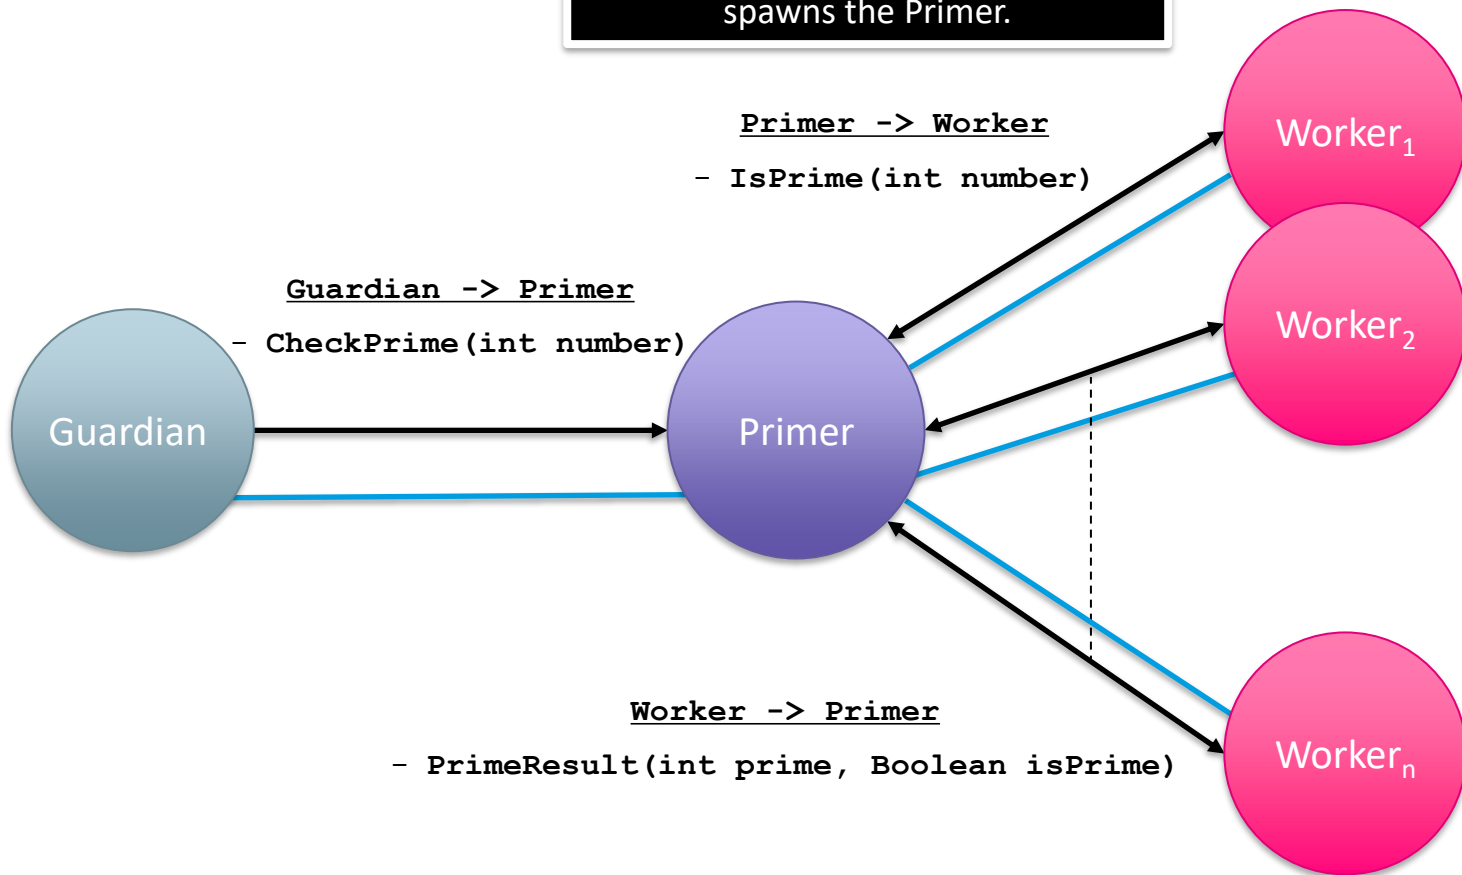
- The actor uses a (fixed) set of worker actors to which it forwards the numbers so that several primes are checked in parallel







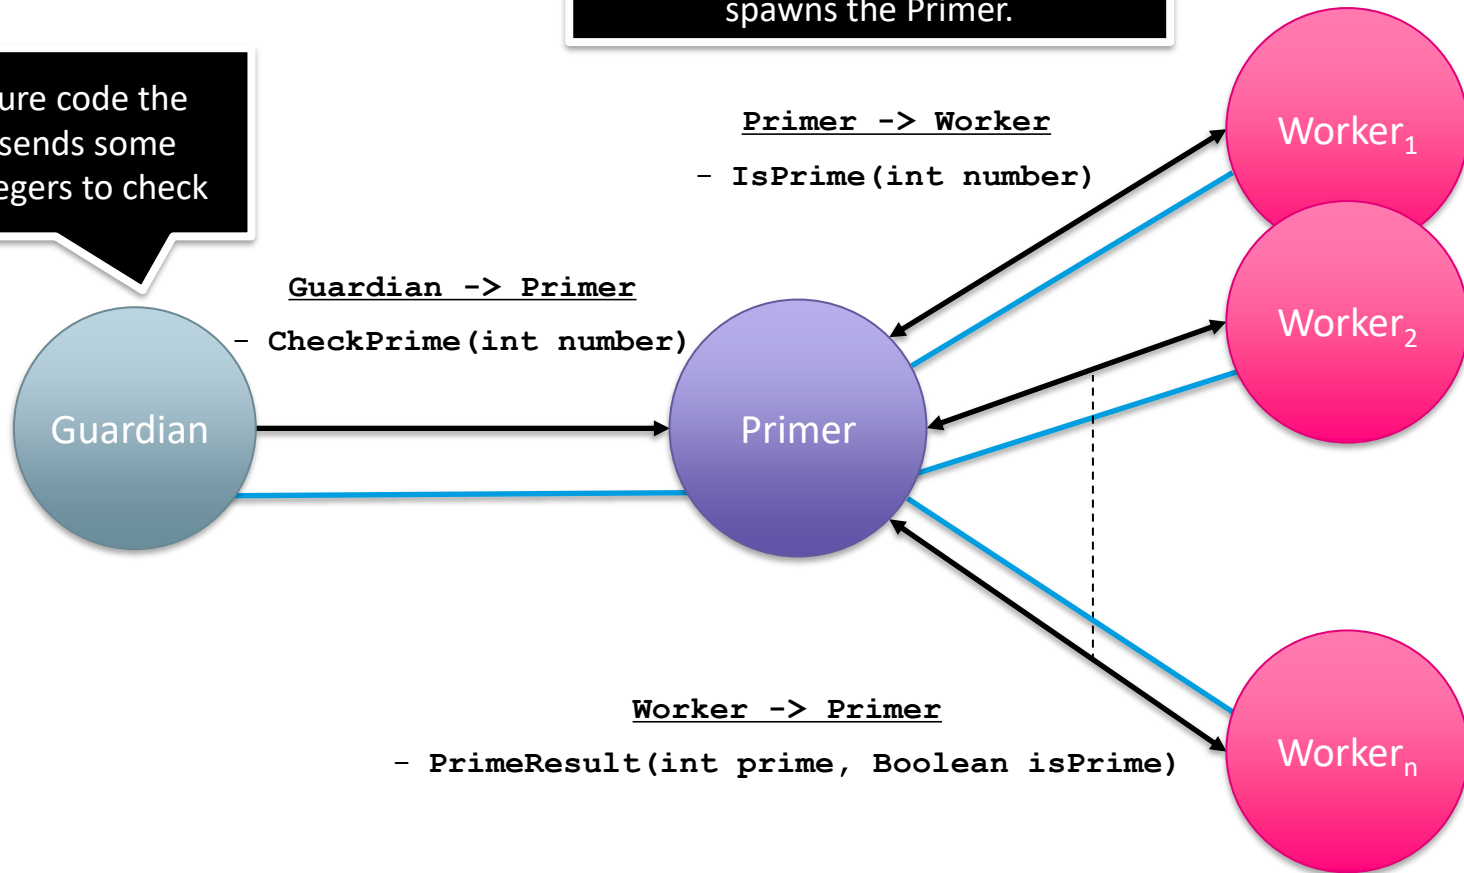
The Primer is in charge of spawning the workers. As usual, the guardian spawns the Primer.





In the lecture code the guardian sends some random integers to check

The Primer is in charge of spawning the workers. As usual, the guardian spawns the Primer.



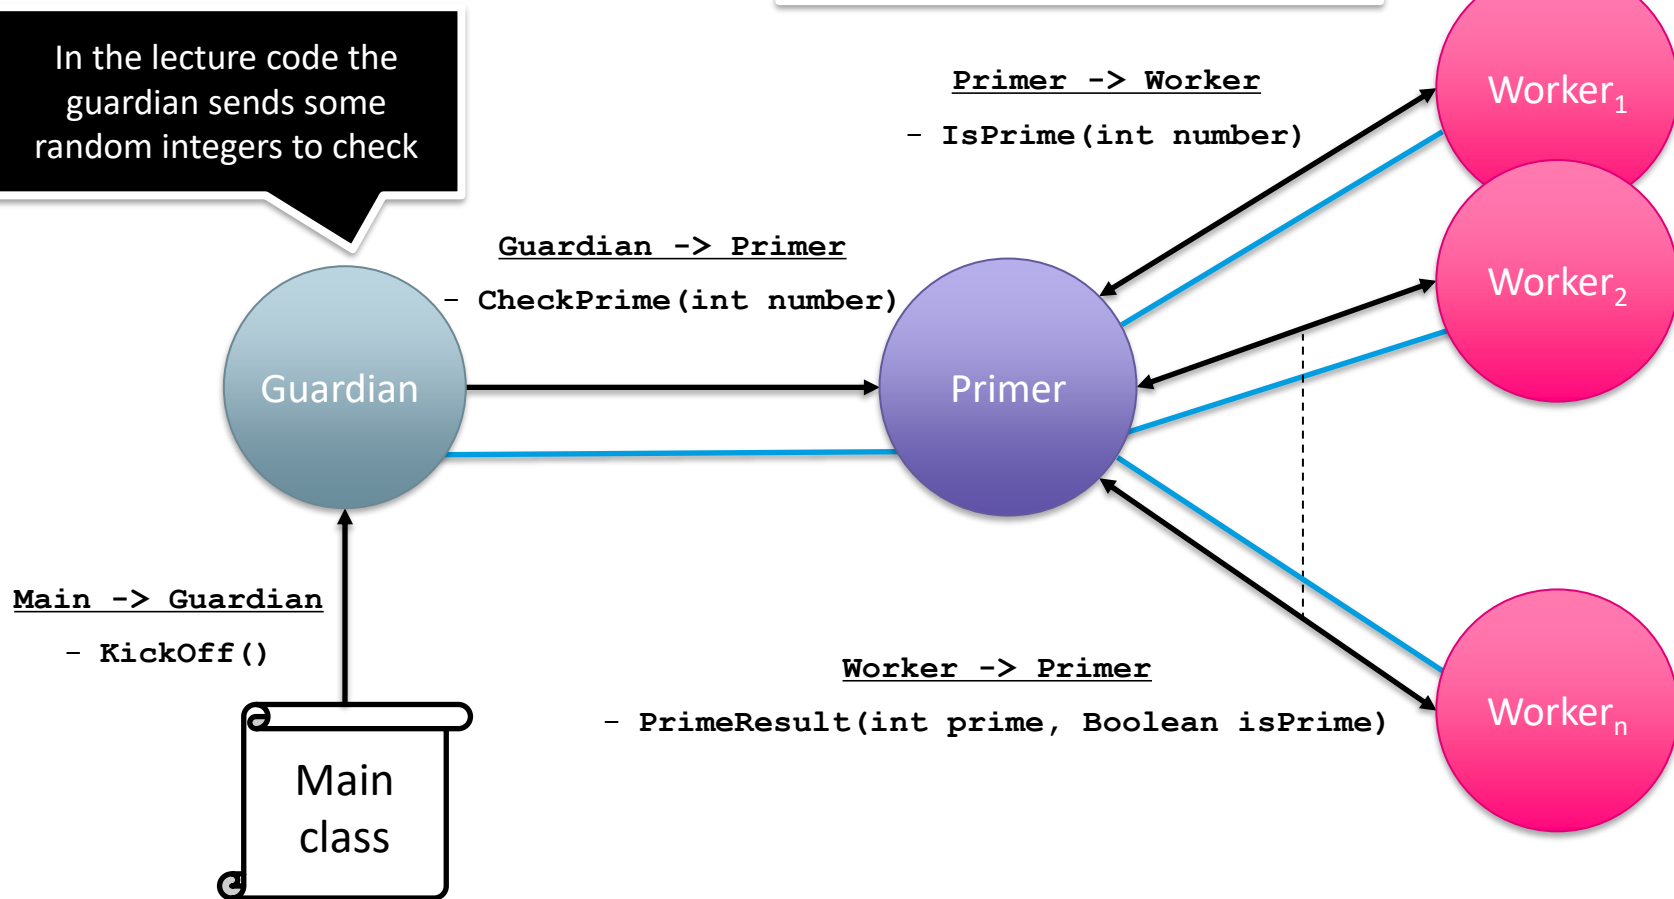
# Primer

· 14



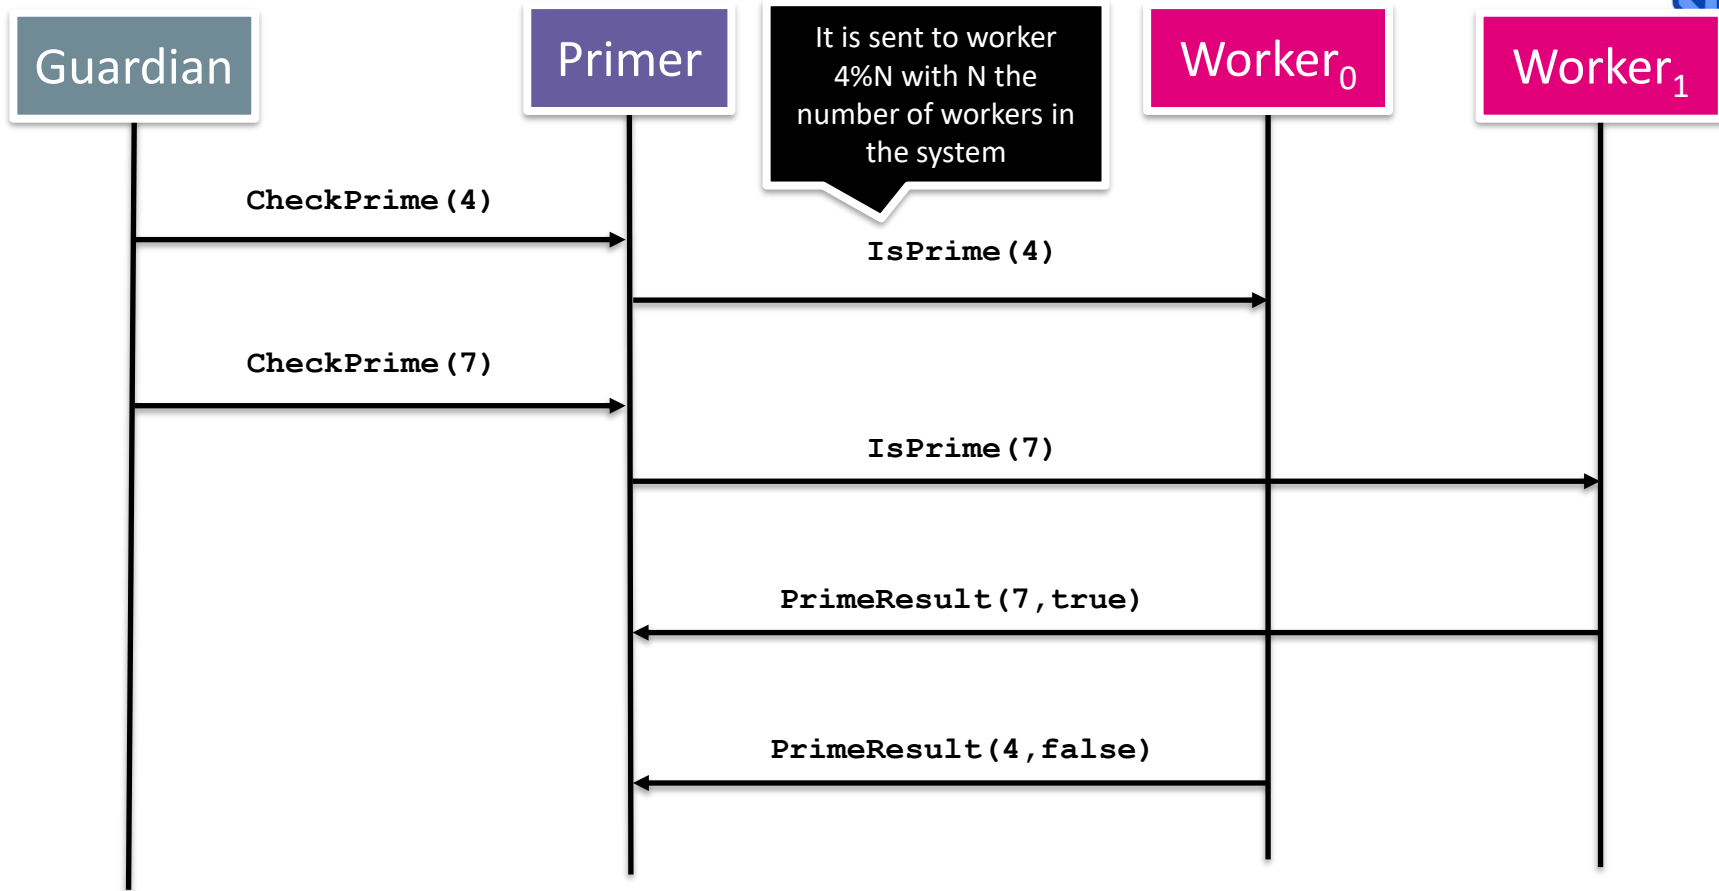
In the lecture code the guardian sends some random integers to check

The Primer is in charge of spawning the workers. As usual, the guardian spawns the Primer.



# Primer – execution example

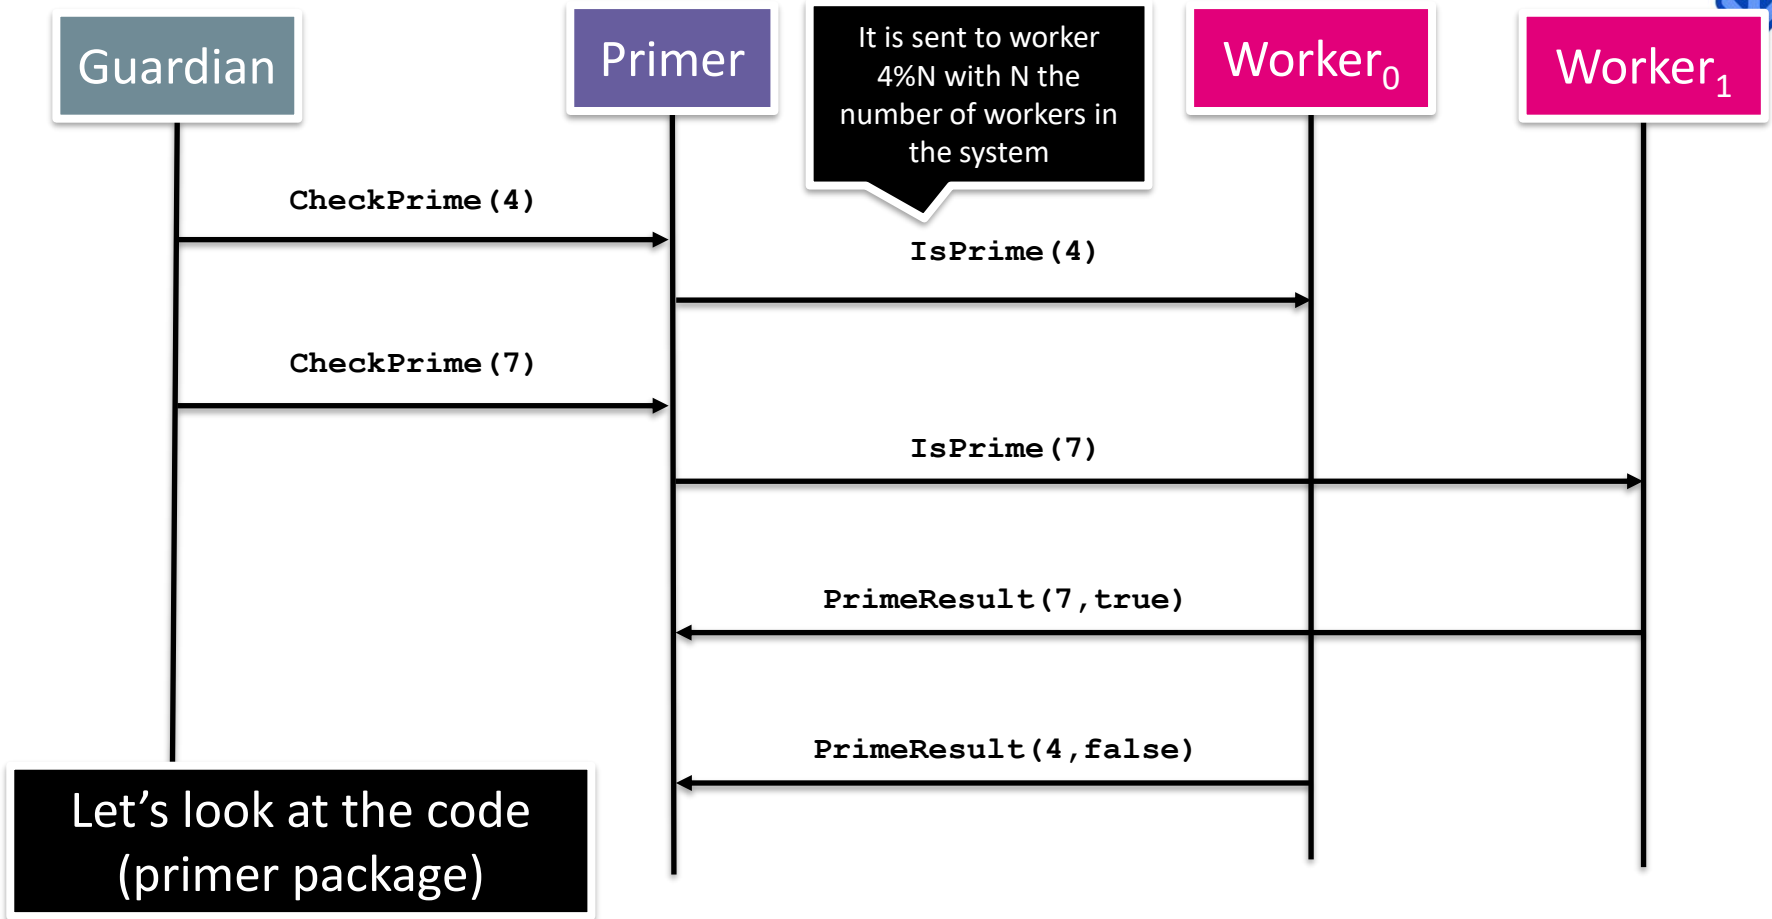
· 15





# Primer – execution example

· 15





- Note that the printing order of the results does not correspond to the order of sending the requests

```
pardo@pardo-work:week10lecture$ gradle run -PmainClass=primer.Main
```

```
> Task :app:run
[primer_system-akka.actor.default-dispatcher-3] INFO akka.event.slf4j.Slf4jLogger - Slf4jLogger started
>>> Press ENTER to exit <<<
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Server and workers started
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 21098598 is prime by worker worker_19
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 1001439 is prime by worker worker_20
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 47257026 is prime by worker worker_7
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 40857223 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 10667083 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 1001439 is not prime. [1/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 21098598 is not prime. [2/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [3/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [4/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 10667083 is not prime. [5/5]
```



- Note that the printing order of the results does not correspond to the order of sending the requests

```
pardo@pardo-work:week10lecture$ gradle run -PmainClass=primer.Main
```

```
> Task :app:run
[primer_system-akka.actor.default-dispatcher-3] INFO akka.event.slf4j.Slf4jLogger - Slf4jLogger started
>>> Press ENTER to exit <<<
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Server and workers started
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 21098598 is prime by worker worker_19
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 1001439 is prime by worker worker_20
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 47257026 is prime by worker worker_7
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 40857223 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 10667083 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 1001439 is not prime. [1/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 21098598 is not prime. [2/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [3/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [4/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 10667083 is not prime. [5/5]
```

How can this ordering happen?



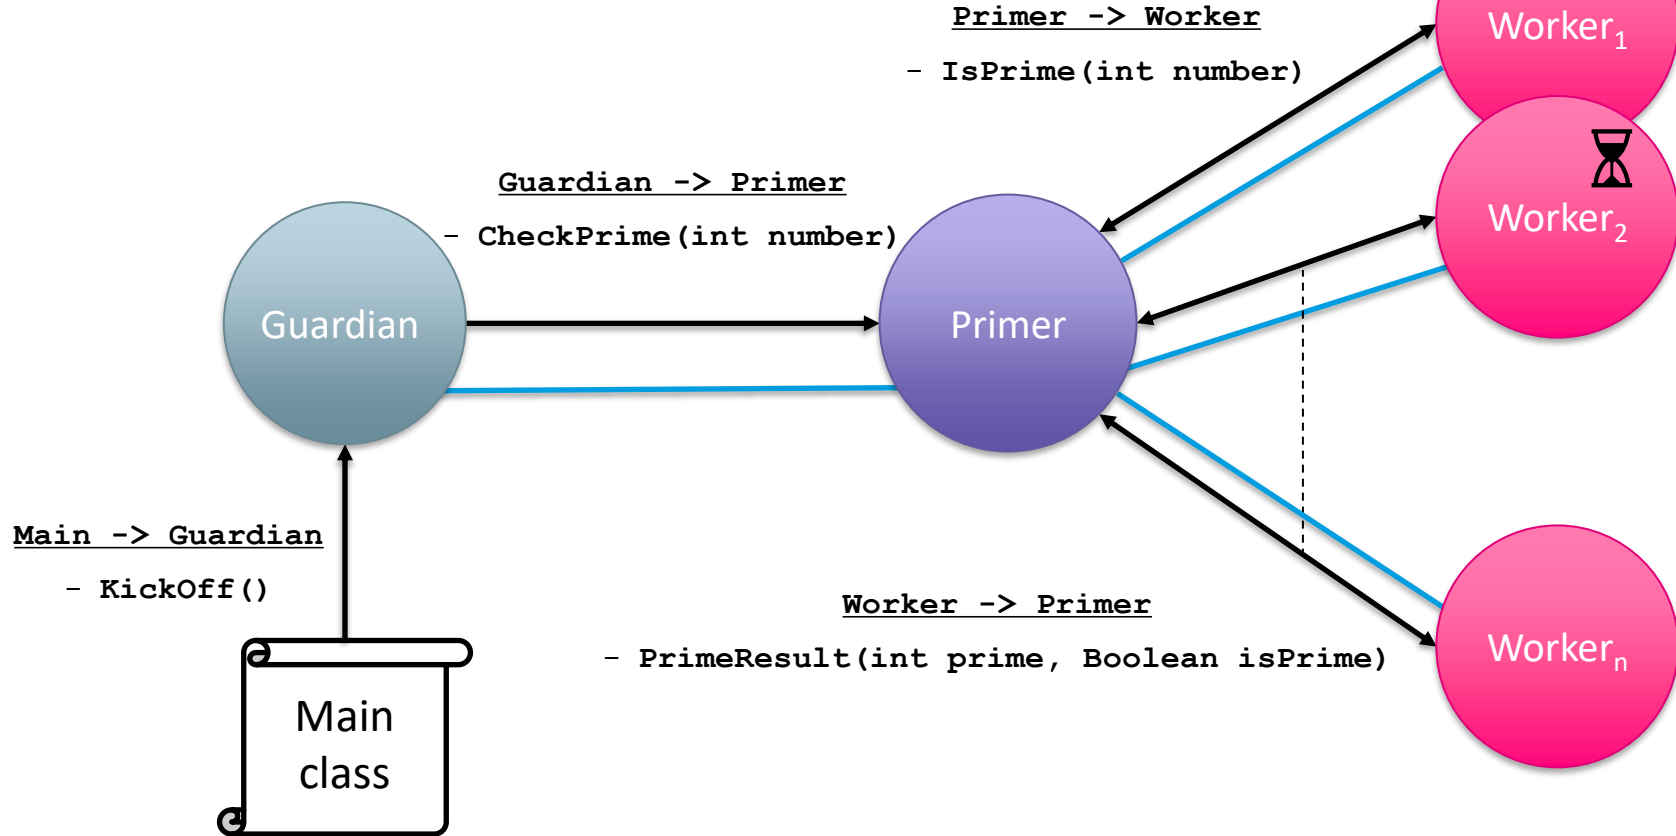
- Note that the printing order of the results does not correspond to the order of sending the requests

```
pardo@pardo-work:week10lecture$ gradle run -PmainClass=primer.Main
```

```
> Task :app:run
[primer_system-akka.actor.default-dispatcher-3] INFO akka.event.slf4j.Slf4jLogger - Slf4jLogger started
>>> Press ENTER to exit <<<
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Server and workers started
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 21098598 is prime by worker worker_19
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 1001439 is prime by worker worker_20
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 47257026 is prime by worker worker_7
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 40857223 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 10667083 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 1001439 is not prime. [1/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 21098598 is not prime. [2/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [3/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [4/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 10667083 is not prime. [5/5]
```

How would you change the system to print the results in the same order as they arrived?

What happens if one of the workers gets stuck working on a difficult prime?



# Actors systems with dynamic topology

# Don't be shy, spawn actors!



- The Actors model encourages creating many actors that perform small tasks and communicate with each other




# Do more actors improve performance?



· 21

- As usual, performance depends on the hardware
- These are the results of running the primer to check 1 million numbers between 1 billion and `Integer.MAX_VALUE`.
  - Not very strong statistics (4 runs for each number of actors).
- Akka implements actors systems using a `ForkJoinPool` (a version of the `ThreadPool`, which is more efficient for tasks with low dependencies)
- However, actor systems can be distributed among many JVMs and computers
  - We are not limited to a single computer throughput
  - See Akka cluster






# Do more actors improve performance?



· 21

That said, distributing computation among actors makes it easy to implement fault-tolerant systems and adaptive load-balancing


- Akka implements actors systems using a ForkJoinPool (a version of the ThreadPool, which is more efficient for tasks with low dependencies)
- However, actor systems can be distributed among many JVMs and computers
  - We are not limited to a single computer throughput
  - See Akka cluster



# Topology (Actor hierarchy)




- We use the term *topology* to refer to the parental structure of actors in the system
  - In Akka, this structure is a tree, and it is called a *hierarchy*.
- The systems we have seen so far feature a *static* topology
  - All the actors in the system are spawned during initialization



Solid lines and actors represent elements that are created during initialization and never change



- Actor systems with static topology may not exploit computational resources effectively
  - As we saw, the system may slow down if some actors are consuming excessive computational resources
  - Actors may also crash, and the system should be able to recover from this (fault-tolerance)
- The advantages of the actors model are better exploited when the system can adaptively decide the number of workers
- Actors should be seen as *nice co-workers*
  - *A group of computational resources that collaborate to achieve a common goal*



Dashed lines and actors represent elements that may be created dynamically (on-demand, after initialization)



- To avoid excessive delays by primes that are difficult to check, we extend the system with dedicated actors whose only task is to check the prime
- After these dedicated actors have finished the computation they report the result and terminate the execution

# Primer with job workers


· 25



Every time that a worker receives a prime to check, it spawns a SingleJobWorker to perform the computation

Primer -> Worker

- IsPrime(int number)



Worker -> Primer

- PrimeResult(int prime, Boolean isPrime)

# Primer with job workers

Worker -> SingJobW

- IsPrime(int number,  
ActorRef server)


After begin spawned, the job is  
forwarded to the single job worker

Primer -> Worker

- IsPrime(int number)

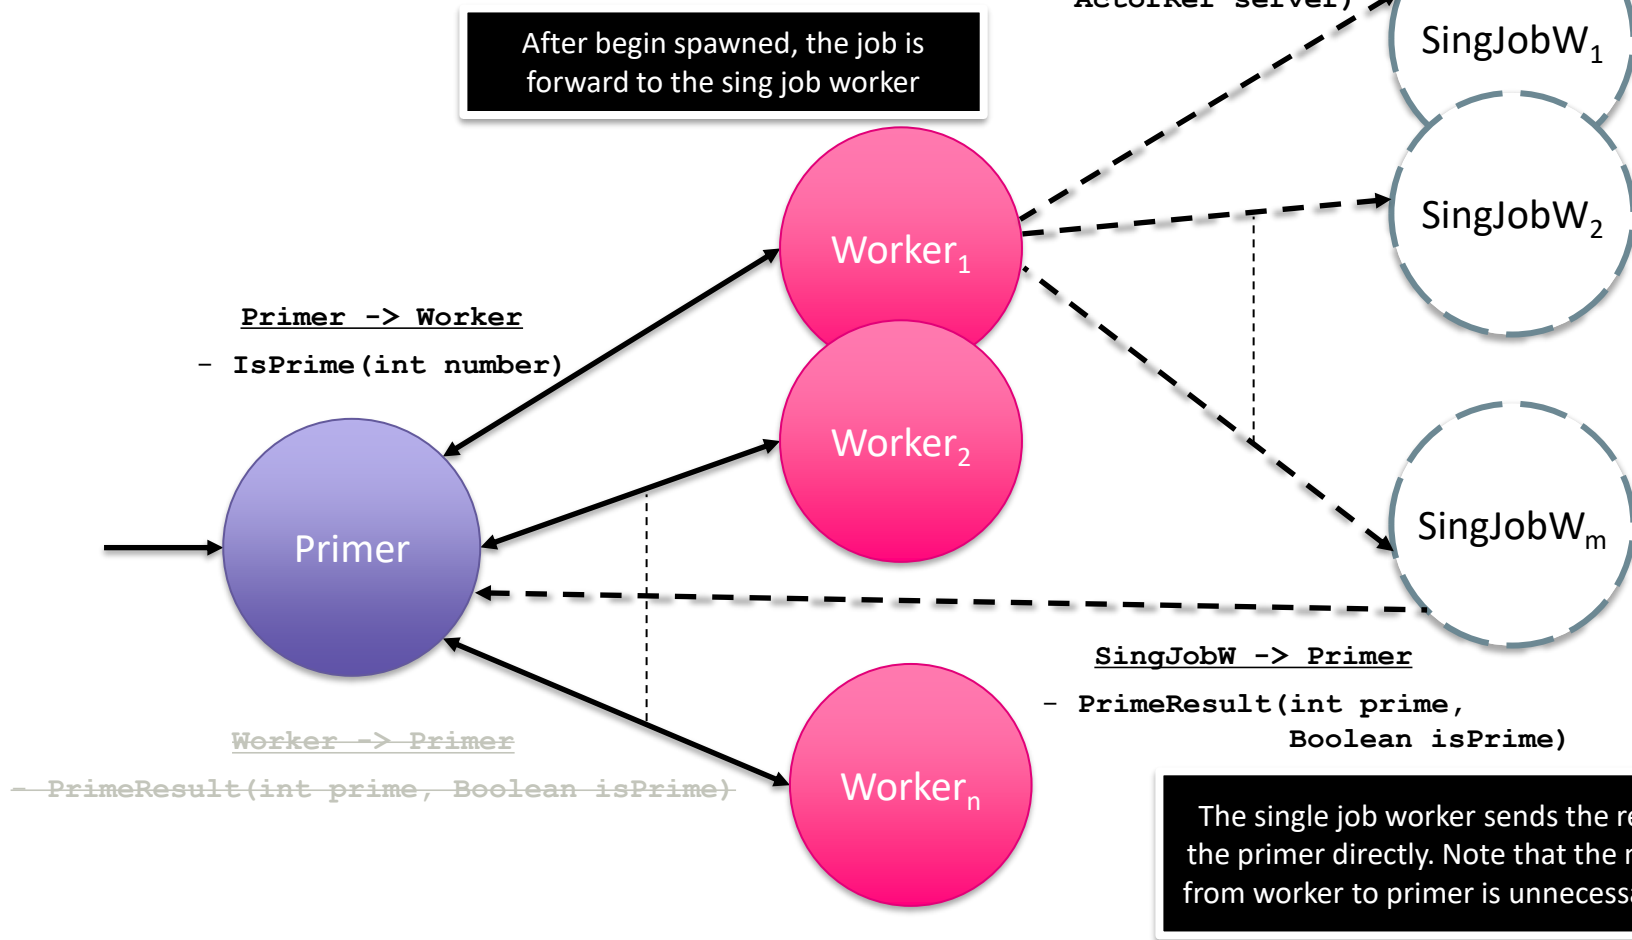
Worker -> Primer

- PrimeResult(int prime, Boolean isPrime)



# Primer with job workers

· 27



# Primer with job workers

Worker -> SingJobW

- IsPrime(int number,  
ActorRef server)

After begin spawned, the job is  
forward to the sing job worker

Primer -> Worker

- IsPrime(int number)

Would there be any problem if we  
send the message to the parent  
worker instead? (and it forwards it  
to the primer?)

SingJobW -> Primer

- PrimeResult(int prime,  
Boolean isPrime)

Worker -> Primer

- PrimeResult(int prime, Boolean isPrime)

The single job worker sends the result to  
the primer directly. Note that the message  
from worker to primer is unnecessary now.





# Primer with job workers

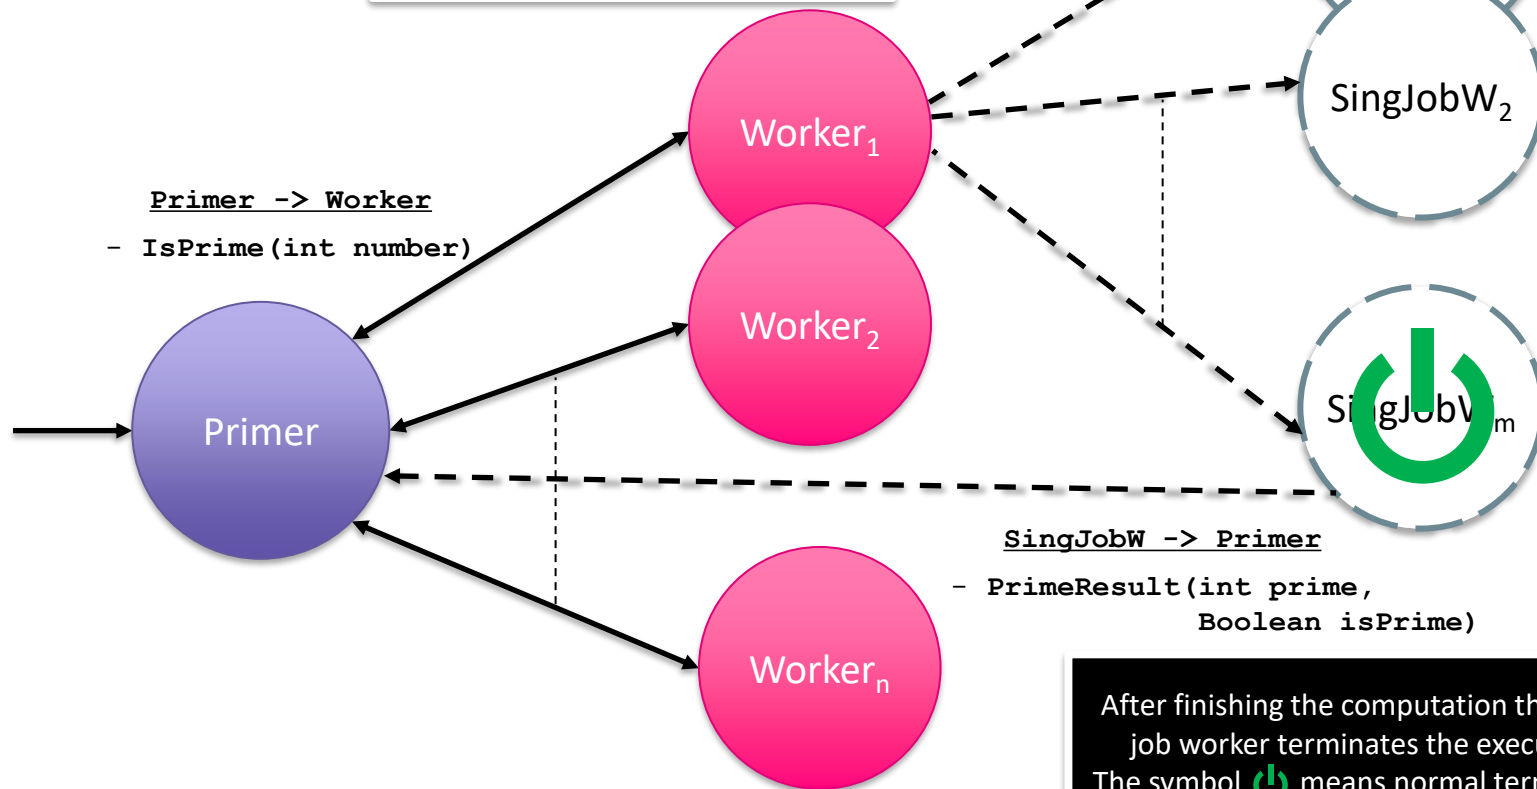
Worker -> SingJobW

- IsPrime(int number,  
ActorRef server)

After begin spawned, the job is  
forward to the sing job worker


Primer -> Worker

- IsPrime(int number)



SingJobW -> Primer

- PrimeResult(int prime,  
Boolean isPrime)

After finishing the computation the single  
job worker terminates the execution.  
The symbol  means normal termination



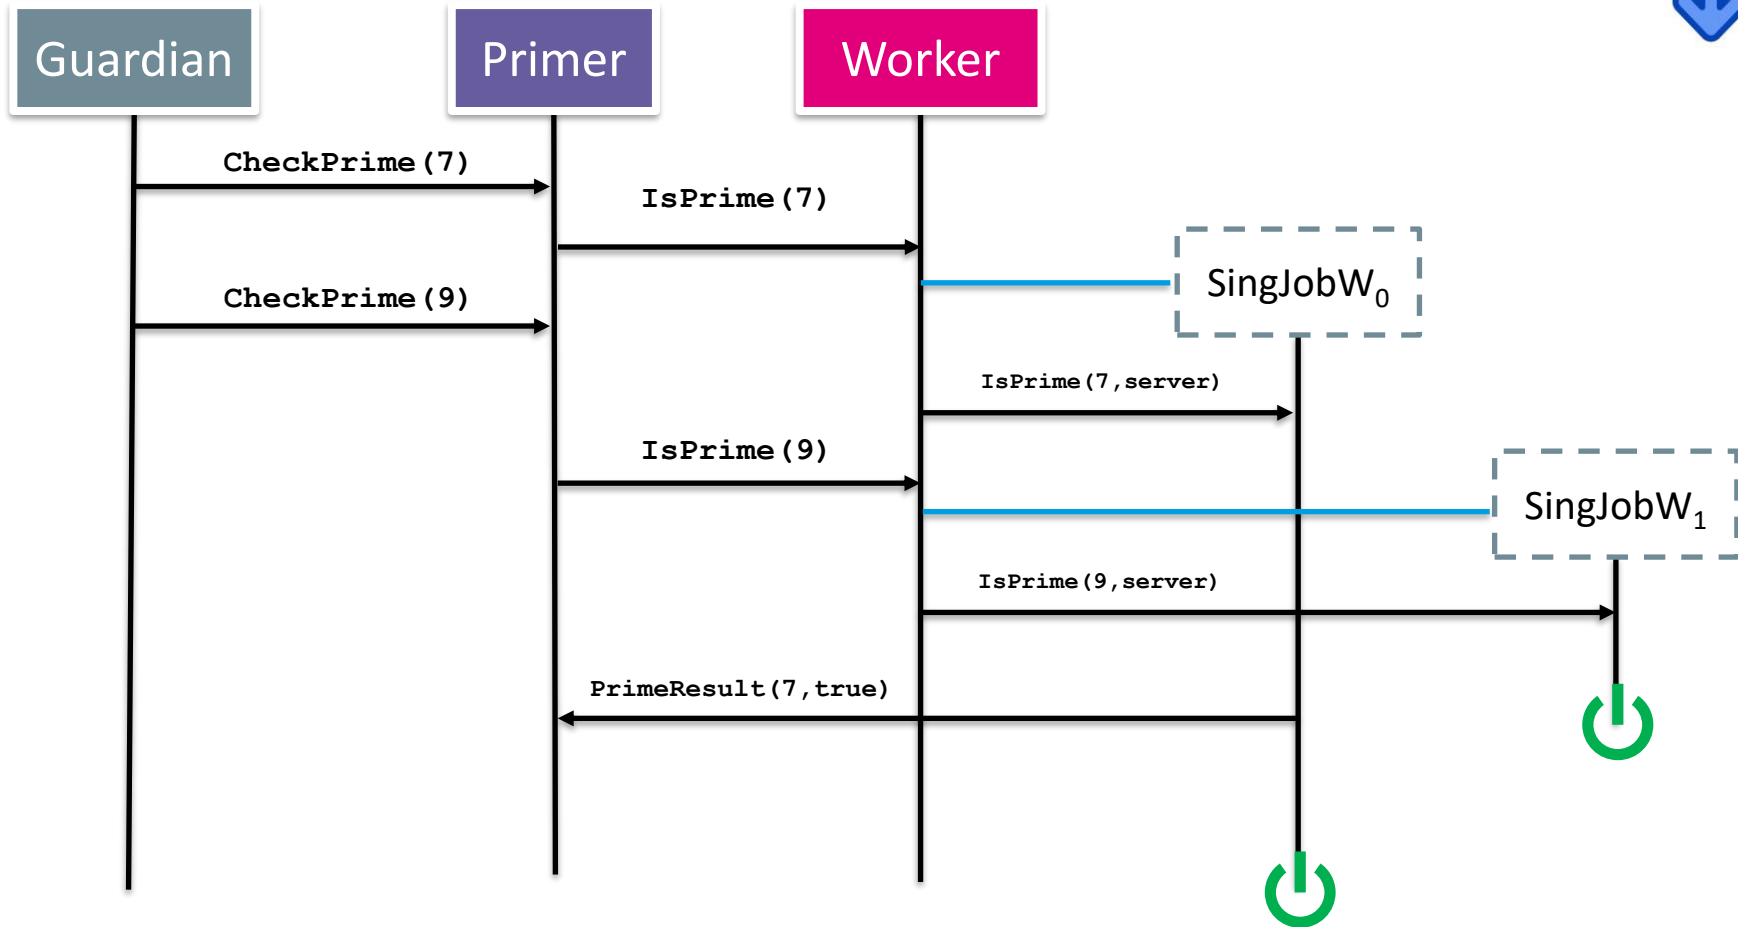
- The Akka API has a specific method to shutdown an actor
  - `Behaviours.stopped()`
- This behaviour can be used as the return object for a message handler

```
public Behavior<IsPrime> onIsPrime(IsPrime msg) {  
    msg.server.tell(new Primer.PrimeResult(msg.number, isPrime(msg.number)));  
    return Behaviors.stopped();  
}
```

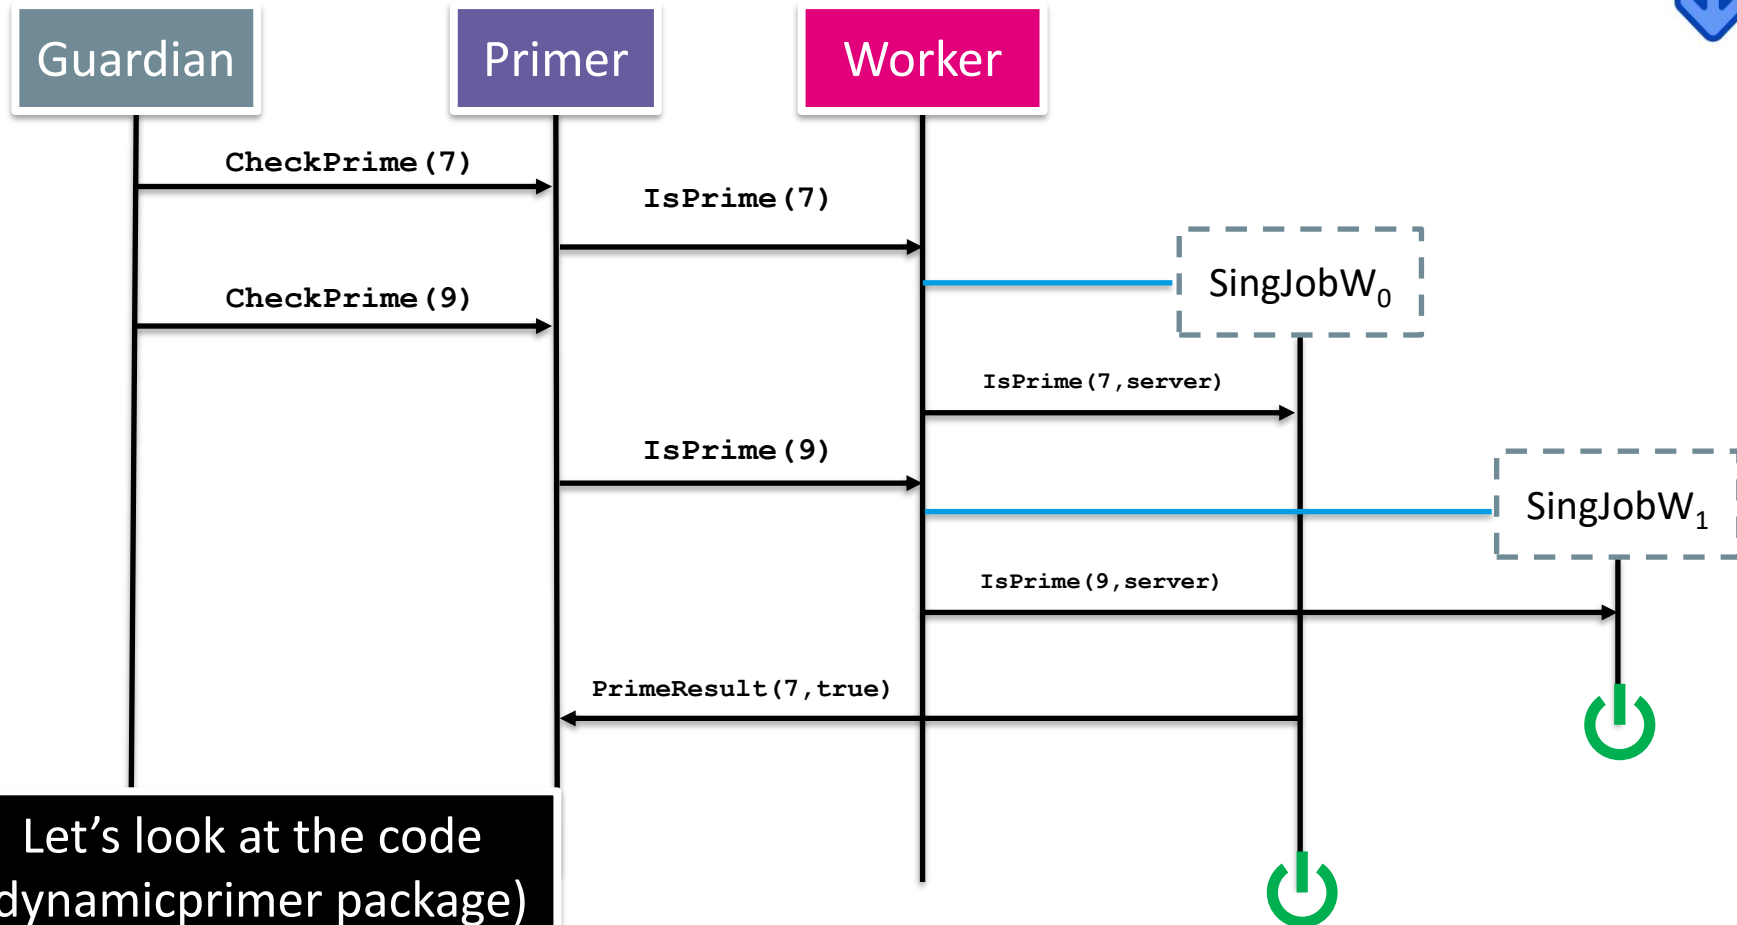
# Primer with job workers – execution example



· 30



# Primer with job workers – execution example

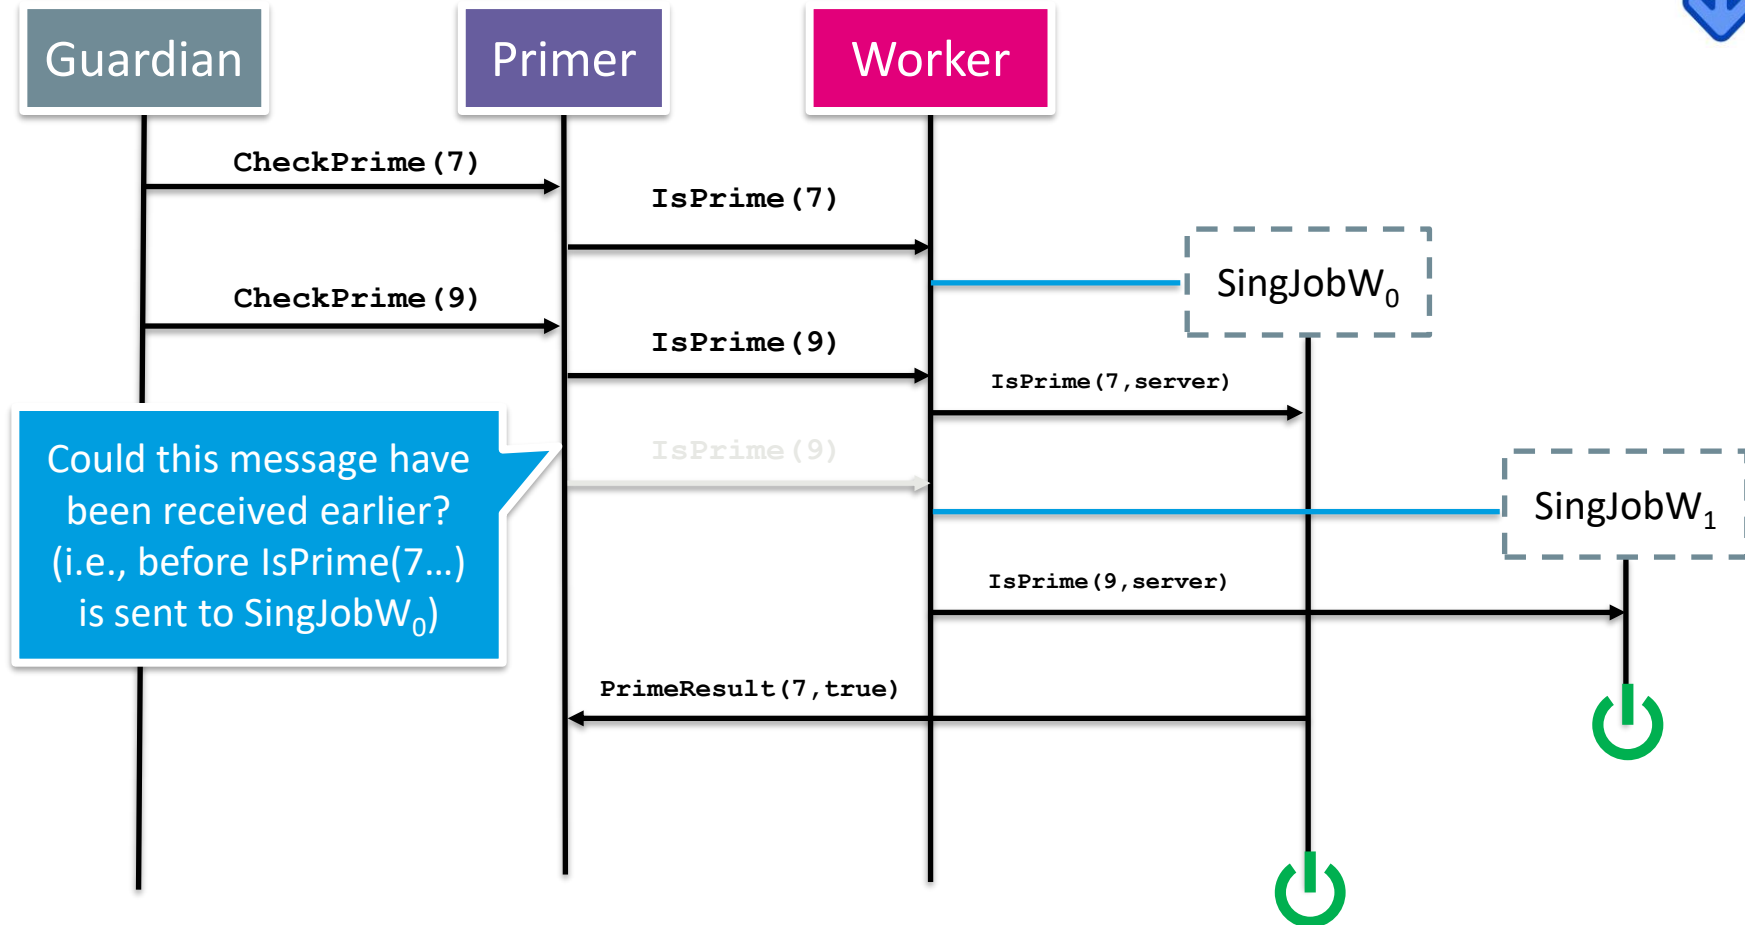


Let's look at the code  
(dynamicprimer package)

# Primer with job workers – execution example



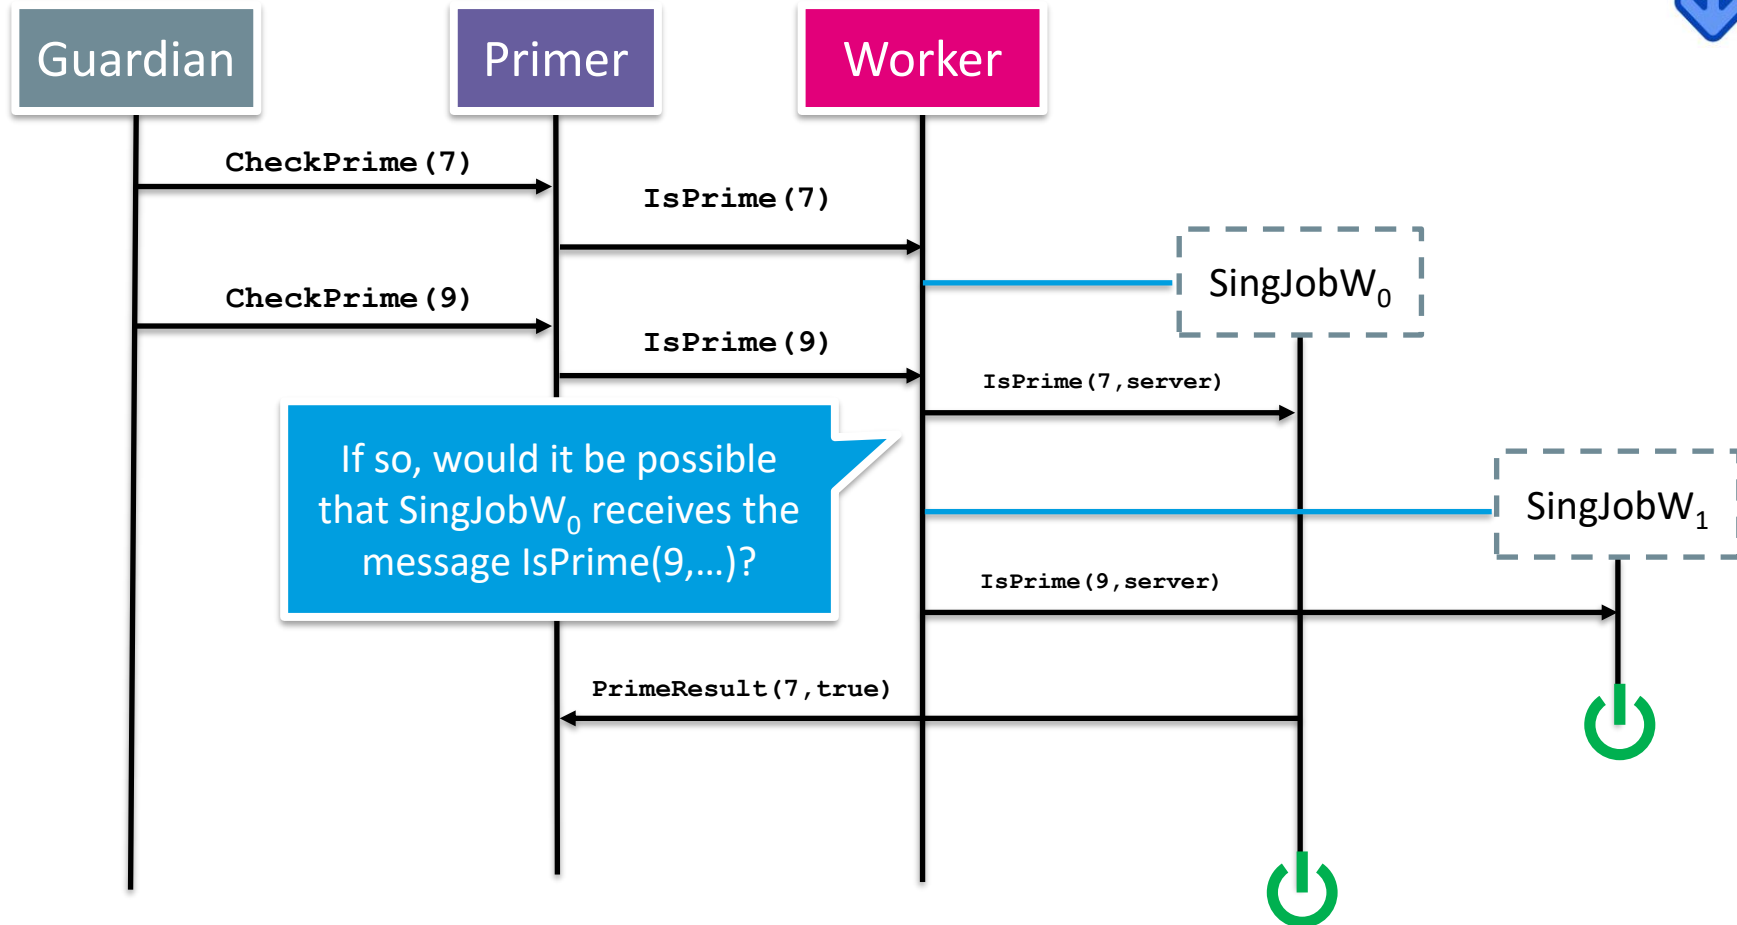
· 31



# Primer with job workers – execution example



· 32



• 33

```
- IsPrime(int number,
          ActorRef server)
```

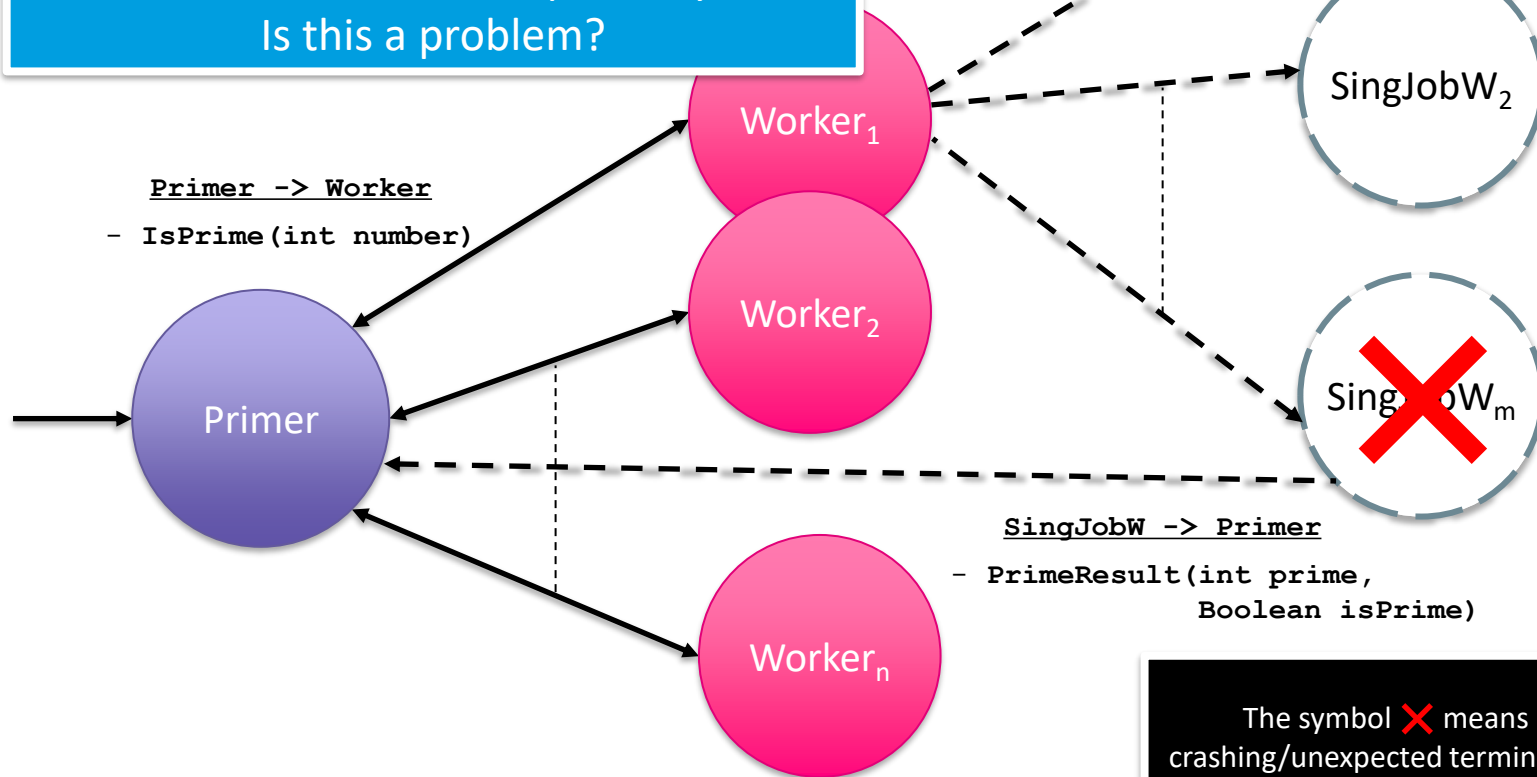


# Primer with job workers

· 34



What happens if one of the single job workers crashes unexpectedly?  
Is this a problem?






# Fault-tolerance in Akka

# Let it crash! model



· 36

- Actor libraries and programming languages encourage a *let it crash* programming model
- Do not put a lot of effort ensuring that actors never crash
  - Assume that things will fail
- Develop actors systems ensuring that if an actors crashes the system can recover
- Specially useful in distributed systems when you cannot predict what type of message you will receive





- Akka implements supervision mechanisms to react to failures
- Children may inform their parents when they terminate or fail
- Actors may use the function `watch(ActorRef<T> actor)` to supervise their children
- If an actor is being supervised by a (parent) watcher, it sends to the watcher
  - A `ChildFailed` signal, if it crashed due to an exception
  - A `Terminated` signal, if it terminates normally

But `ChildFailed` extends  
from `Terminated`

- A *signal* can be seen as a message that is automatically sent by Akka
- For a watcher to handle signals, it *must* have a message handler with `onSignal(Signal.class, Function f)`
- The message send by the signal contains a reference to the sender actor. It can be accessed with `getRef()`.

```
/* --- Message handling ----- */  
@Override  
public Receive<T> createReceive() {  
    return newReceiveBuilder()  
        .onMessage(Message.class, this::onMessage)  
        // Here order matters `ChildFailed` extends `Terminated`  
        .onSignal(ChildFailed.class, this::onChildFailed)  
        .onSignal(Terminated.class, this::onTerminated)  
        .build();  
}
```


Note I: When processing a message/signal, the message handler picks the handler that first matches the class of the message.

Note II: Since `ChildFailed` extends `Terminated`, if `onSignal(Terminated,...)` appears before `onSignal(ChildFailed,...)`, when a `ChildFailed` signal arrives, the latter `onSignal` will not be triggered.


# Actor supervision (graphically)




· 39



# Actor supervision (graphically)




· 40



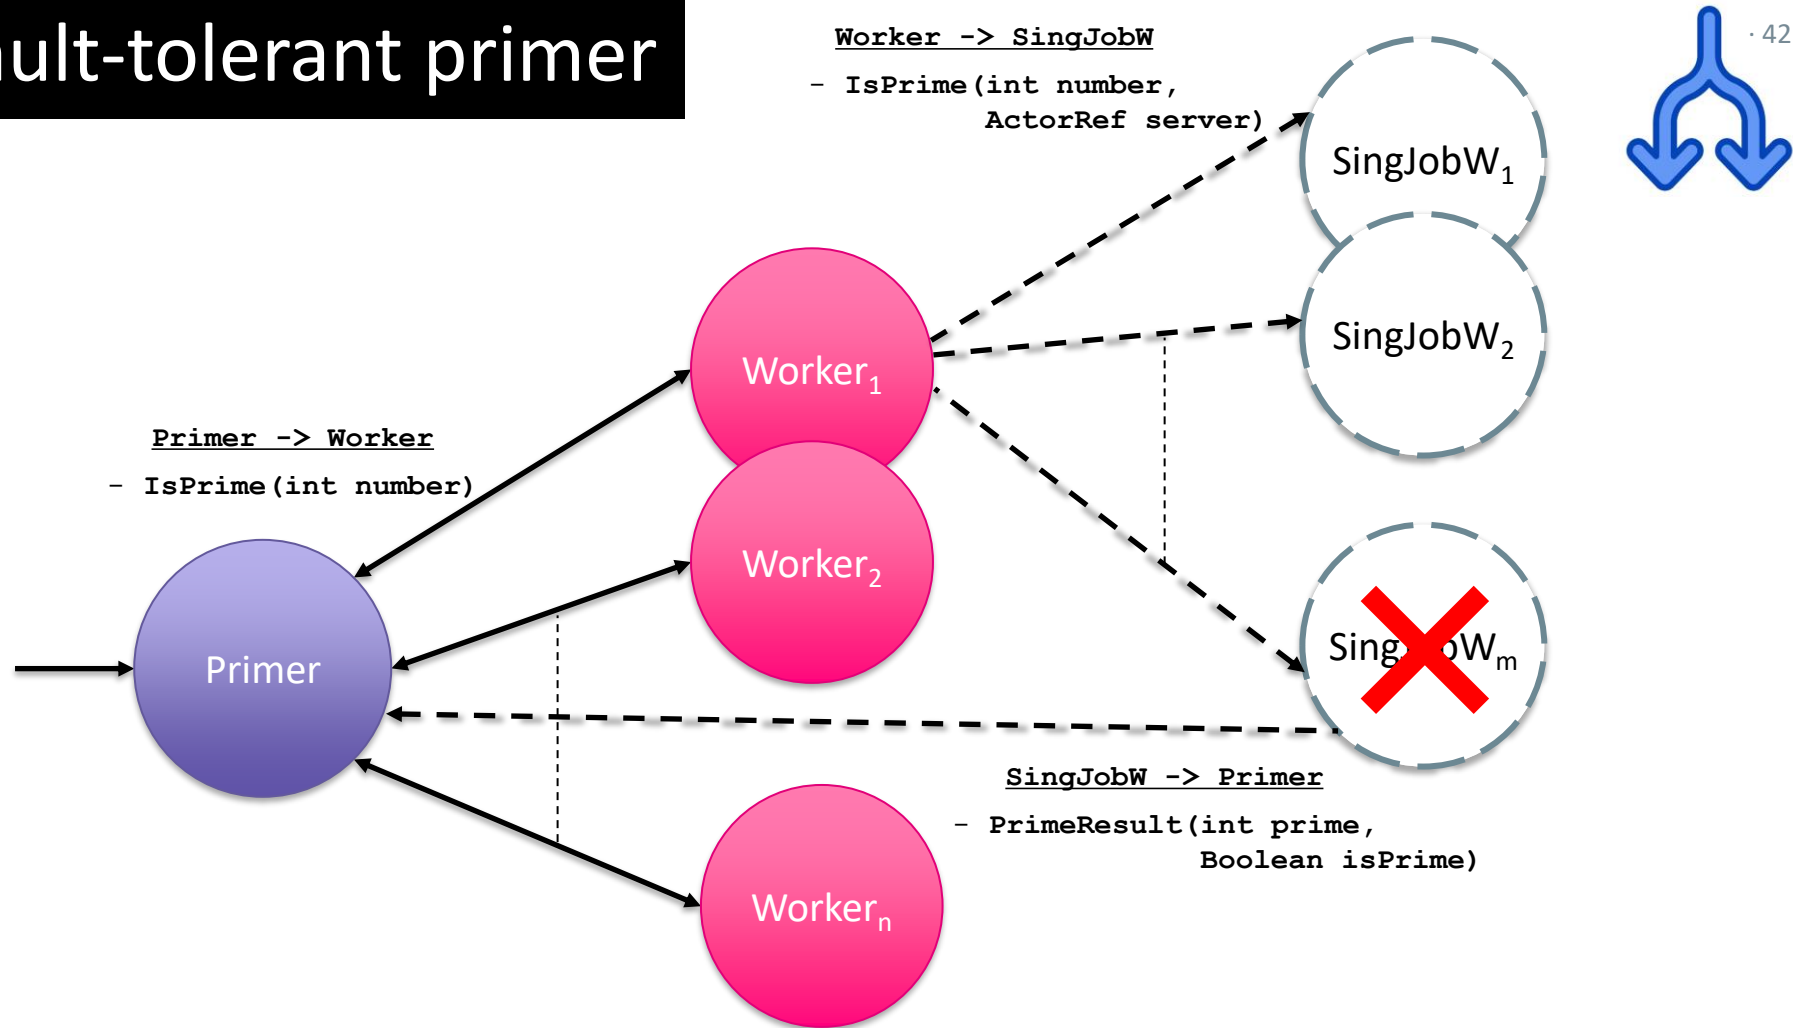
The dashed-dotted yellow arrow indicates the sending of a signal. These are sent automatically by Akka as part of the supervision functionality. If the child finishes normally, it sends a Terminated signal.

# Actor supervision (graphically)



If the child finishes with an exception, it sends a `ChildFailed` signal.

# Fault-tolerant primer






# Fault-tolerant primer

· 43

- We extend the primer to handle the case when a single job worker fails
- To this end, the worker needs to:
  1. Watch all the actors it spawns
  2. Handle ChildFailed signals
    - The handler spawns a new worker and sends the number again to check whether it is prime
  3. Handle Terminated signals
    - No more computation needed, we can mark the number as checked




# Fault-tolerant primer

· 43

- We extend the primer to handle the case when a single job worker fails
- To this end, the worker needs to:
  1. Watch all the actors it spawns
  2. Handle ChildFailed signals
    - The handler spawns a new worker and sends the number again to check whether it is prime
  3. Handle Terminated signals
    - No more computation needed, we can mark the number as checked

Do we need to extend the state of Worker actors to handle fault-tolerance?

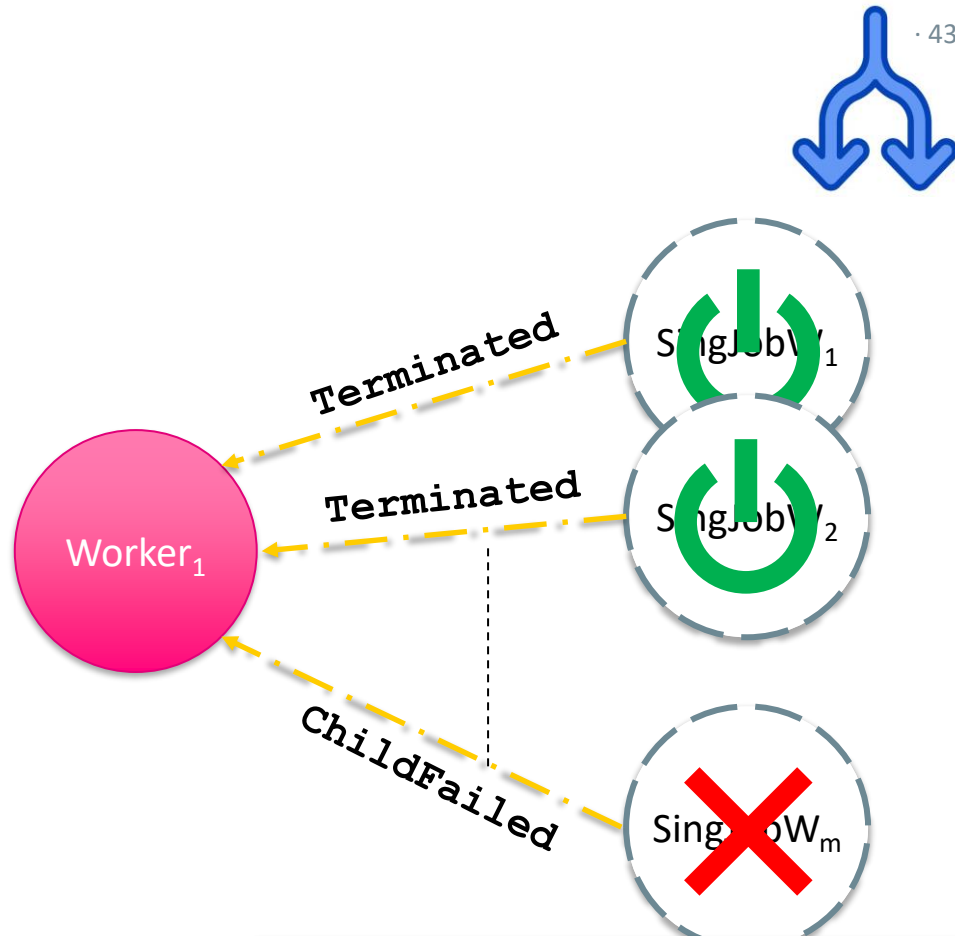


# Fault-tolerant primer

· 43

- We extend the primer to handle the case when a single job worker fails
- To this end, the worker needs to:
  1. Watch all the actors it spawns
  2. Handle ChildFailed signals
    - The handler spawns a new worker and sends the number again to check whether it is prime
  3. Handle Terminated signals
    - No more computation needed, we can mark the number as checked

Do we need to extend the state of Worker actors to handle fault-tolerance?



Let's look at the code (faulttolerantprimer package)

# Adaptive load balancing



- Load balancing refers to the process of distributing a set of tasks over a set of resources (computing units), with the aim of making their overall processing more efficient. [Wikipedia]
- In the (static) primer system, we indiscriminately spawned processes to perform tasks
  - This may cause sending tasks to busy workers while other idle workers could be processing them
- There exists some patterns that aim at distributing computation fairly among actors.
  - For instance, the scatter-gather pattern



- Scatter-Gather is a common design pattern in distributed systems that can be easily implemented with actors
- Typically, the level of scattering (i.e., number of spawned actors) depends on the size of the problem to solve (dynamic load balancing)
  - But it can also be limited by other factors, e.g., CPU or memory usage
- A scatter-gather system contains two main types of actors
  - Scatterer: if possible, it splits computation in smaller units. Otherwise, it may perform a processing step in the atomic piece of data and send it to a gatherer.
  - Gatherer: Receives pieces of data from scatterers, and combines them into a single piece of data performing

# Average and scatter-gather



· 47

- A problem for which this pattern is suitable is computing the average of a list of numbers
- Given a set of natural numbers  $a_1, a_2, \dots, a_n$ , the average is  $\frac{1}{n} \sum_i a_i$ 
  - Note that this is equivalent to  $\sum_i \frac{a_i}{n}$
- In a nutshell, we can have scatterer actors splitting computation and computing each factor  $\frac{a_i}{n}$ , and gatherers summing up the results

# Averager



1	2	3	4
---	---	---	---



· 48


Consider a system that computes  
the average of a list of numbers



# Averager



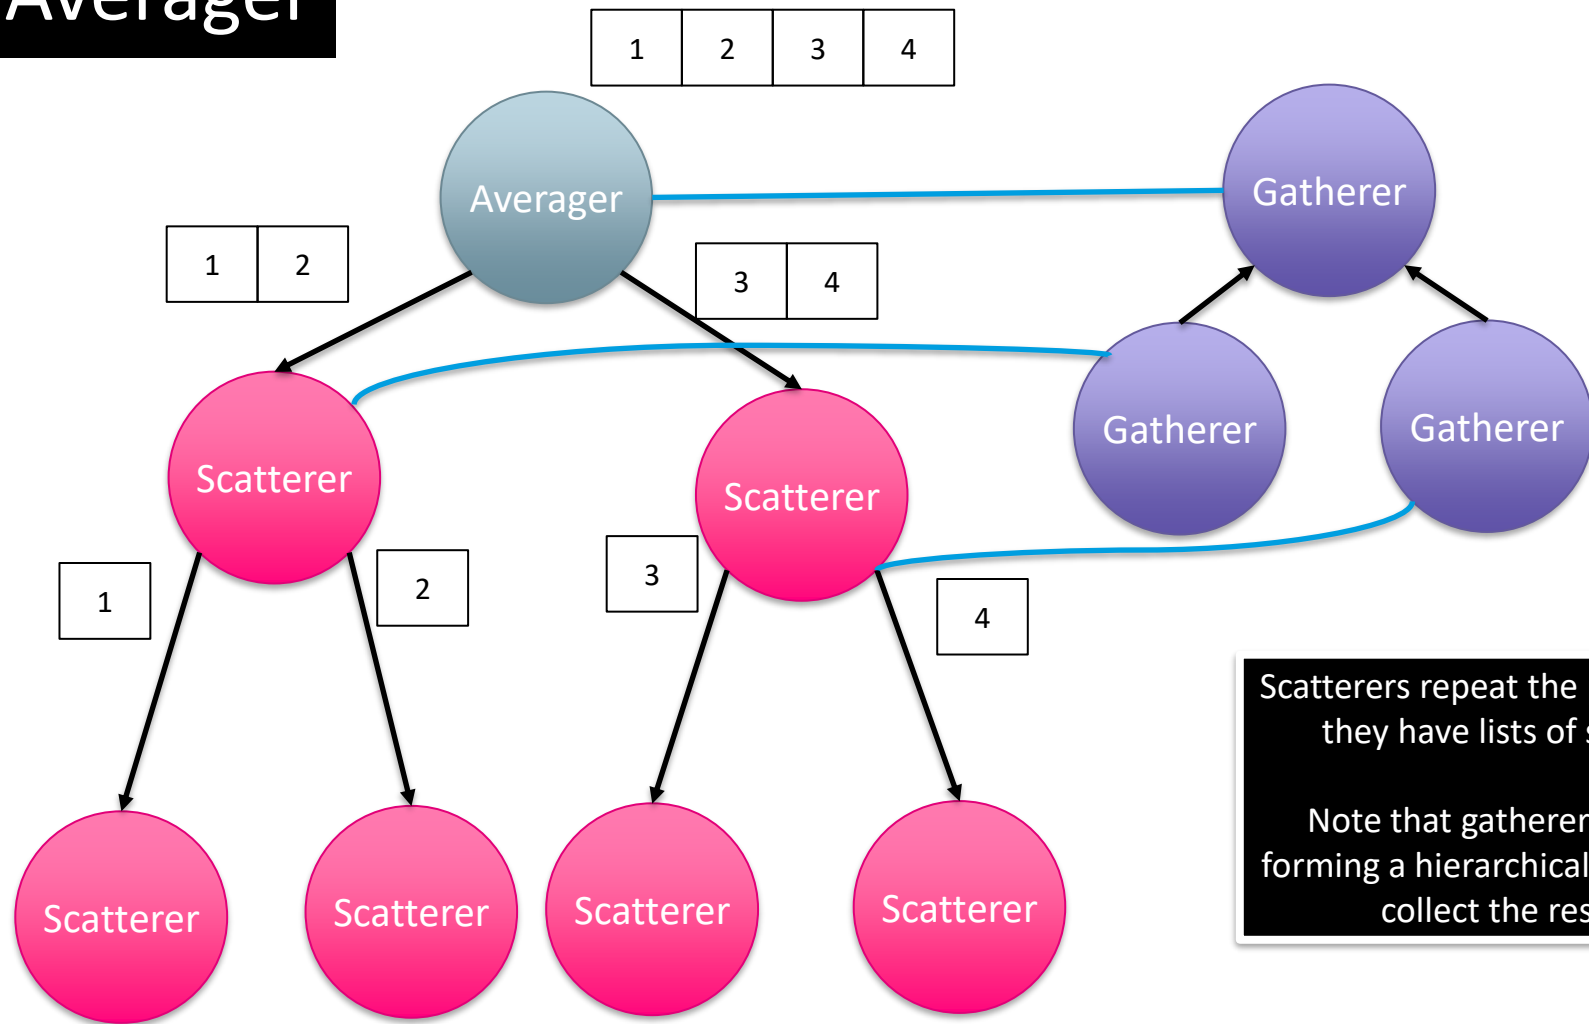
· 49



In the first step, we split the computation into two sublists, and assign them to separate scatterer workers.

Also, we spawn a gatherer worker that will receive merge the average of each sublist

# Averager



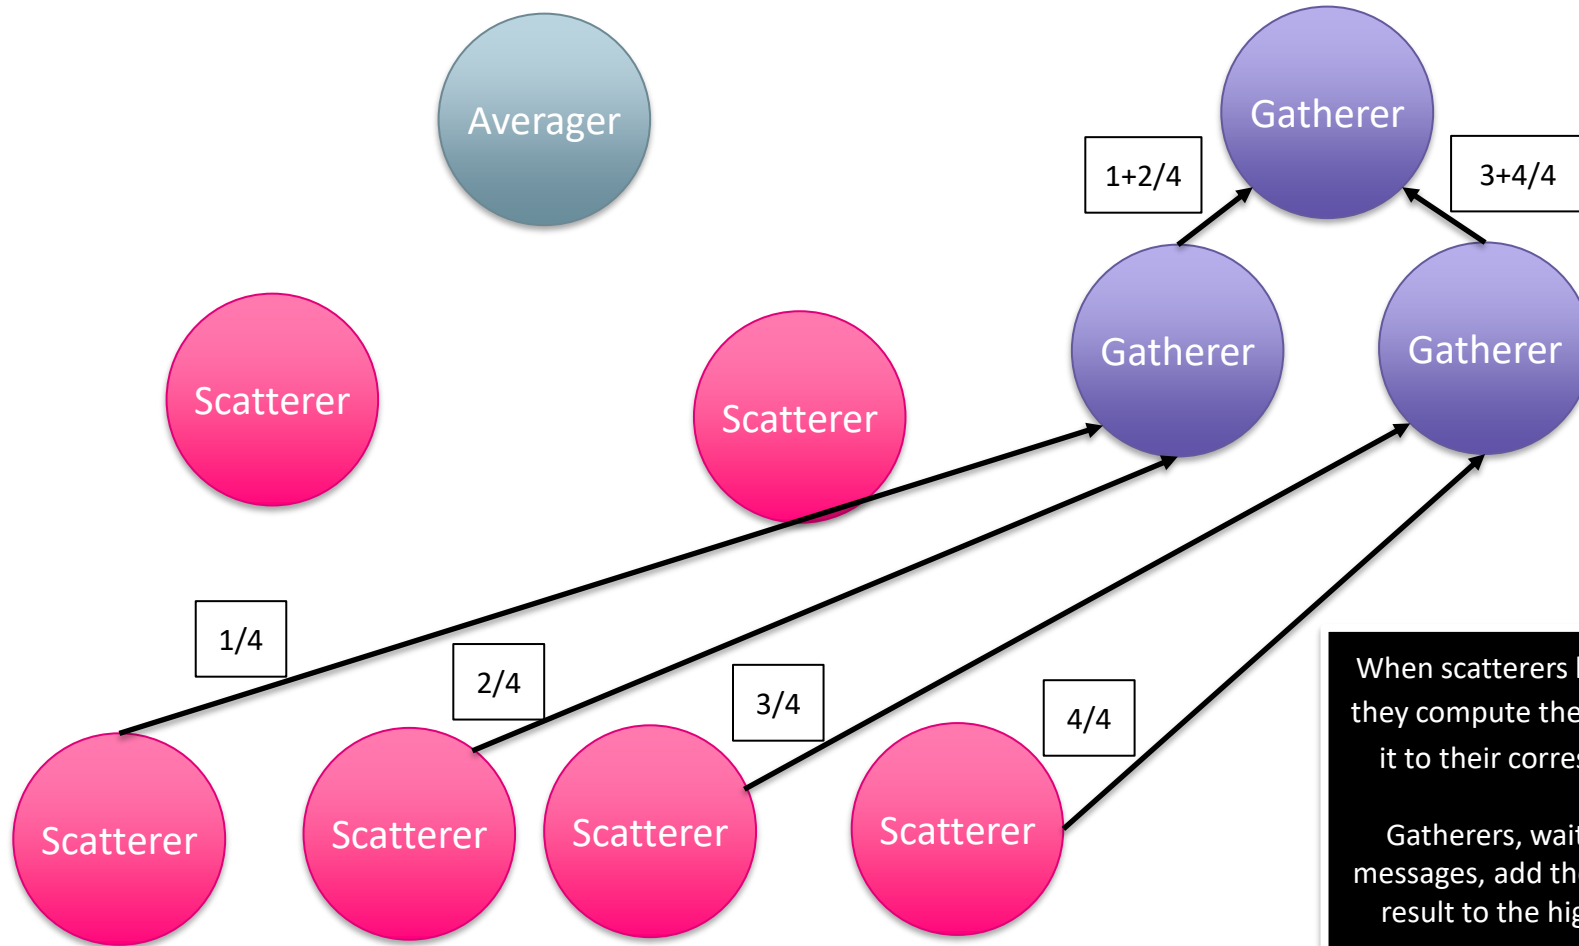
Scatterers repeat the process until they have lists of size one.

Note that gatherers are also forming a hierarchical structure to collect the results.

# Averager



· 51









When scatterers have lists of size one they compute the fraction  $\frac{a_i}{n}$  and send it to their corresponding gatherer.

Gatherers, wait for two scatterer messages, add them up, and send the result to the higher level gatherer

# Averager

Proof by example ☺

Output:  $1+2+3+4/4$



$1/4$

$2/4$

$3/4$

$4/4$

$1+2/4$

$3+4/4$

Finally, the top level gatherer adds up the two last partial results and outputs the complete result.

# Averager

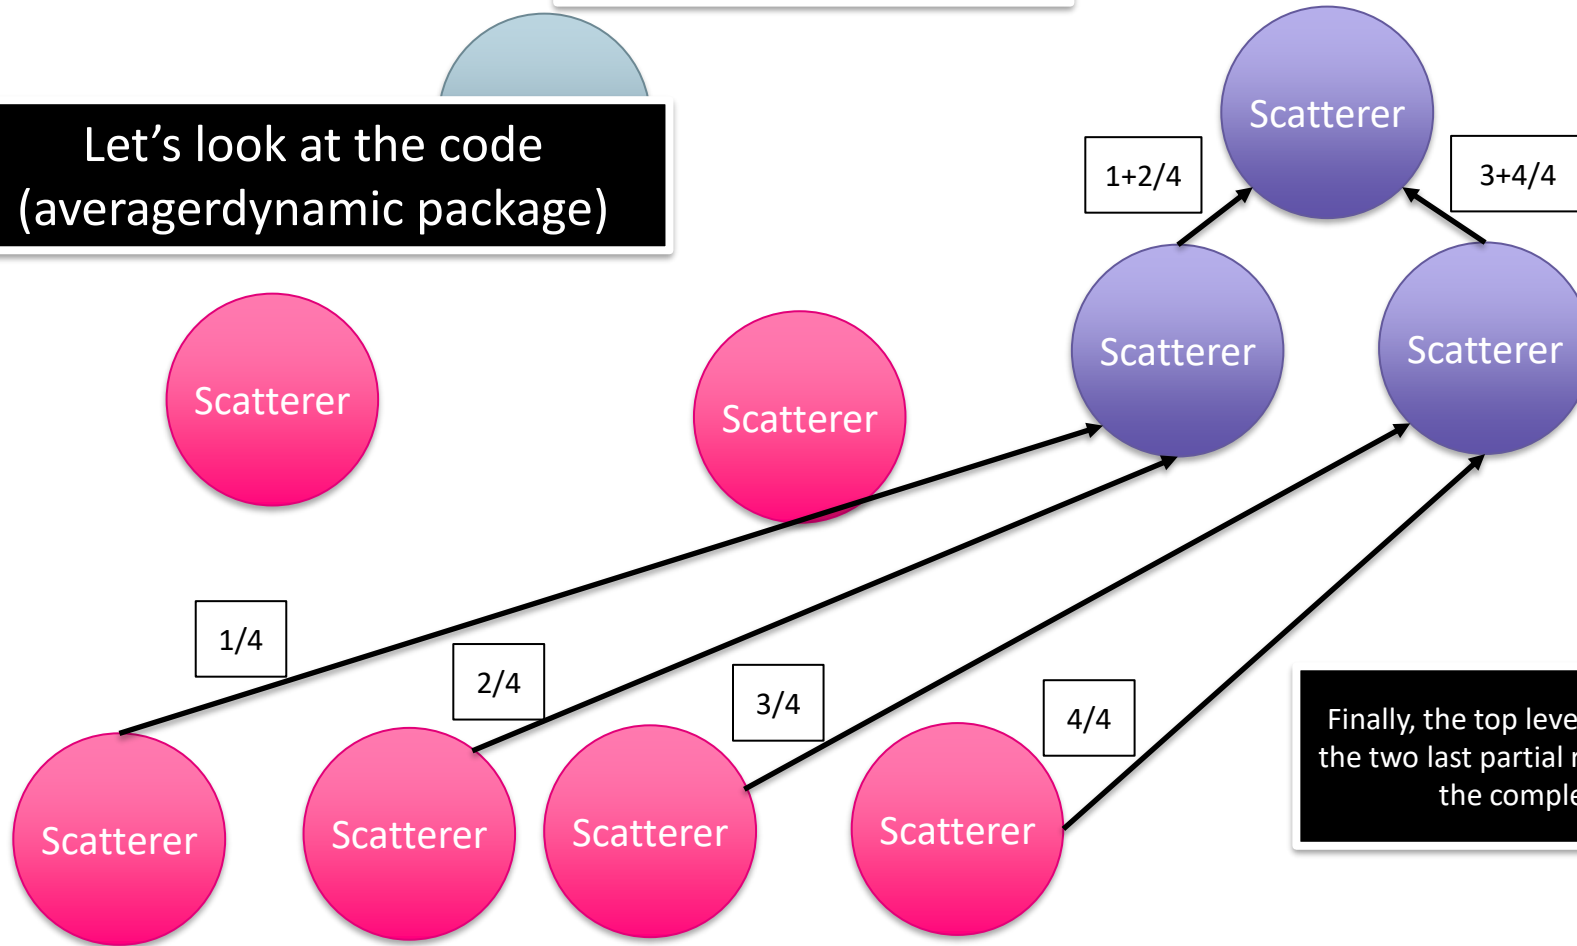
Proof by example ☺

Output:  $1+2+3+4/4$



· 52

Let's look at the code  
(averagerdynamic package)



Finally, the top level gatherer adds up the two last partial results and outputs the complete result.


- The size of the problem does not necessarily need to determine the distribution of computation
- One may have HW restrictions
  - As we saw, actors systems running in a single machine may not scale well beyond the number of processors
- Another example of adaptive load balancing are elastic systems
  - Elastic systems try to keep a number of active actors proportional to the workload
  - Several exercises for this week target implementing an elastic server

## Changing behaviour

# Actors with changing behaviour




- The actors model states that *“upon receiving a message an actor may change its behaviour”*
- So far we have considered change in behaviour as change in state



The actor receives a message and updates its internal state. The arrow ➡ models state transition



- The actors model states that *“upon receiving a message an actor may change its behaviour”*
- However, actors may also change the *functions* to process messages (i.e., message handlers)



In this example, the actor that was executing  $f()$  when a message comes, now executes  $g()$ . The new function  $g()$  could be completely different, e.g., changing how messages are processed or even waiting for different type of messages!

- In Akka, we can change the behaviour of an actor by defining a function that returns the new behaviour
  - Like in the behaviour defined in `createReceive()`
- In fact, we have already done this when we return `Behaviors.stopped()` to terminate an actor

- The packages `averagerdynamic` and `averagerbehavior` implement the same system, but the latter uses changing behaviour
- In `averagerdynamic`, we use a Boolean variable (`receivedFirstNumber`) to determine whether we have received the first or second `GathererCommand` message
- In `averagebehavior`, we define a new behaviour (`waitForSecond`) to which the actor transitions after receiving the first `GathererCommand` message.
  - In this way we can do without the Boolean variable mentioned above

# Averager with changing behaviour



- The packages `averagerdynamic` and `averagerbehavior` implement the same system, but the latter uses changing behaviour
- In `averagerdynamic`, we use a Boolean variable (`receivedFirstNumber`) to determine whether we have received the first or second `GathererCommand` message
- In `averagebehavior`, we define a new behaviour (`waitForSecond`) to which the actor transitions after receiving the first `GathererCommand` message.
  - In this way we can do without the Boolean variable mentioned above

Let's look at the code  
(`averagerbehaviors` package)

# Actors in distributed systems



The actors model has natural mapping in distributed systems


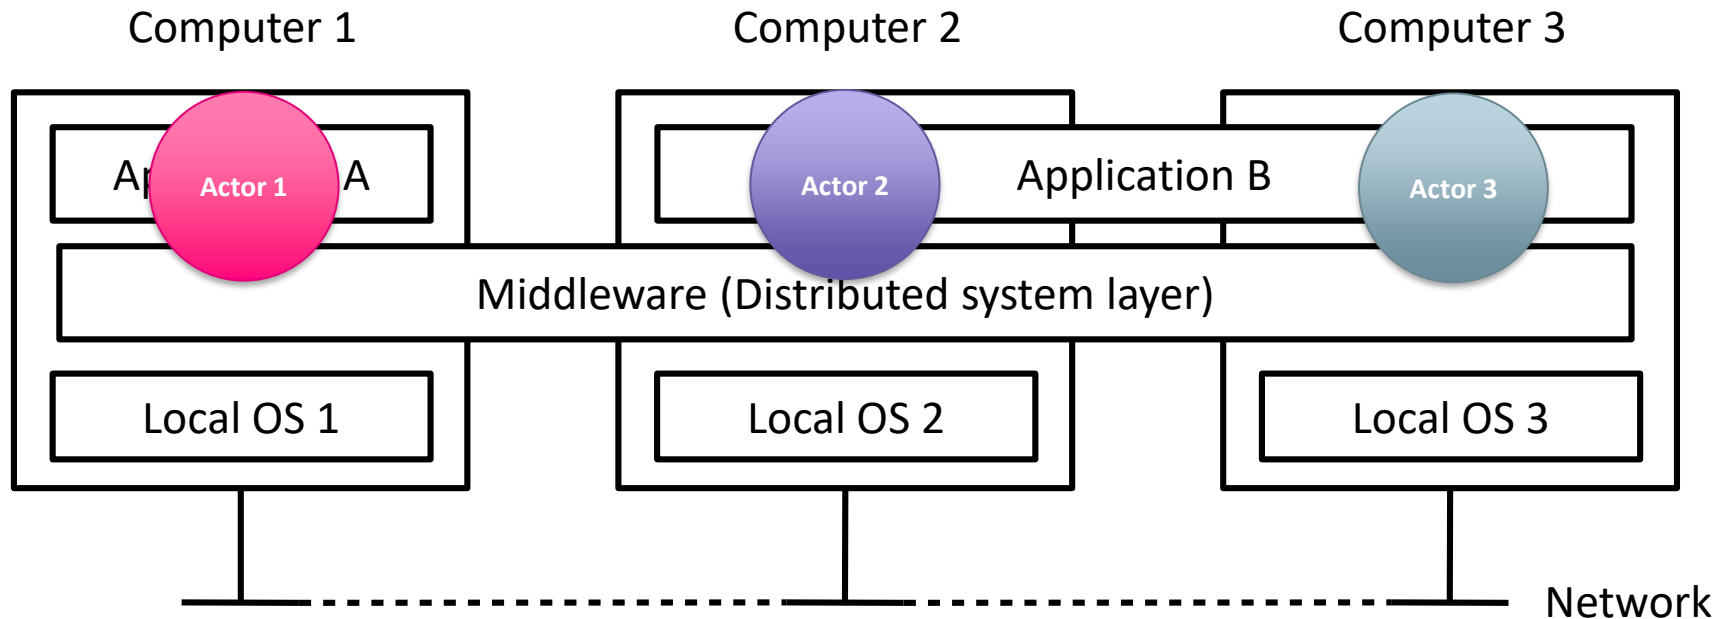


Figure taken from -> Distributed Systems: Principles and Paradigms. Andrew S. Tanenbaum and Maarten Van Steen. 2007.

# Actors in distributed systems




The actors model has natural mapping in distributed systems



# Actors in a single computer



The actors model is applicable in a single computer as well



In this course, we focus on this type of actor system

- Actors model (revisited)
  - Bounded Buffer
  - Primer
- Dynamic topology
- Fault-tolerance
  - Supervision
- Adaptive load balancing
  - Scatter-Gather
- Changing behaviour